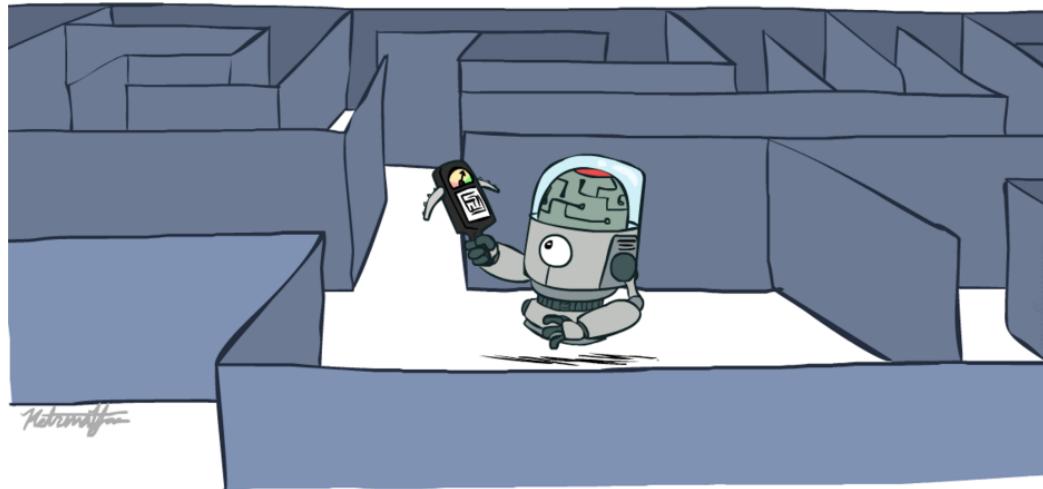


CS 188: Artificial Intelligence

Search Continued

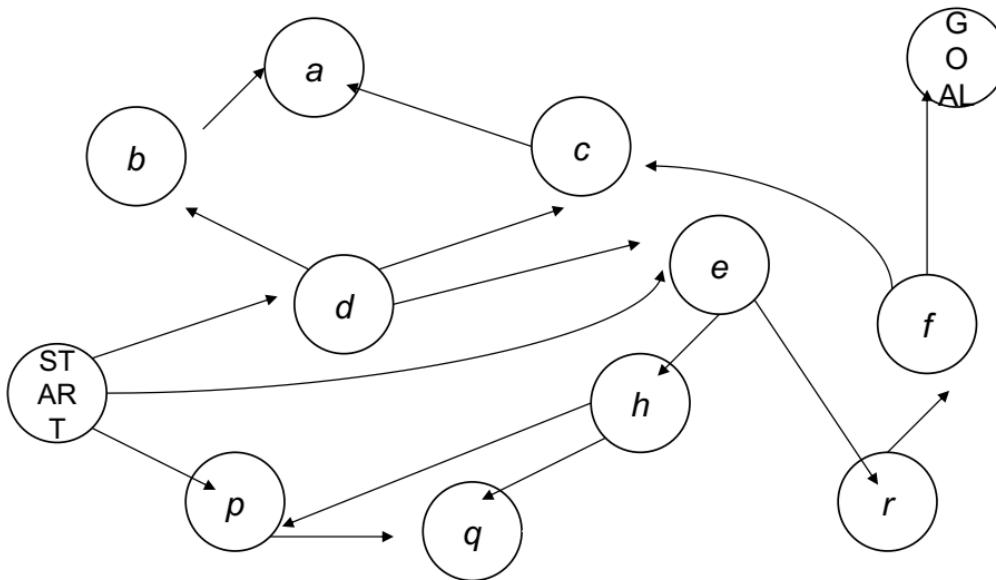


Instructors: Anca Dragan

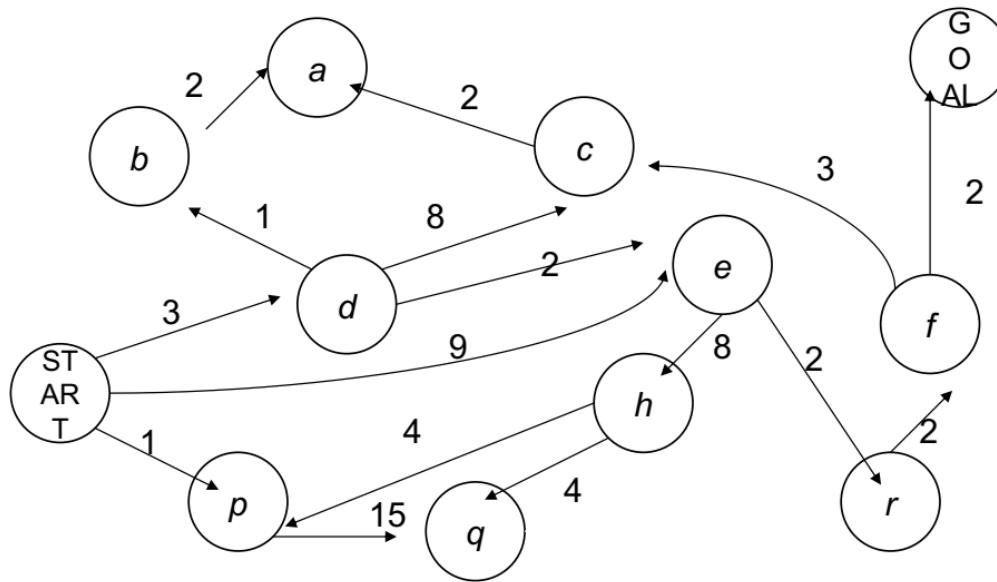
University of California, Berkeley

[These slides adapted from Dan Klein and Pieter Abbeel; ai.berkeley.edu]

Cost-Sensitive Search



Cost-Sensitive Search



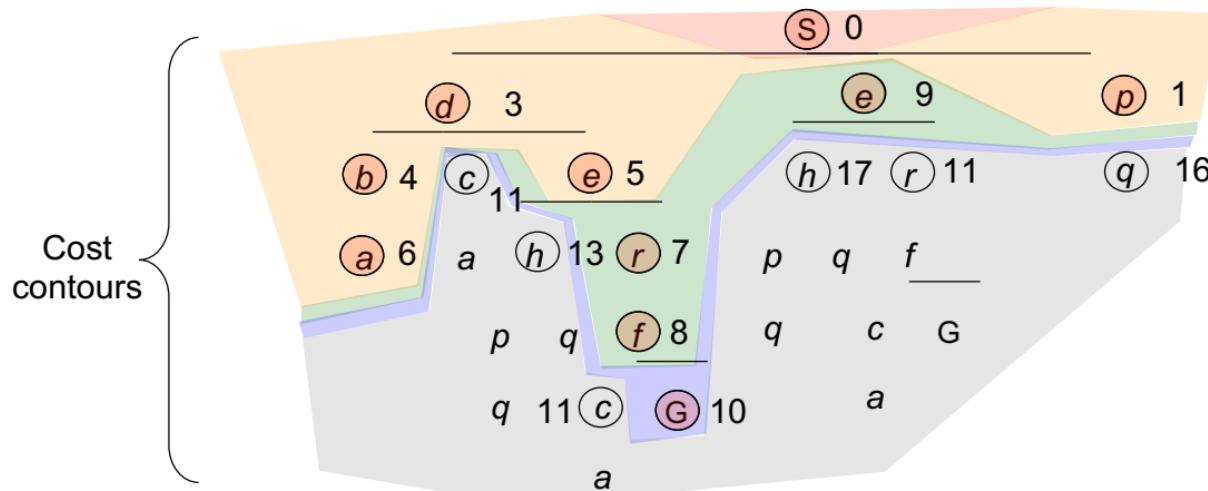
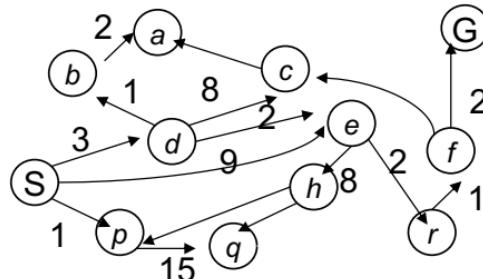
BFS finds the shortest path in terms of number of actions.
It does not find the least-cost path. We will now cover
a similar algorithm which does find the least-cost path.

How?

Uniform Cost Search

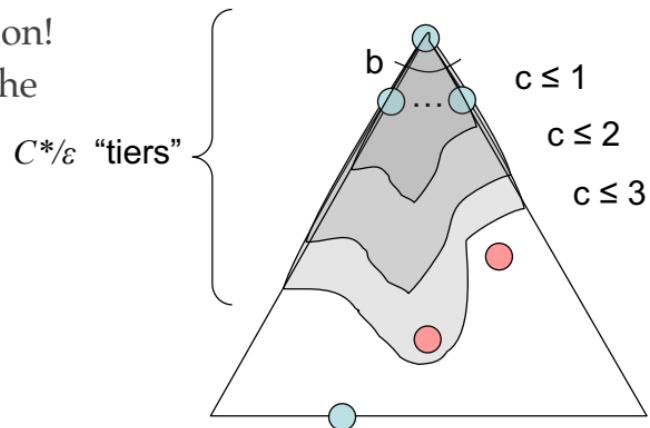
Strategy: expand a cheapest node first:

Fringe is a priority queue
(priority: cumulative cost)



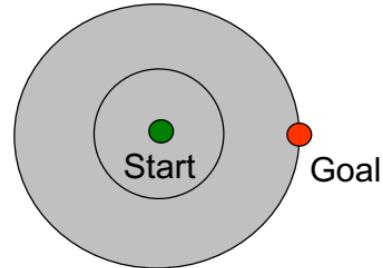
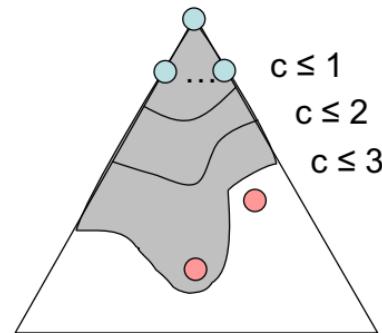
Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
 - Processes all nodes with cost less than cheapest solution!
 - If that solution costs C^* and arcs cost at least ε , then the “effective depth” is roughly C^*/ε
 - Takes time $O(b^{C^*/\varepsilon})$ (exponential in effective depth)
- How much space does the fringe take?
 - Has roughly the last tier, so $O(b^{C^*/\varepsilon})$
- Is it complete?
 - Assuming best solution has a finite cost and minimum arc cost is positive, yes! (if no solution, still need depth $\neq \infty$)
- Is it optimal?
 - Yes! (Proof via A*)



Uniform Cost Issues

- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
 - Explores options in every “direction”
 - No information about goal location
- We'll fix that soon!

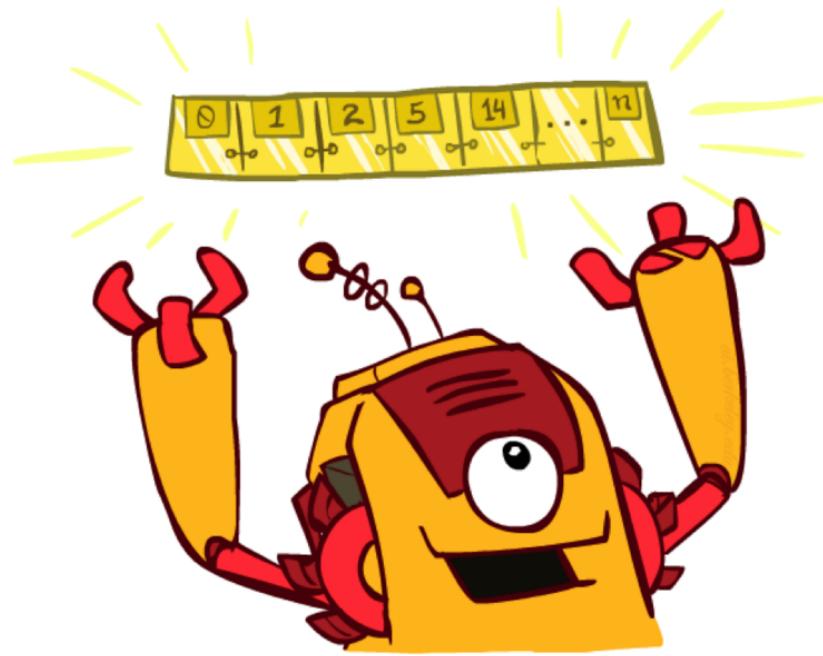


[Demo: empty grid UCS (L2D5)]

[Demo: maze with deep/shallow water DFS/BFS/UCS (L2D7)]

The One Queue

- All these search algorithms are the same except for fringe strategies
 - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
 - Practically, for DFS and BFS, you can avoid the $\log(n)$ overhead from an actual priority queue, by using stacks and queues
 - Can even code one implementation that takes a variable queuing object



???

Up next: Informed Search

- Uninformed Search
 - DFS
 - BFS
 - UCS



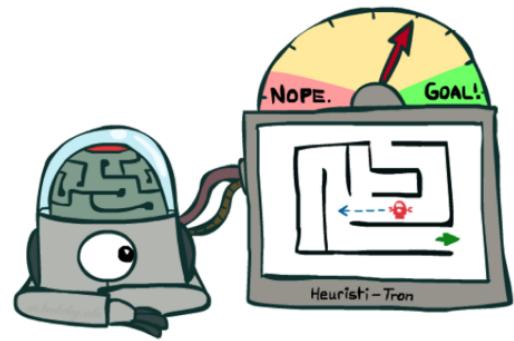
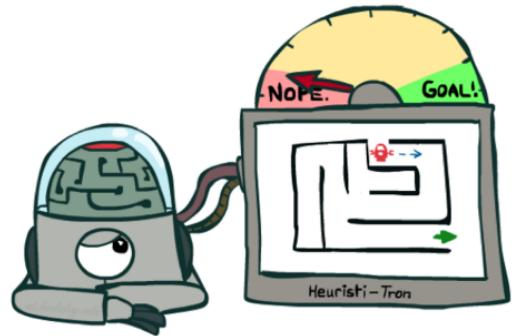
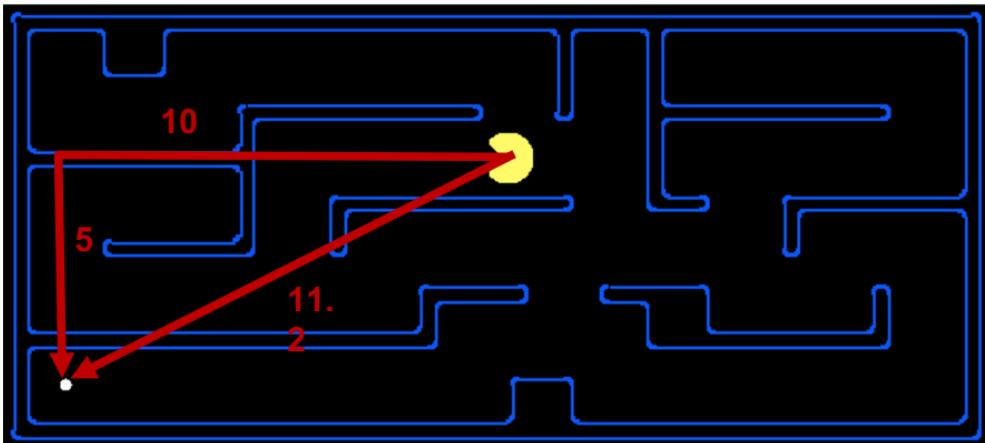
- Informed Search
 - Heuristics
 - Greedy Search
 - A* Search
 - Graph Search



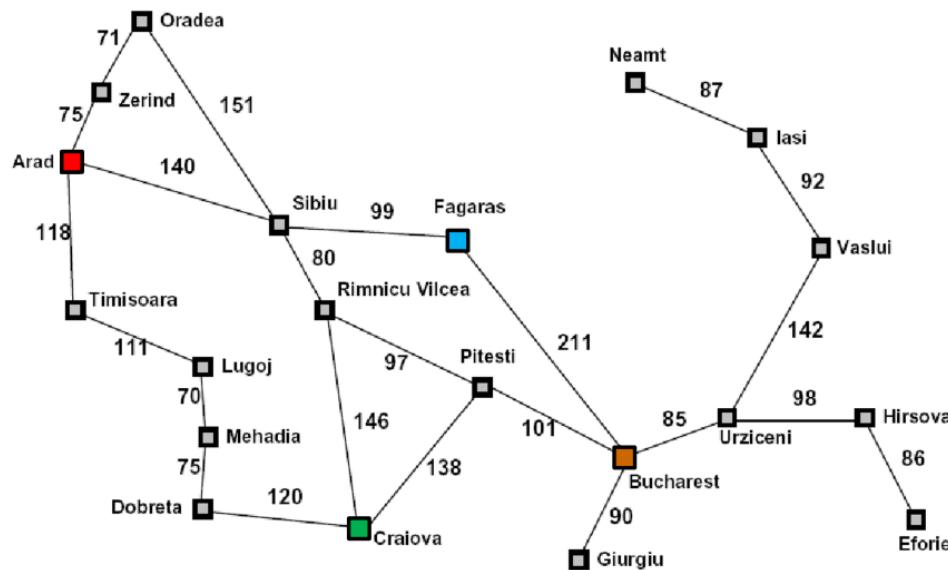
Search Heuristics

h - function

- A heuristic is:
 - A function that *estimates* how close a state is to a goal
 - Designed for a particular search problem
 - Pathing?
 - Examples: Manhattan distance, Euclidean distance for pathing



Example: Heuristic Function



Straight-line distance to Bucharest	
Arad	36
Bucharest	
Craiova	16
Dobreta	24
Eforie	16
Fagaras	17
Giurgiu	7
Hirsova	15
Iasi	22
Lugoj	24
Mehadia	24
Neamt	23
Oradea	38
Pitesti	9
Rimnicu Vilcea	19
Sibiu	25
Timisoara	32
Urziceni	8
Vaslui	19
Zerind	37

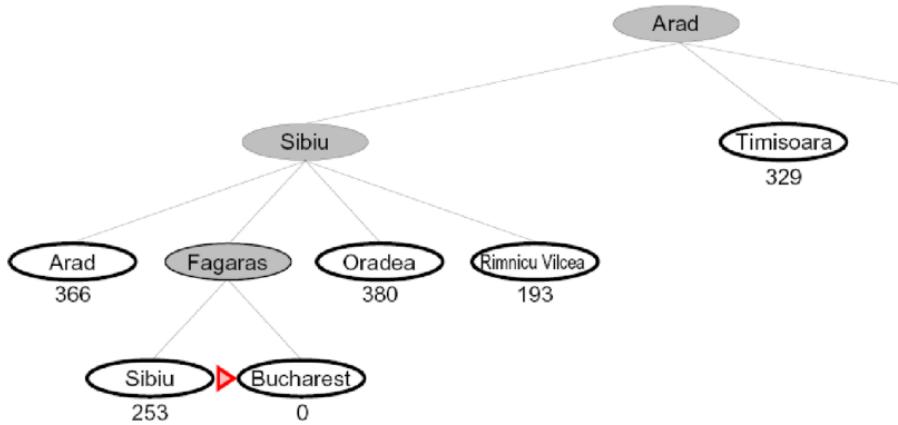
$h(x)$

Greedy Search

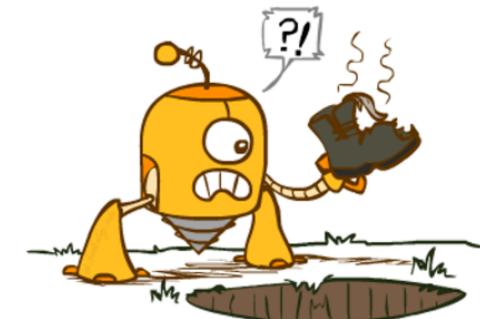
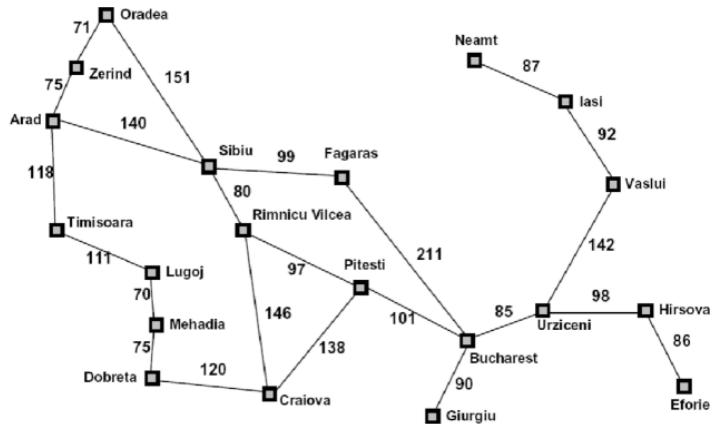


Greedy Search

- Expand the node that seems closest..

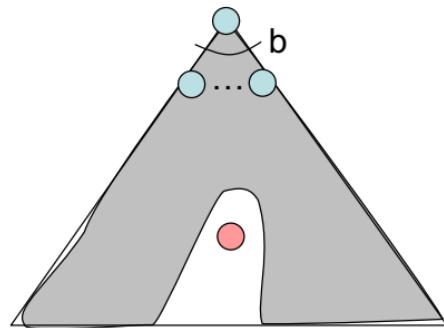
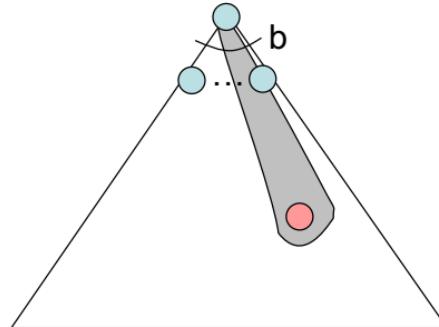


- Is it optimal?
 - No. Resulting path to Bucharest is not the shortest!



Greedy Search

- Strategy: expand a node that you think is closest to a goal state
 - Heuristic: estimate of distance to nearest goal for each state
- A common case:
 - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS



[Demo: contours greedy empty (L3D1)]

[Demo: contours greedy pacman small maze (L3D4)]

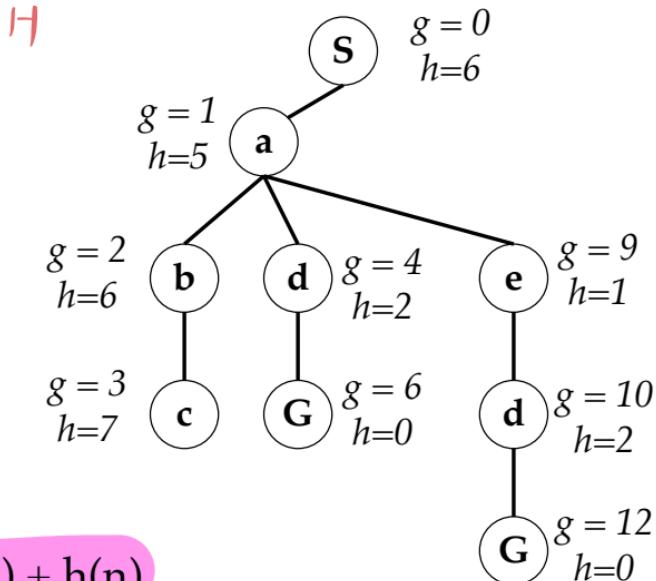
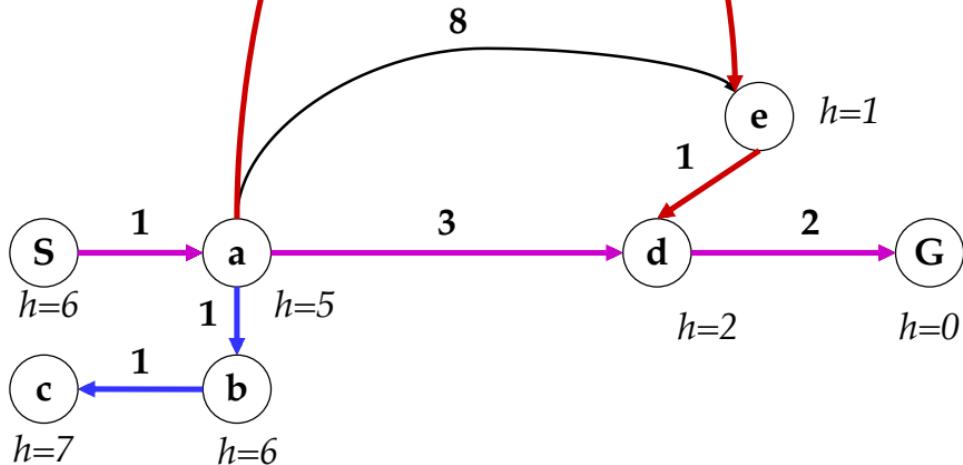
A* Search



Combining UCS and Greedy

are cost

- Uniform-cost orders by path cost, or *backward cost* $g(n)$
- Greedy orders by goal proximity, or *forward cost* $h(n)$



- A* Search orders by the sum: $f(n) = g(n) + h(n)$

Example: Teg

Algorithm Description

Suppose we are finding the shortest path from vertex a to a vertex z

The A* search algorithm initially:

- Marks each vertex as unvisited
- Starts with a priority queue containing only the initial vertex a
 - The priority of any vertex v in the queue is the weight $w(v)$ which assumes we have found the shortest path to v (initialize it to be infinity except for the initial vertex a)
 - Shortest weights have highest priority
- For each vertex v , $d(a, v)$ is the shortest known distance from a to v , $d(a, a) = 0$ and $d(a, v) = \text{infinity}$ for all $v \neq a$
- For each vertex v , $h(v, z)$ is the heuristic distance from v to z

估值
 $f(n) = g(n) + h(n)$

当前已到目标 (已达)
(各种距离)
从 n 到 z / $f(n)$ 里数

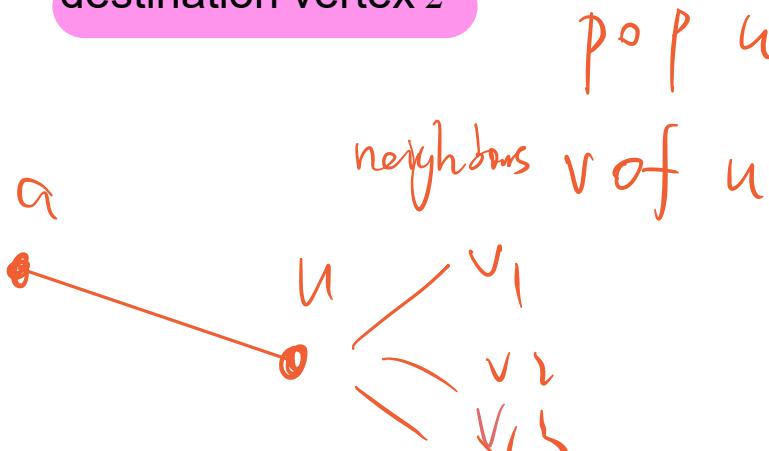
Algorithm Description I (Tree Search)

The algorithm then iterates:

- Pop the vertex u with highest priority
- For each adjacent vertex (neighbor) v of u :
 - If $w(v) = d(a, u) + d(u, v) + h(v, z)$ is less than the current weight/priority of v , update the path leading to v and its priority
 - If v is not in the queue, push v into the queue

Update Tree weight priority
TAK TAK

Continue iterating until the item popped from the priority queue is the destination vertex z



Algorithm Description II (*Graph Search*)

Visited / Unvisited

The algorithm then iterates:

- Pop the vertex u with highest priority
 - mark u as visited
- For each unvisited adjacent vertex (neighbor) v of u :
 - If $w(v) = d(a, u) + d(u, v) + h(v, z)$ is less than the current weight/priority of v , update the path leading to v and its priority
 - If v is not in the queue, push v into the queue

Continue iterating until the item popped from the priority queue is the destination vertex z

1. 首先需要创建两个集合，一个存储待访问的节点 (nodeLists)，一个存储已经访问过的节点 (visitedNodeLists)

2. 添加起点 (2,1) 到 nodeLists 列表，并且计算该点的预估值。

Weight 距离 + 估价
这里就需要用到上面讲到的预估值函数了，

$$f(n) = g(n) + h(n)$$

我们可以很清楚的看到，起点到当前节点的距离为0，即： $g(n) = 0$ 。

而当前顶点 (2,1) 到目标顶点 (2,5) 的距离为4，即： $h(n) = 4$ 。

这里的h值可以采用曼哈顿距离公式计算得来：

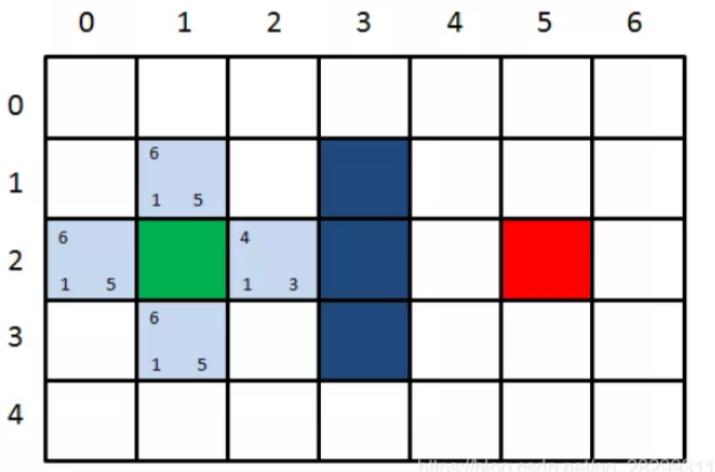
$$h = |x_1 - x_2| + |y_1 - y_2|$$

这里需要注意的是：h 值是在不考虑障碍的情况下计算的

3. 查找nodeLists 里预估值最小的节点，作为当前访问的节点，并且从nodeLists 删除该

4. 获取当前节点的邻居节点，计算出他们的预估值 并且添加到nodeLists列表中。

如图：



function startSearch() {
 // 把起点添加到nodeLists list
 nodeLists.push(startNode);
 // 循环
 while (nodeLists.length > 0) {
 // 查找预估值最小的节点
 let currentNode = findMinNode();
 // 获取并且添加邻居节点到待访问列表
 findAndAddNeighborNode(currentNode);
 // 把当前节点添加到已访问列表
 visitedNodeLists.push(currentNode);
 // 判断是否是目标节点
 if (this.isTarget(currentNode)) {
 return currentNode
 }
 }
 return null
}

这里只考虑上下左右4个方位，并且分别计算了他们的各个预估值，

右下角： $h(n)$ 值

左下角： $g(n)$ 值

左上角： $f(n)$ 值

除了计算预估值，我们还需要把当前节点作为每个邻居节点的父节点，以便为了后面确定最终的路线。

这里需要注意的是：我们在添加到待访问列表之前需要处理一下边界问题，并且判断一下是否是障碍物，如果是就不需要添加到列表中。

5. 把当前节点添加到 visitedNodeLists 中，代表已经访问过了。

6. 重复以上步骤 3 - 5，直到找到目标节点位置为止。

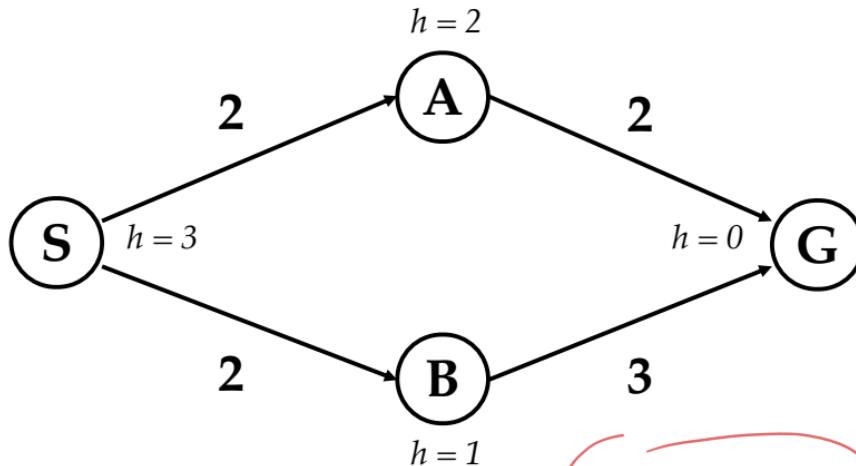
7. 循环输出最终节点的父节点，就是我们需要的路径了。



end 2. 从 start 到 target

When should A* terminate?

- Should we stop when we enqueue a goal?

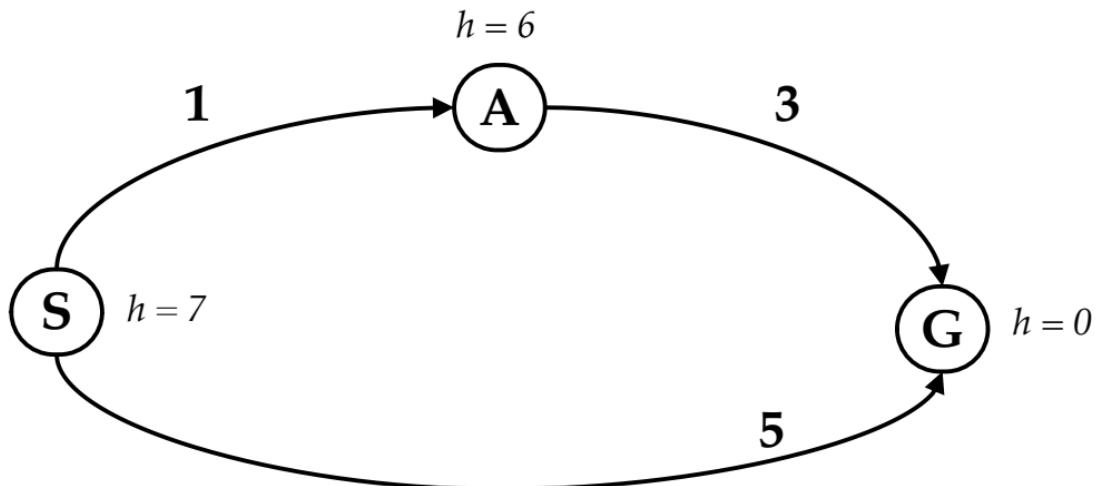


- No: only stop when we dequeue a goal

g	h	+
S	0	3
S->A	2	2
S->B	2	1
S->B->G	5	0
S->A->G	4	0

Annotations:
Row 5: $h=0$ circled in pink, labeled "not deg".
Row 6: Boxed, labeled "4 0 4".

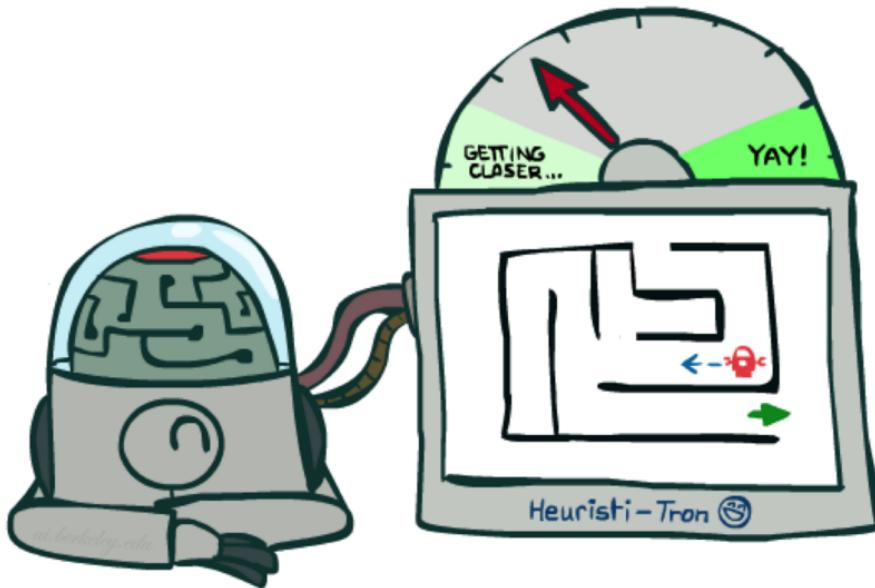
Is A* Optimal?



g	h	$+$
S	0	7
S->A	1	6
S->G	5	0

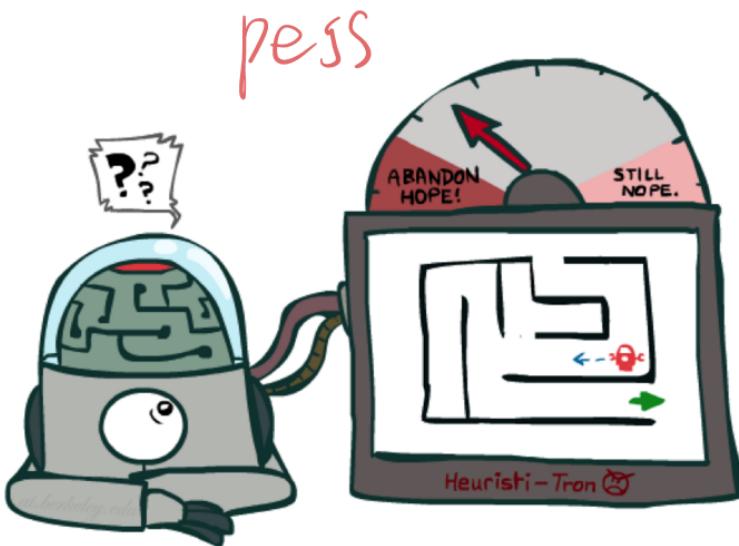
- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!

Admissible Heuristics

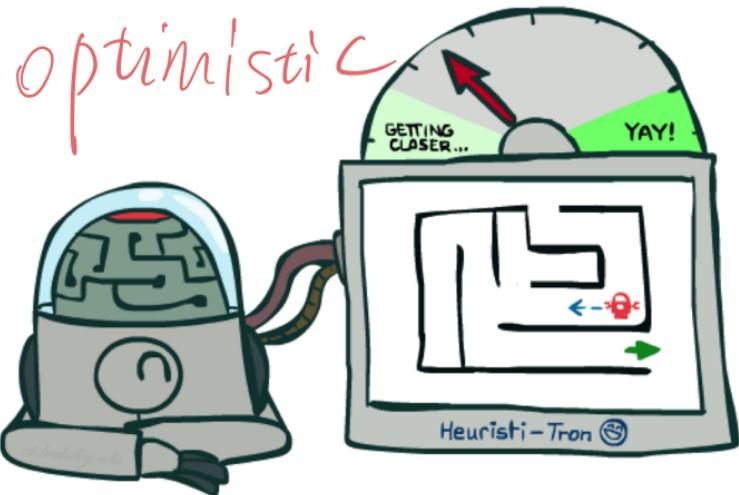


可接返生

Idea: Admissibility acceptable?



Inadmissible (pessimistic) heuristics
break optimality by trapping
good plans on the fringe



Admissible (optimistic) heuristics
slow down bad plans but
never outweigh true costs

Admissible Heuristics

Pretend < actual

Admissible heuristics h must always be optimistic:

- Let $d(u, v)$ represent the actual shortest distance from u to v
- A heuristic $h(u, v)$ is **admissible** if $h(u, v) \leq d(u, v)$
- The heuristic is *optimistic* or a *lower bound* on the distance

Using the Euclidean distance between two points on a map is clearly an admissible heuristic

- The flight of the crow is shorter than the run of the wolf

A problem with fewer restrictions on the actions is called a **relaxed problem**

- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem

Theorem: If $h(n)$ is admissible, A* using TREE-SEARCH is optimal

Consistent \rightarrow non-decreasing

Consistent Heuristics

\hookrightarrow Graph

一致者

A heuristic is **consistent** if for every node n , every successor n' of n generated by any action a , we have

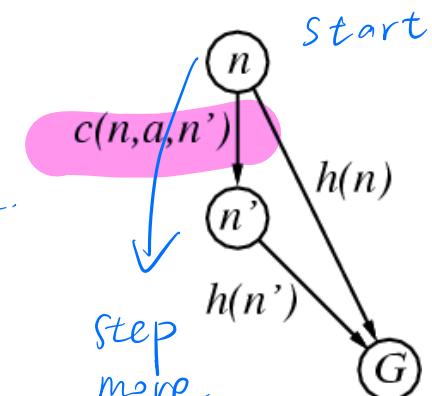
$$h(n) \leq c(n, a, n') + h(n')$$

consistent
admissible.

If h is consistent, we have

$$\begin{aligned} w(n') &= d(n') + h(n') && \text{(by def.)} \\ &= d(n) + c(n, a, n') + h(n') && (d(n') = d(n) + c(n, a, n')) \\ &\geq d(n) + h(n) = w(n) && \text{(consistency)} \\ w(n') &\geq w(n) \end{aligned}$$

i.e., $w(n)$ is non-decreasing along any path.



keeps all checked nodes
in memory to avoid repeated states

tree-search

Theorem:

If $h(n)$ is consistent, A* using GRAPH-SEARCH is optimal

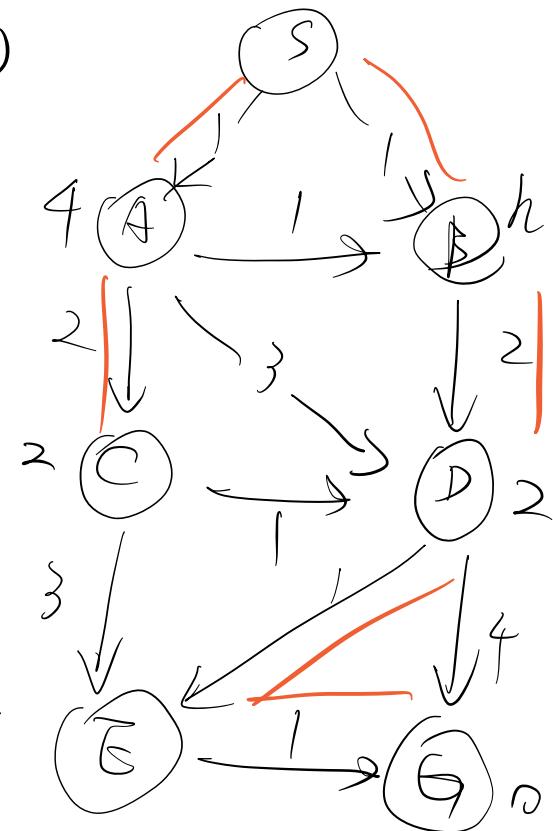
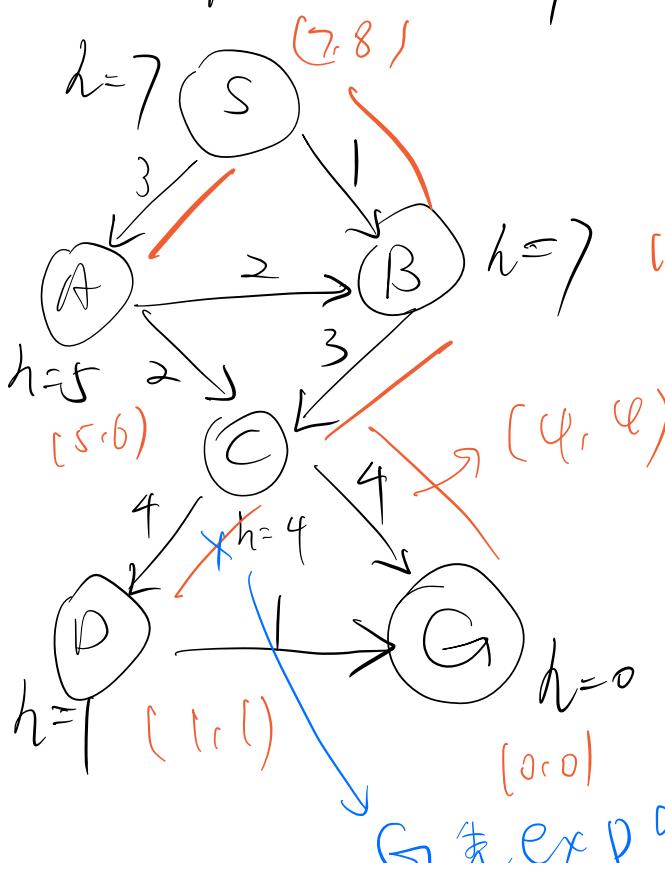
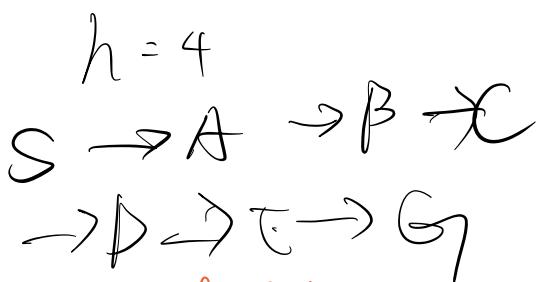
$$h(n) \leq \underbrace{c(n; g, n')}_{\sim} + h(n')$$

$$h \leq 2+2.$$

$$4 \leq 1+h$$

$$5 \leq 1+h$$

start



ad: $h \leq d_c$ ✓

cons: $g \leq d_c$, ok

↪ * . ex D and, D $\overline{\text{tr}} \leq 3$

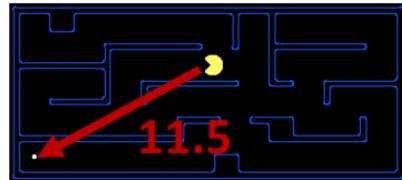
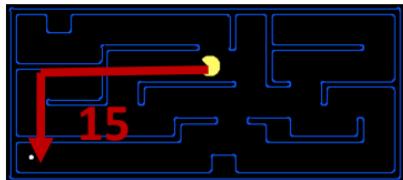
Admissible Heuristics

- A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

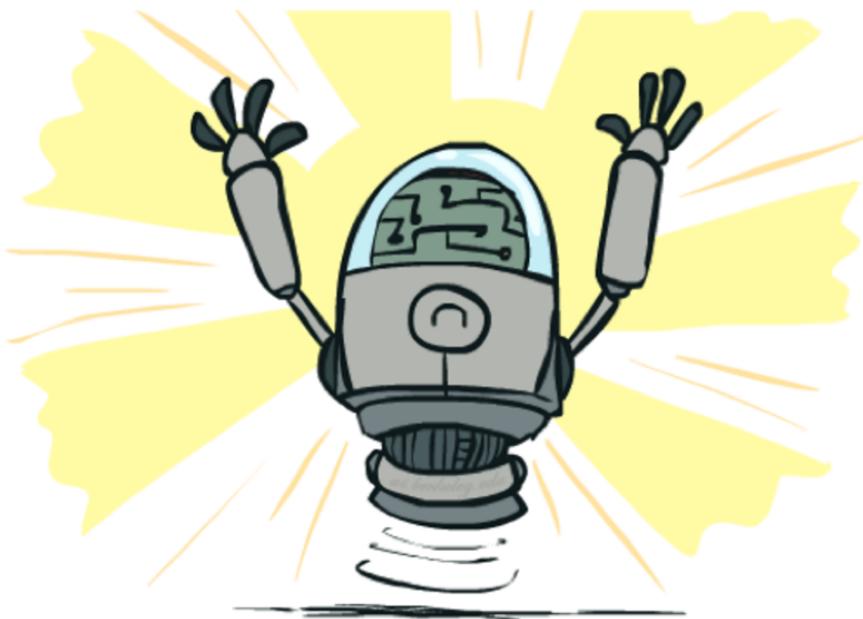
- Examples:



0.0

- Coming up with admissible heuristics is most of what's involved in using A* in practice.

Optimality of A* Tree Search



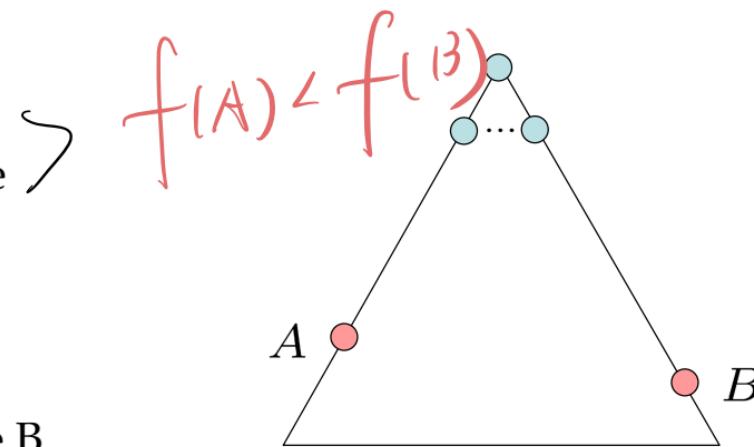
Optimality of A* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:

- A will exit the fringe before B

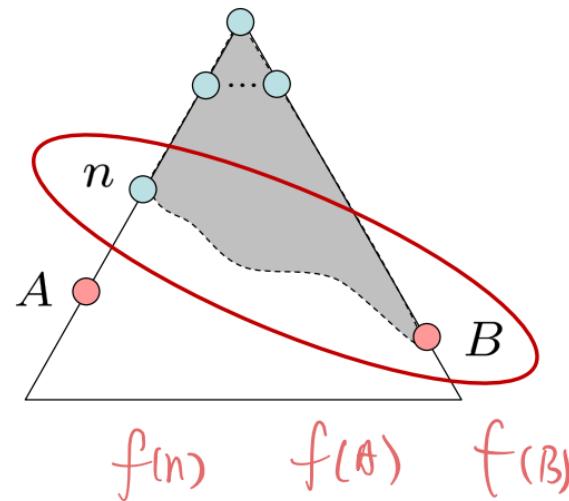


admissible \rightarrow find optimal
faster than suboptimal

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$
 3. n expands before B
- All ancestors of A expand before B
- A expands before B
- A* search is optimal

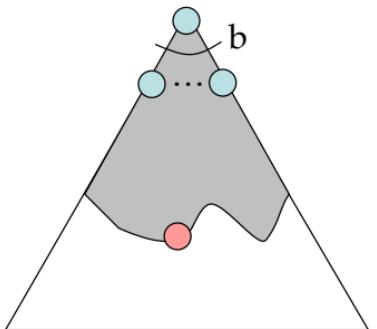


未遍
all A's ancestor

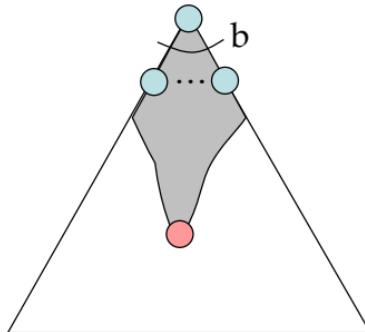
Properties of A*

expanded
 $\rightarrow A^*$ expanded

Uniform-Cost

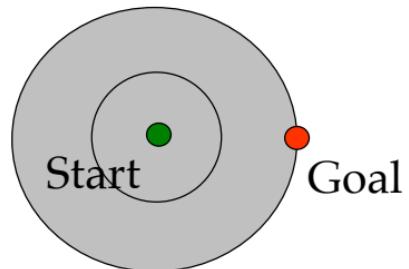


A^*

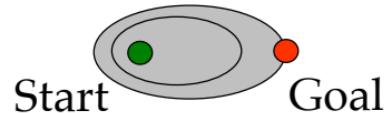


UCS vs A* Contours

- Uniform-cost expands equally in all “directions”



- A* expands mainly toward the goal, but does hedge its bets to ensure optimality *(Greedy) — [A]*



[Demo: contours UCS / greedy / A* empty (L3D1)]
[Demo: contours A* pacman small maze (L3D5)]

Comparison



Greedy



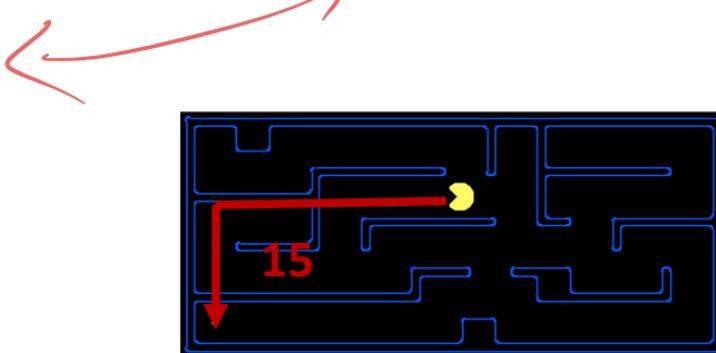
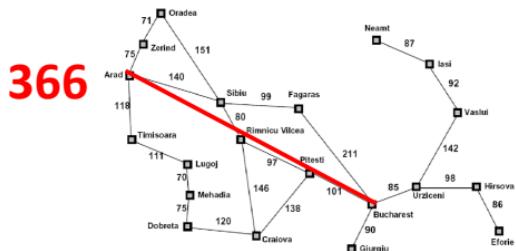
Uniform Cost



A*

Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available

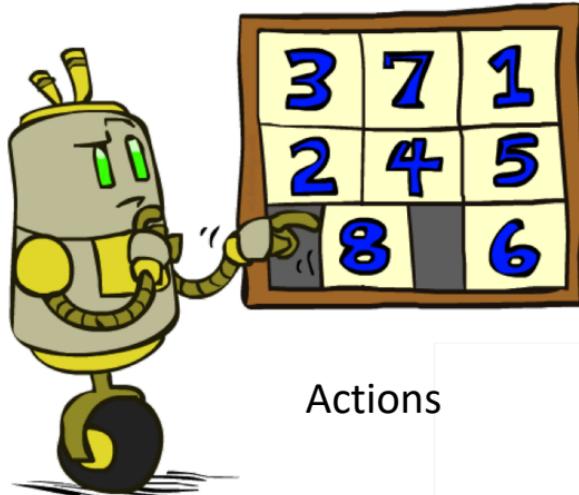


- Inadmissible heuristics are often useful too

Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



Actions

3	7	1
2	4	5
8	6	

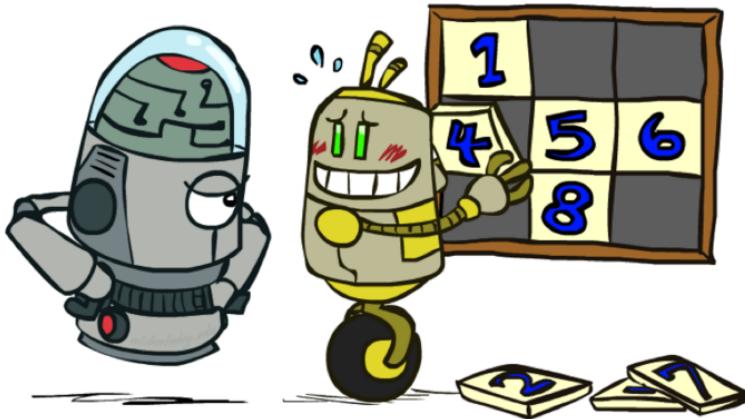
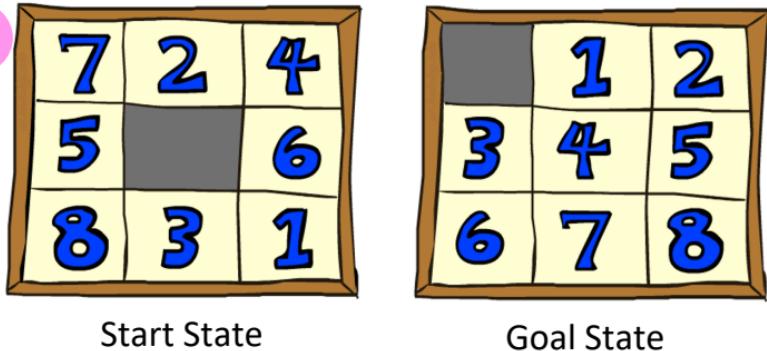
Goal State

- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

Admissible
heuristics?

8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) = 8$
- This is a *relaxed-problem* heuristic



Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	3.6×10^6
TILES	13	39	227

8 Puzzle II

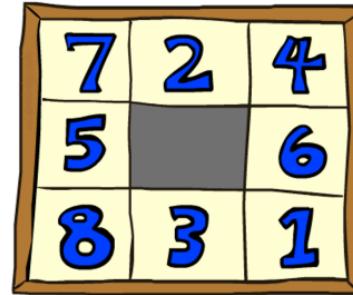
- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total *Manhattan* distance
- Why is it admissible?

$$h(\text{start}) = 3 + 1 + 2 + \dots = 18$$

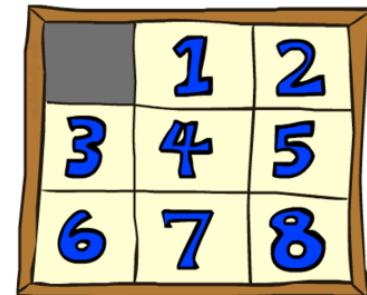
Euclidean

$$\sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$$

$$L_m \left(|x_1 - x_2|^m + |y_1 - y_2|^m \right)$$



Start State



Goal State

$$|x_1 - x_2| + |y_1 - y_2|$$

Average nodes expanded
when the optimal path has...

	...4 steps	...8 steps	...12 steps
TILES	13	39	227
MANHATTA N	12	25	73

8 Puzzle III

- How about using the *actual cost* as a heuristic?

- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?



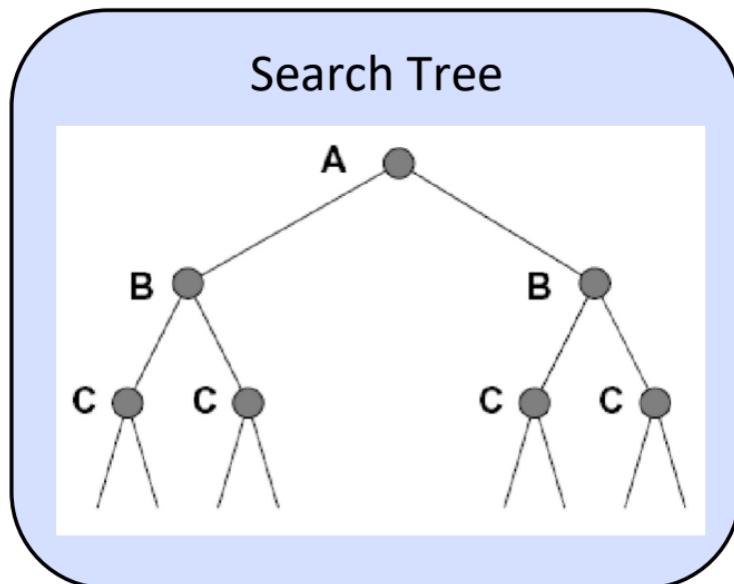
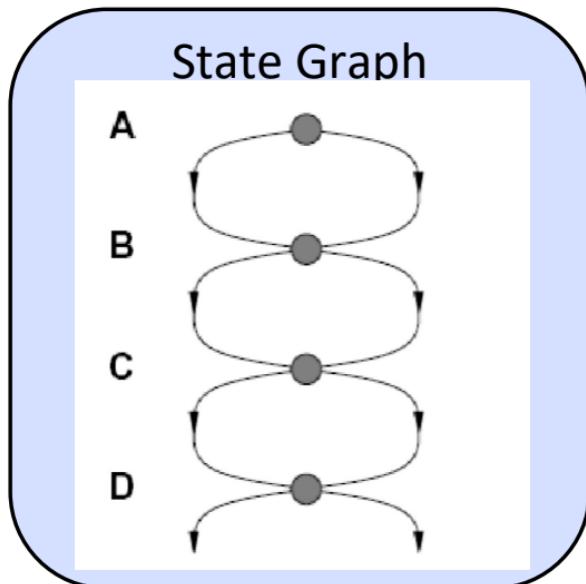
- With A*: a trade-off between quality of estimate and work per node

- As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

of compute
Heur

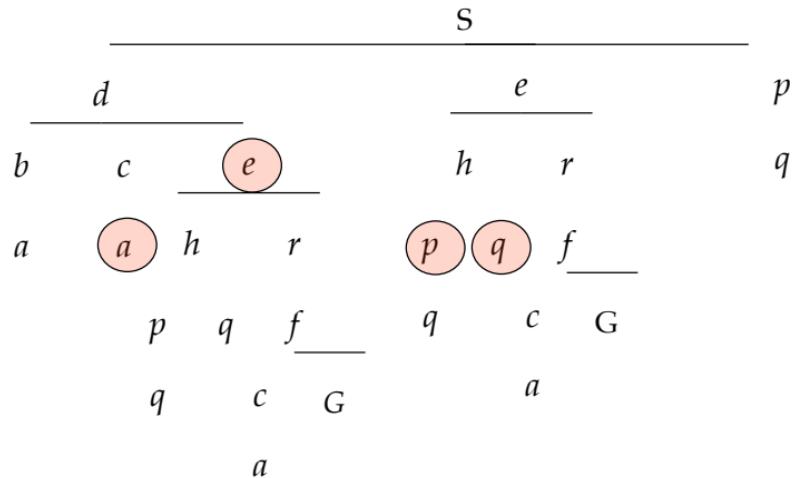
Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work.



Graph Search

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)

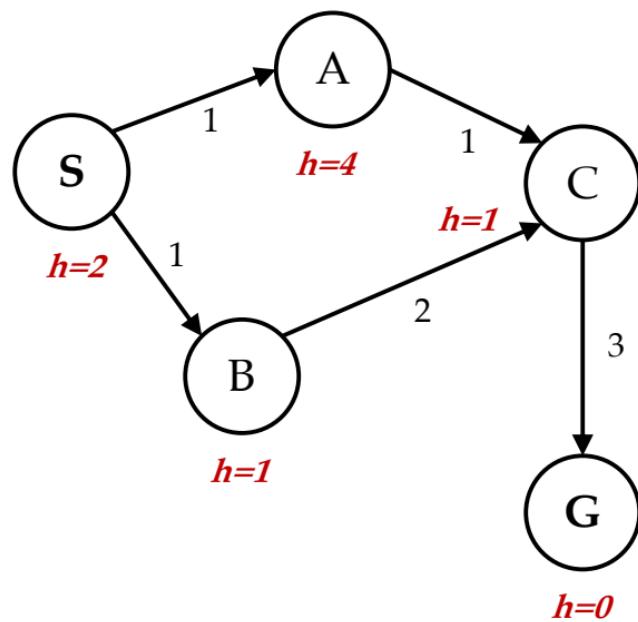


Graph Search

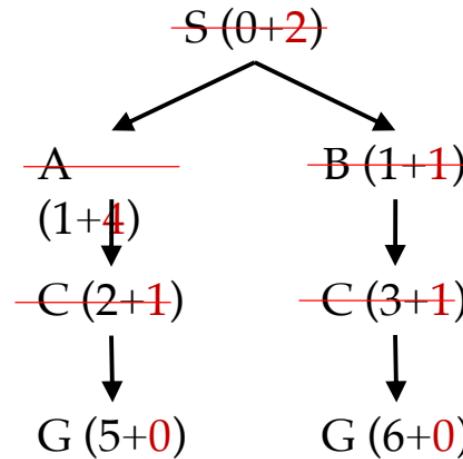
- Idea: never **expand** a state twice
- How to implement:
(visited tag)
 - Tree search + set of expanded states ("closed set")
 - Expand the search tree node-by-node, but...
 - Before expanding a node, check to make sure its state has never been expanded before
 - If not new, skip it, if new add to closed set
- Important: **store the closed set as a set**, not a list
not a list
- Can graph search ~~wreck~~ completeness? Why/why not?
- How about optimality?

A* Graph Search Gone Wrong?

State space graph

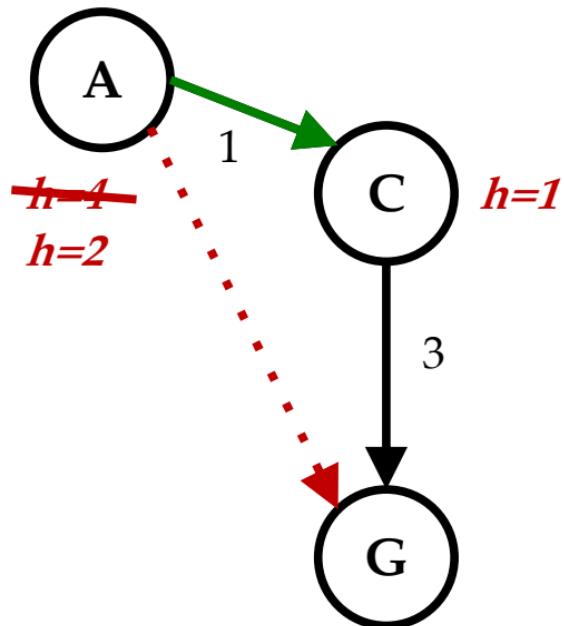


Search tree



Closed Set: S B C A

Consistency of Heuristics



- Main idea: estimated heuristic costs \leq actual costs
 - Admissibility: heuristic cost \leq actual cost to goal
$$h(A) \leq \text{actual cost from A to G}$$
 - Consistency: heuristic “arc” cost \leq actual cost for each arc
$$h(A) - h(C) \leq \text{cost}(A \text{ to } C)$$
- Consequences of consistency:
 - The f value along a path never decreases
$$h(A) \leq \text{cost}(A \text{ to } C) + h(C)$$
 - A* graph search is optimal

Consistent \rightarrow non-decreasing

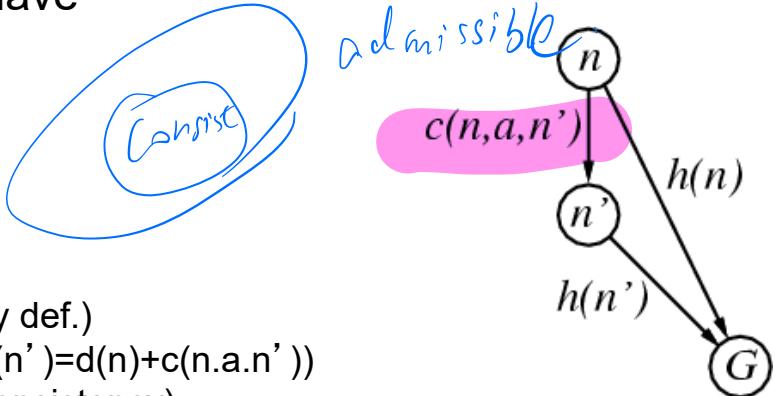
Consistent Heuristics

\hookrightarrow Graph

一致者

A heuristic is **consistent** if for every node n , every successor n' of n generated by any action a , we have

$$h(n) \leq c(n, a, n') + h(n')$$



If h is consistent, we have

$$\begin{aligned} w(n') &= d(n') + h(n') && \text{(by def.)} \\ &= d(n) + c(n, a, n') + h(n') && (d(n') = d(n) + c(n, a, n')) \\ &\geq d(n) + h(n) = w(n) && \text{(consistency)} \end{aligned}$$

$$w(n') \geq w(n)$$

i.e., $w(n)$ is non-decreasing along any path.

It's the triangle inequality !

keeps all checked nodes in memory to avoid repeated states

tree-search

Theorem:

If $h(n)$ is consistent, A* using GRAPH-SEARCH is optimal

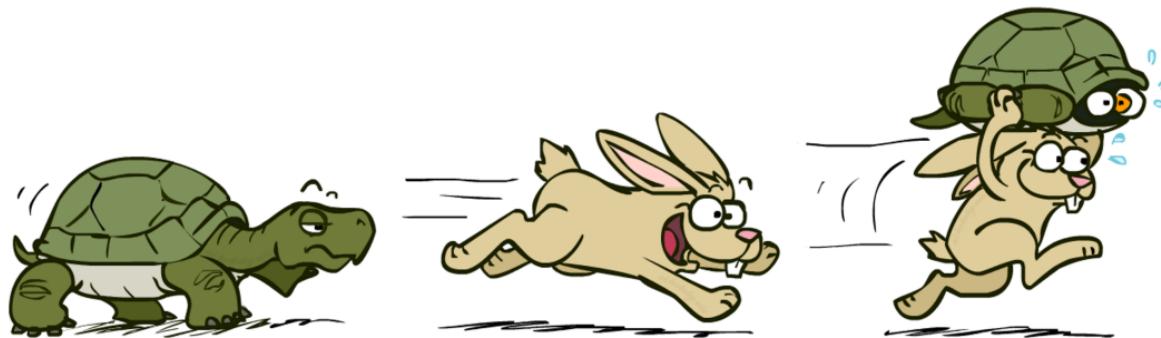
Optimality of A* Search

- With a admissible heuristic, Tree A* is optimal.
- With a consistent heuristic, Graph A* is optimal. *tree is also.*
- See slides, also video lecture from past years for details.
- With $h=0$, the same proof shows that UCS is optimal.

UCS : $h=0$ A*
(all)

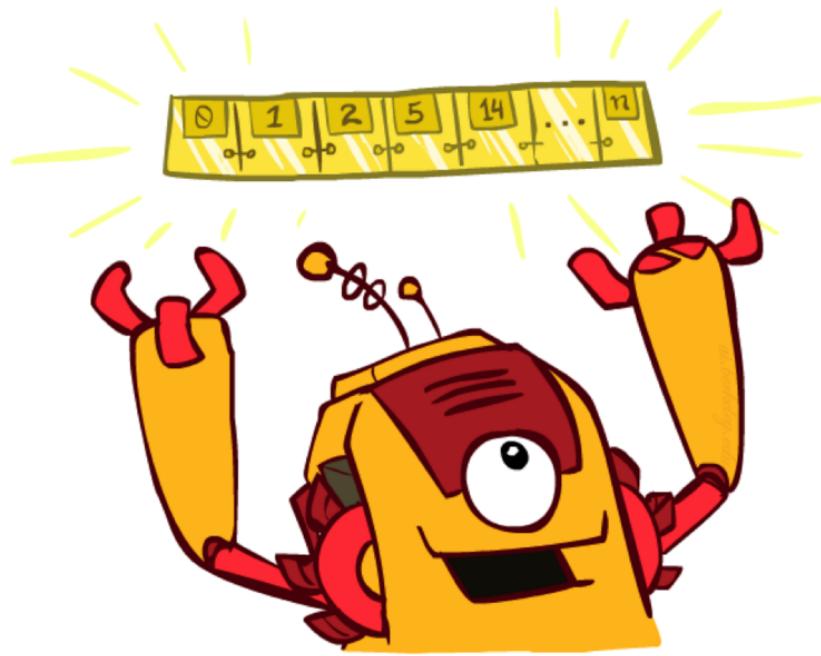
A*: Summary

- A* uses both backward costs and (estimates of) forward costs
count
$$g + h$$
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



The One Queue

- All these search algorithms are the same except for fringe strategies
 - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
 - Practically, for DFS and BFS, you can avoid the $\log(n)$ overhead from an actual priority queue, by using stacks and queues
 - Can even code one implementation that takes a variable queuing object



Search and Models

- Search operates over models of the world
 - The agent doesn't actually try all the plans out in the real world!
 - Planning is all "in simulation"
 - Your search is only as good as your models...

