

Announcement

- TA office hour
 - Wed evening: 20:00-21:00
 - Location: SIST 1B-101
- Programming Assignment 1:
 - Released on Thursday
- Autolab registration
 - See email or BB announcement

Constraint Satisfaction Problems



AIMA Chapter 6

What is Search For?

- Assumptions about the world: a single agent, deterministic actions, fully observed state, discrete state space

Type

Type A Planning: (optimal) sequences of actions

- The path to the goal is the important thing
- Paths have various costs, depths
- Heuristics give problem-specific guidance



Type B Identification ; assignments to variable

- The goal itself is important, not the path
- All paths may be at the same depth (for some formulations)

goal



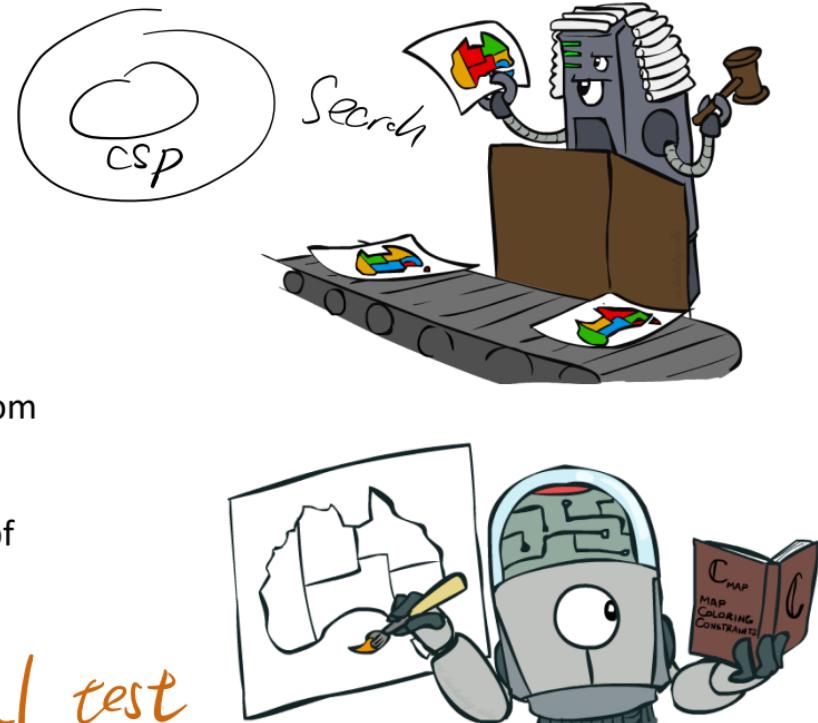
Search Problems

- Standard search problems:
 - State is a “black box”: arbitrary data structure
 - Goal test can be any function over states
 - Successor function can also be anything

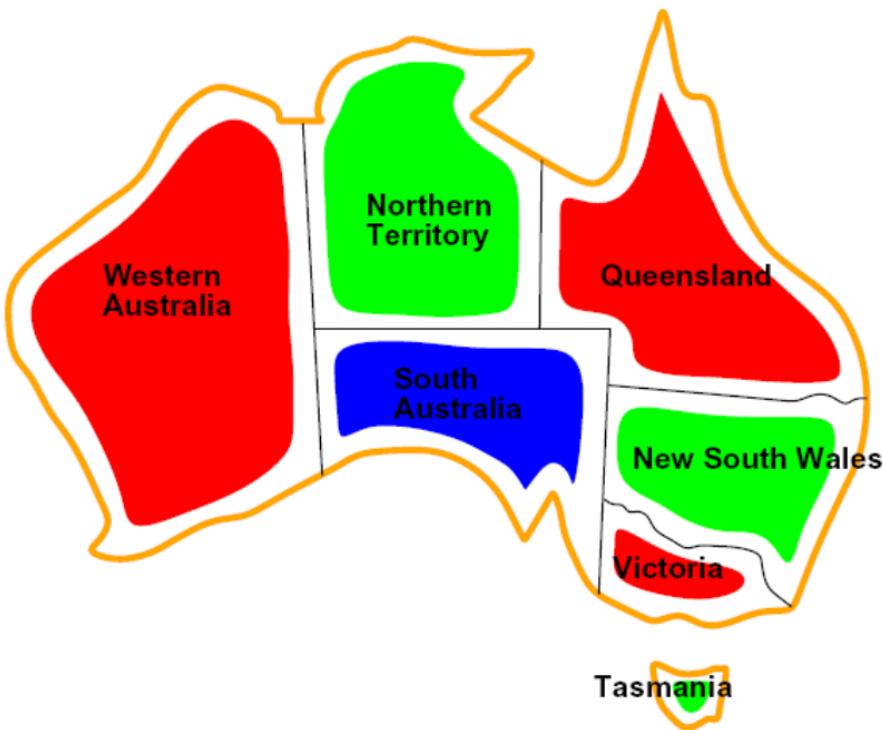
- Constraint satisfaction problems (CSPs):
 - A special subset of search problems
 - State is defined by variables X_i with values from a domain D (sometimes D depends on i)
 - Goal test is a set of constraints specifying allowable combinations of values for subsets of variables

- CSPs are specialized for identification problems

X_i state domain D relation goal test



CSP Examples



Example: Map Coloring

state
↓

- Variables: WA, NT, Q, NSW, V, SA, T x_i

- Domains: $D = \{\text{red, green, blue}\}$

"~~try~~"
~~colors~~

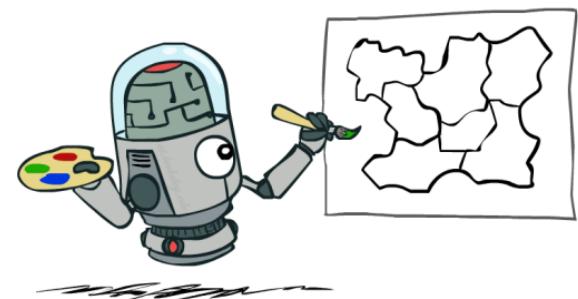
- Constraints: adjacent regions must have different colors

Implicit: $WA \neq NT$

Explicit: $(WA, NT) \in \{(red, green), (red, blue), \dots\}$

- Solutions are assignments satisfying all constraints, e.g.

{WA=red, NT=green, Q=red, NSW=green, V=red, SA=blue, T=green}

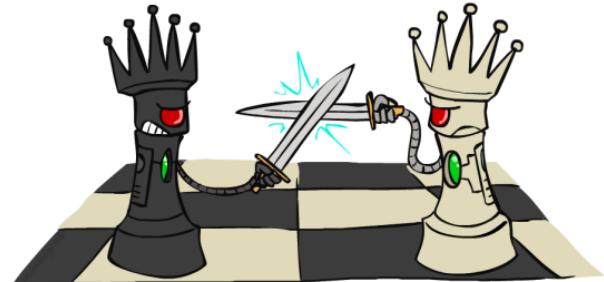
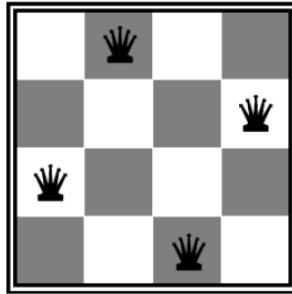


Example: N-Queens

- Formulation 1:

- Variables: X_{ij}
- Domains: $\{0, 1\}$
- Constraints

棋盤上皇后



$$\forall i, j, k \quad (X_{ij}, X_{ik}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{kj}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j+k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j-k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\sum_{i,j} X_{ij} = N$$

Example: N-Queens

- Formulation 2: *assume queen ¹ in every Line*

- Variables: Q_k

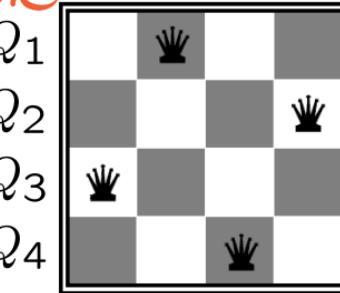
- Domains: $\{1, 2, 3, \dots, N\}$

- Constraints:

Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$

...



Constraint Graphs

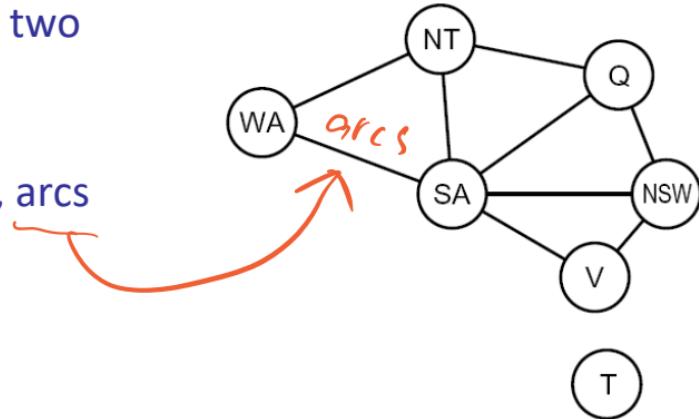
- Binary CSP: each constraint relates (at most) two variables

- Binary constraint graph: nodes are variables, arcs show constraints

Q why c. color?

- Now we can develop general-purpose CSP algorithms on the constraint graph
- What if there are constraints relating more than two variables?

Use graph to speed up search and find cases like isolated nodes.



Example: Cryptarithmetic

- Variables:

$F \ T \ U \ W \ R \ O \ X_1 \ X_2 \ X_3$

Carries

- Domains:

$\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

- Constraints:

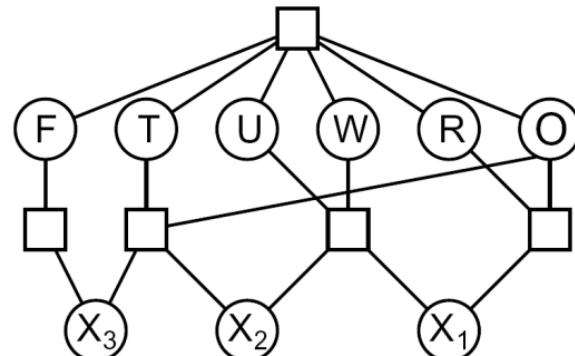
$\text{alldiff}(F, T, U, W, R, O)$

$$O + O = R + 10 \cdot X_1$$

...

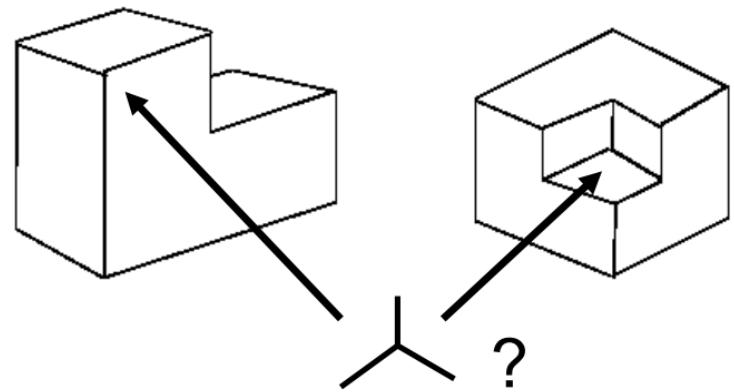
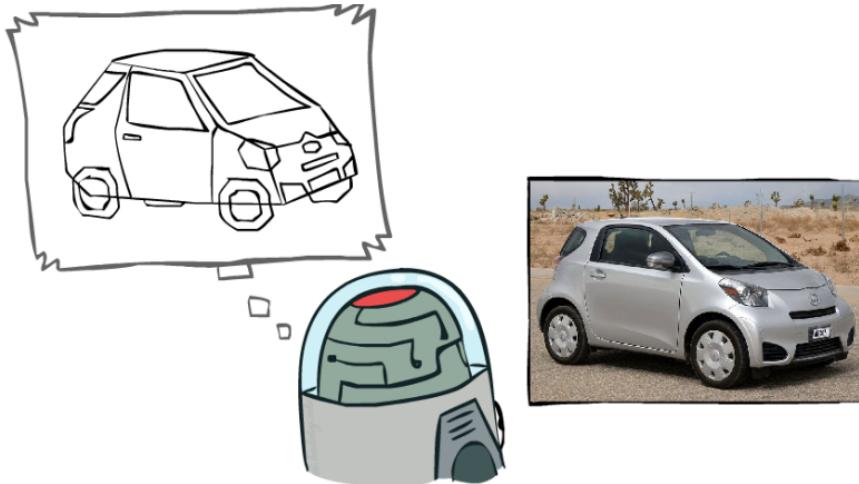
U

$$\begin{array}{r} \text{T} \ \text{W} \ \text{O} \\ + \ \text{T} \ \text{W} \ \text{O} \\ \hline \text{F} \ \text{O} \ \text{U} \ \text{R} \end{array}$$



Example: The Waltz Algorithm

- The Waltz algorithm is for interpreting line drawings of solid polyhedra as 3D objects
- An early example of an AI computation posed as a CSP



- Approach:
 - Each intersection is a variable
 - Adjacent intersections impose constraints on each other
 - Solutions are physically realizable 3D interpretations

Varieties of CSPs and Constraints

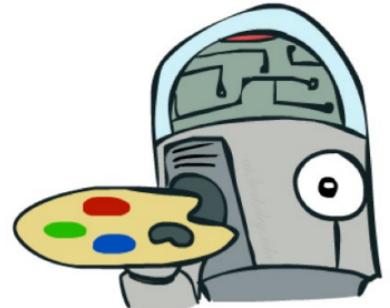


Varieties of CSPs

■ Discrete Variables

- Finite domains $n \uparrow \text{sized}$
 - Size d means $O(d^n)$ complete assignments
 - E.g., Boolean CSPs, including Boolean satisfiability (NP-complete) $d=2$
- Infinite domains (integers, strings, etc.)
 - E.g., job scheduling variables are start/end dates for each job
 - Linear constraints solvable, nonlinear undecidable

Size $\Leftrightarrow \text{domain}_n$



■ Continuous variables

- E.g., start/end times for Hubble Telescope observations
- Linear constraints solvable in polynomial time by LP methods

↓
{ Ap (1/2)}



Varieties of Constraints

- Varieties of Constraints

- Unary constraints involve a single variable (equivalent to reducing domains), e.g.:

$$SA \neq \text{green}$$

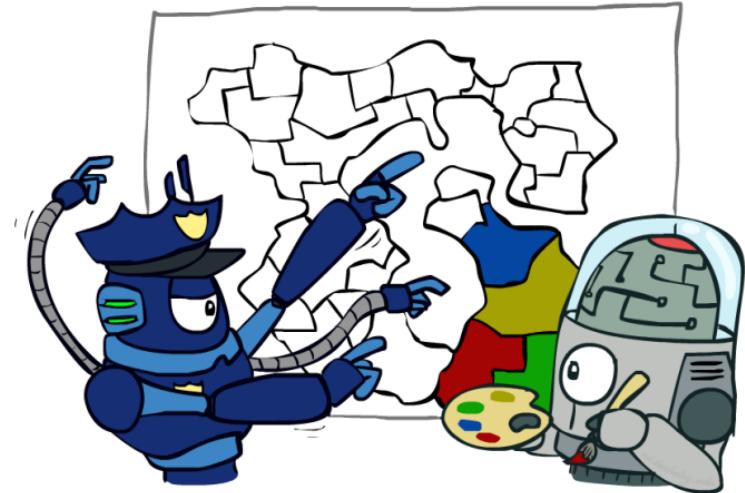
- Binary constraints involve pairs of variables, e.g.:

$$SA \neq WA$$

- Higher-order constraints involve 3 or more variables:

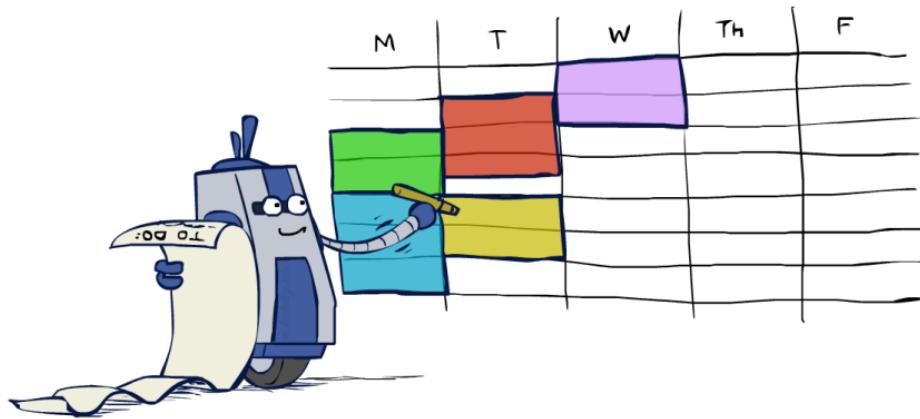
- Preferences (soft constraints):

- E.g., red is better than green
- Often representable by a cost for each variable assignment
- Gives constrained optimization problems
- (We'll ignore these until we get to Bayes' nets)



Real-World CSPs

- Assignment problems: e.g., who teaches what class
- Timetabling problems: e.g., which class is offered when and where?
- Hardware configuration
- Transportation scheduling
- Factory scheduling
- Circuit layout
- Fault diagnosis
- ... lots more!



- Many real-world problems involve real-valued variables...

Solving CSPs



Standard Search Formulation

- Standard search formulation of CSPs

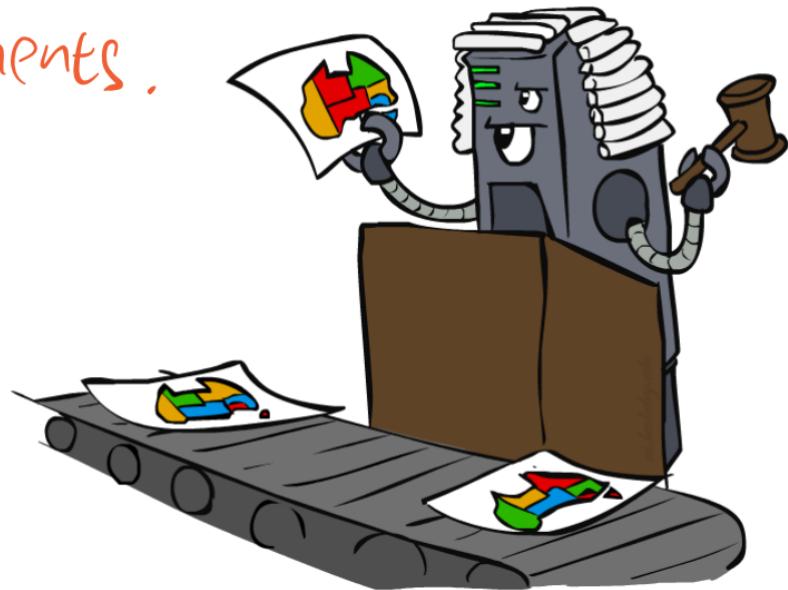
Partial Assignments.

- States defined by the values assigned so far (partial assignments)

- Initial state: the empty assignment, {}
- Successor function: assign a value to an unassigned variable
- Goal test: the current assignment is complete and satisfies all constraints

① {}

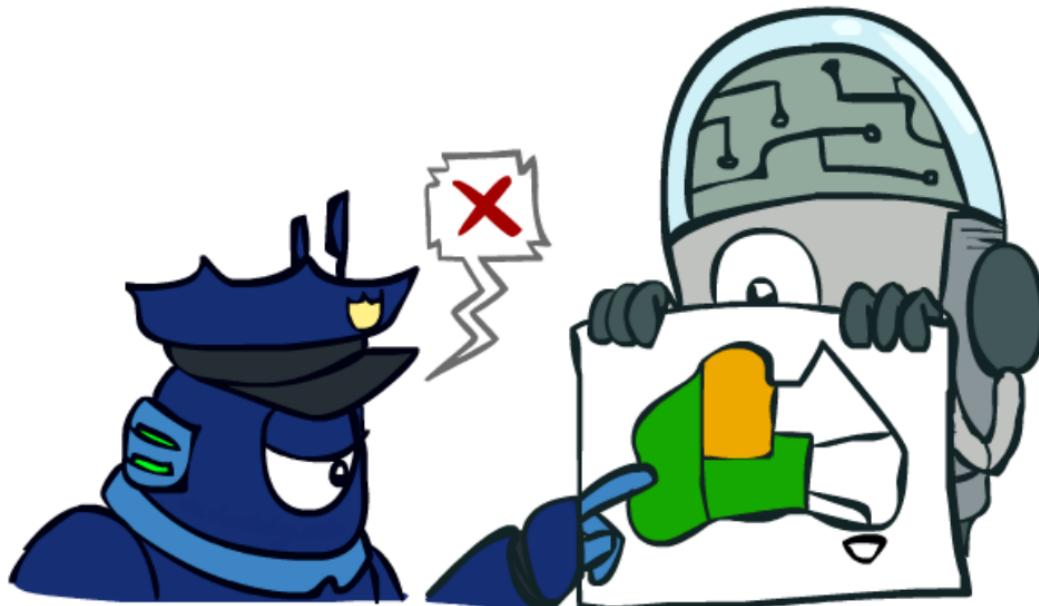
② {A=a ...}



Search Methods

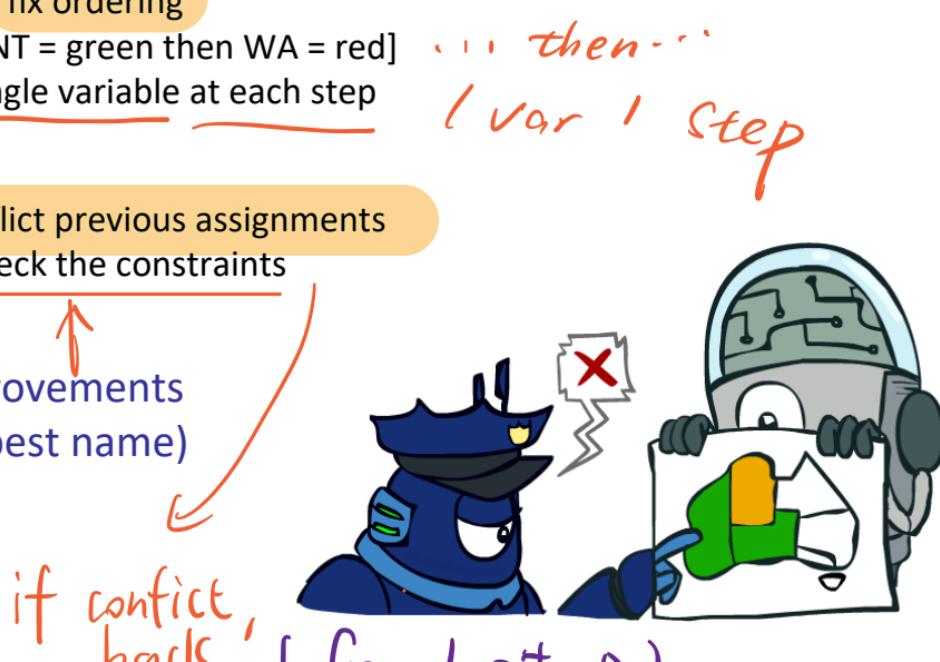
- What would DFS/BFS do?
 - Demo
 - What's wrong?

Backtracking Search

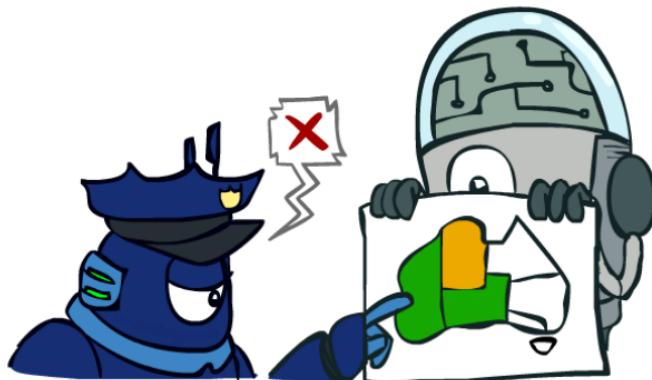
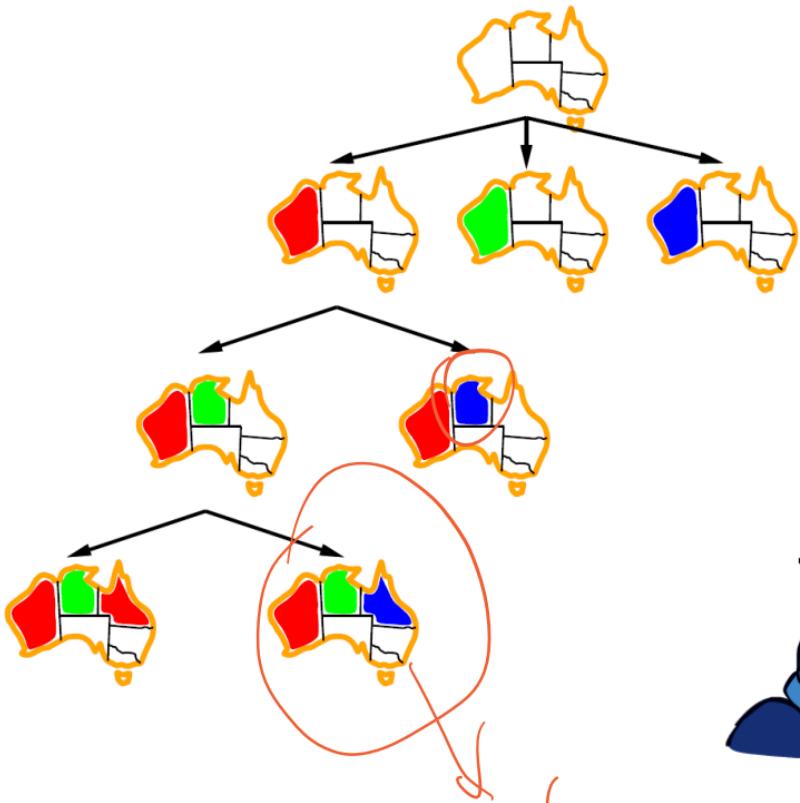


Backtracking Search

- Backtracking search is the basic uninformed algorithm for solving CSPs
- Idea 1: One variable at a time
 - Variable assignments are commutative, so fix ordering
 - I.e., [WA = red then NT = green] same as [NT = green then WA = red]
 - Only need to consider assignments to a single variable at each step
- Idea 2: Check constraints as you go
 - I.e. consider only values which do not conflict previous assignments
 - Might have to do some computation to check the constraints
 - “Incremental goal test”
逐步(逐步)的
- Depth-first search with these two improvements
if conflict, back is called *backtracking search* (not the best name)
- Can solve n-queens for $n \approx 25$



Backtracking Example



Backtracking Search

assignment

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
    return RECURSIVE-BACKTRACKING({ }, csp)
function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
    if assignment is complete then return assignment
    var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment given CONSTRAINTS[csp] then
            add {var = value} to assignment
            result ← RECURSIVE-BACKTRACKING(assignment, csp)
            if result ≠ failure then return result
            remove {var = value} from assignment
    return failure ↴
```

Sols.

↙↙

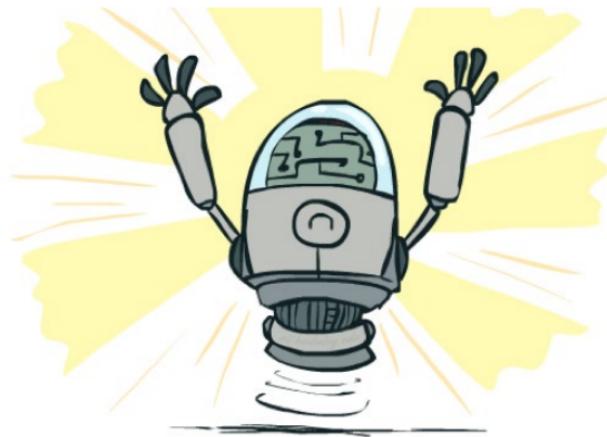
DFS!

- Backtracking = DFS + variable-ordering + fail-on-violation
- What are the choice points?

objects value, object (MRV, LCV ...)

Improving Backtracking

- General-purpose ideas give huge gains in speed
- Filtering: Can we detect inevitable failure early?
- Ordering:
 - Which variable should be assigned next?
 - In what order should its values be tried?
- Structure: Can we exploit the problem structure?



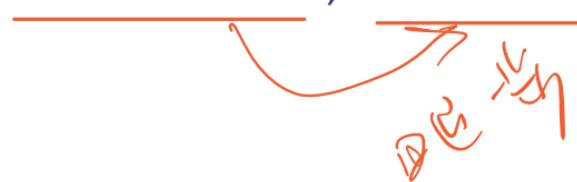
Filtering



filter out the nodes not considered

Filtering: Forward Checking

- Filtering: Keep track of domains for unassigned variables and cross off bad options
- Forward checking: Cross off values that violate a constraint when added to the existing assignment; whenever any variable has no value left, we backtrack



DBS
how to?

WA	NT	Q	NSW	V	SA
█ Red	█ Green	█ Blue	█ Red	█ Green	█ Blue

Forward: assigned value.

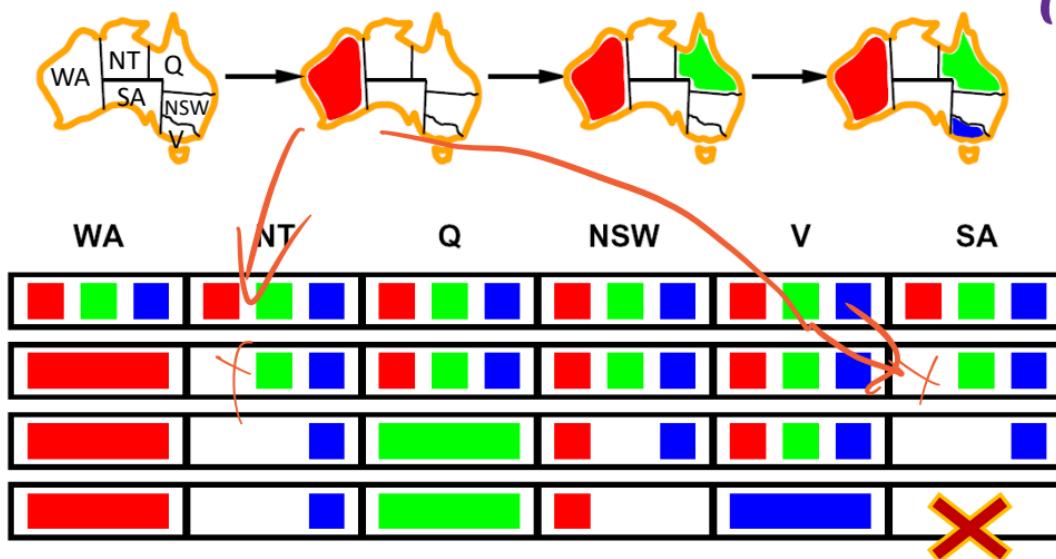
value
✓

look into if contradiction in

Filtering: Forward Checking

rest
(neighbours)

- Filtering: Keep track of domains for unassigned variables and cross off bad options
- Forward checking: Cross off values that violate a constraint when added to the existing assignment; whenever any variable has no value left, we backtrack



"T - t"

demo

Demo

- Backtracking
- Backtracking with Forward Checking

only immediate violation

consider between
unassigned

unassigned



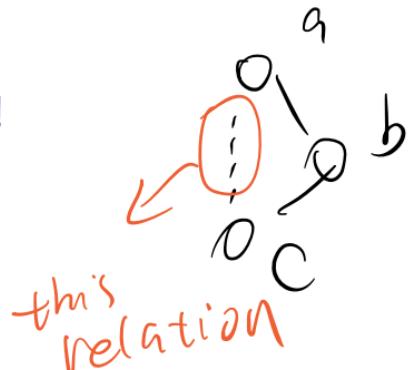
continues to

Filtering: Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:



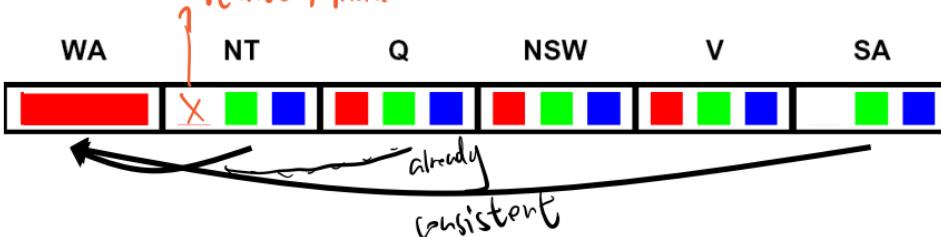
WA	NT	Q	NSW	V	SA
Red	Green	Blue	Red	Green	Blue
Red	White	Green	Red	Green	Blue



- NT and SA cannot both be blue!
- Can we detect this early?

Consistency of A Single Arc

- An arc $X \rightarrow Y$ is **consistent** iff for every x in the tail there is some y in the head which could be assigned without violating a constraint



Forward checking?

Enforcing consistency of arcs pointing to each new assignment

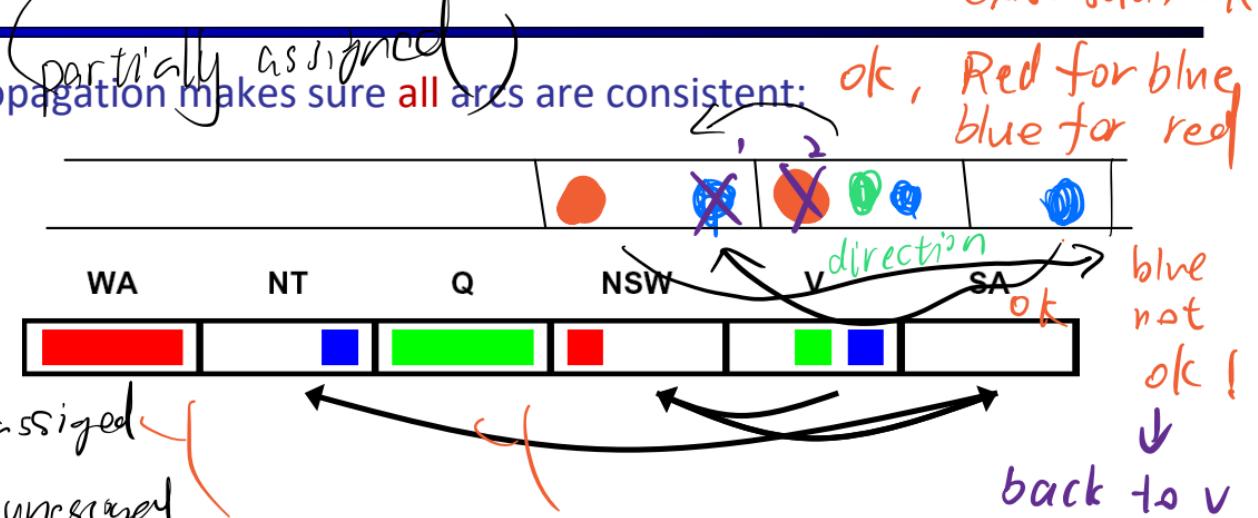
head tail?

Arc Consistency of an Entire CSP

no matter config

exist solu, ok

- A simple form of propagation makes sure all arcs are consistent:



- Important: If Y loses a value, then arc X → Y needs to be rechecked!
- Arc consistency detects failure earlier than forward checking
- What's the downside of enforcing arc consistency?

Remember: Delete from the tail!

failure \Rightarrow backtracks

tail \rightarrow head
arc is:
tail \rightarrow head
break off, head \leftarrow tail

Wa	N7	Q	NSW	V	S/D	
R	b	G	r x	X ^{gb}	b	

3. NSW should be checked!

judge
(forwarded)

1. delete from tail
2. tail → head

3. head changed

recheck, head ←
all

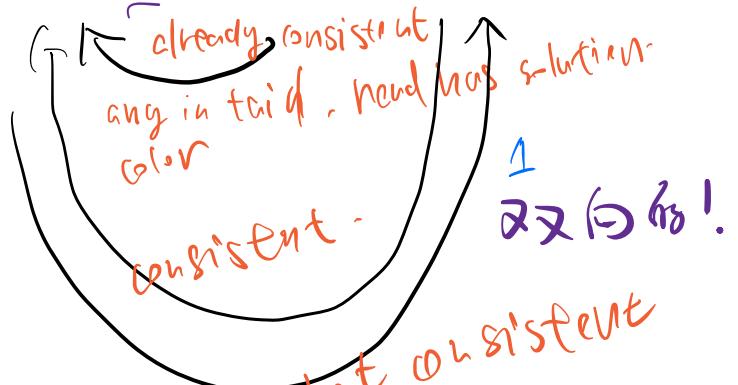
4. for any color in tail

exist sth in head

consist

5. when fail back track

to last time we choose
a value



Enforcing Arc Consistency in a CSP

n^2
back cause d
 d^2

function AC-3(*csp*) returns the CSP, possibly with reduced domains
inputs: *csp*, a binary CSP with variables $\{X_1, X_2, \dots, X_n\}$
local variables: *queue*, a queue of arcs, initially all the arcs in *csp*
while *queue* is not empty do *queue*: inconsistent arc
 $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$ *queue*.
 if REMOVE-INCONSISTENT-VALUES(X_i, X_j) then
 for each X_k in NEIGHBORS[X_i] do
 add (X_k, X_i) to *queue* through neighbour back

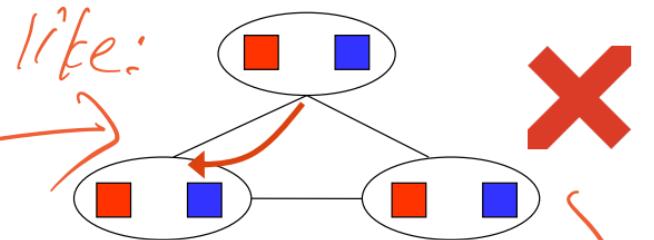
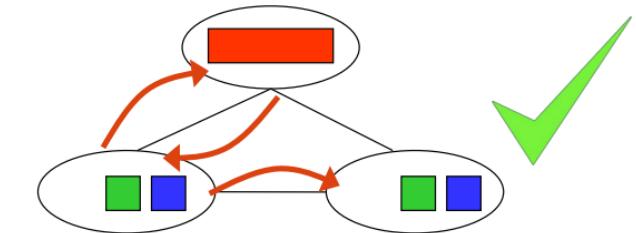
function REMOVE-INCONSISTENT-VALUES(X_i, X_j) returns true iff succeeds
 $\text{removed} \leftarrow \text{false}$
for each x in DOMAIN[X_i] do tail
 if no value y in DOMAIN[X_j] allows (x, y) to satisfy the constraint $X_i \leftrightarrow X_j$
 then delete x from DOMAIN[X_i]; $\text{removed} \leftarrow \text{true}$
return removed

- Runtime: $O(n^2d^3)$, can be reduced to $O(n^2d^2)$

Domain: y from
 X_j

Limitations of Arc Consistency

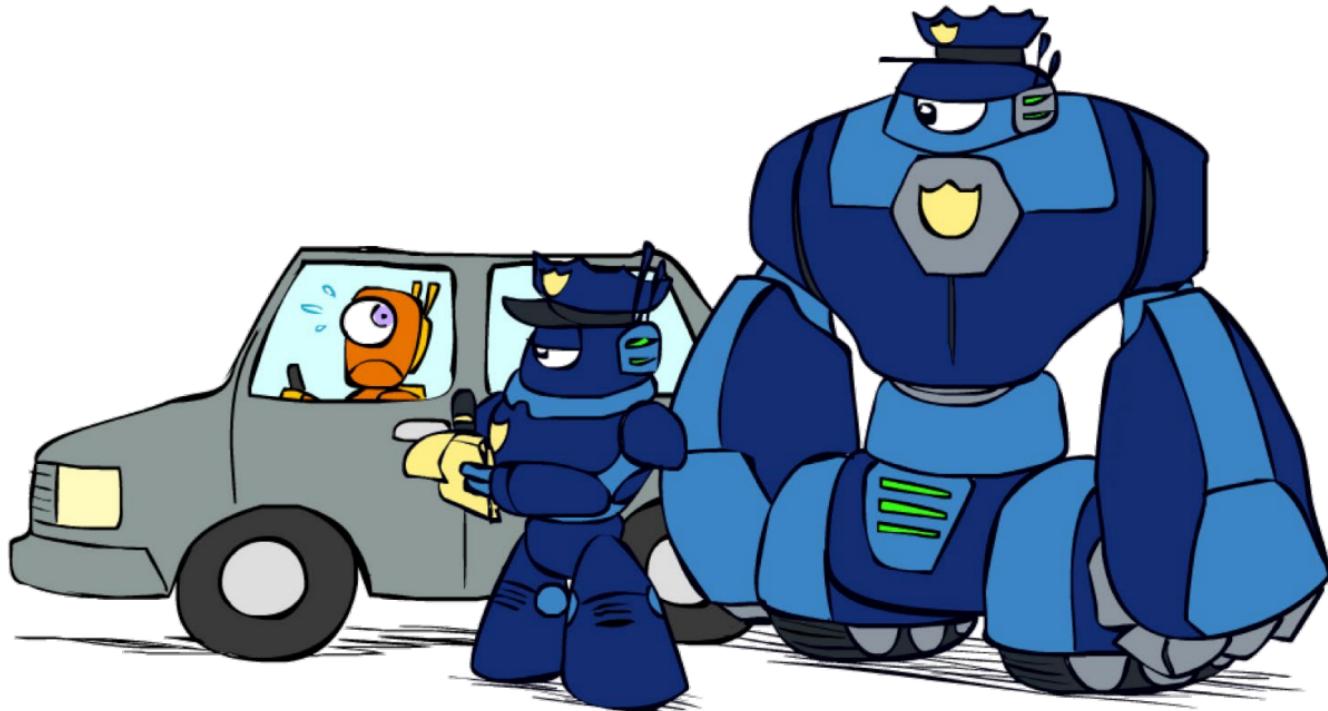
- After enforcing arc consistency:
 - Can have one solution left
 - Can have multiple solutions left
 - Can have no solutions left (and not know it)
- Arc consistency still runs inside a backtracking search!



Demo

- Backtracking with Forward Checking
- Backtracking with Arc Consistency

K-Consistency



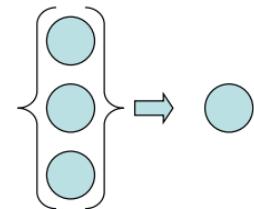
K-Consistency

- Increasing degrees of consistency

- 1-Consistency (Node Consistency): Each single node's domain has a value which meets that node's unary constraints
- 2-Consistency (Arc Consistency): For each pair of nodes, any consistent assignment to one can be extended to the other
- K-Consistency: For each k nodes, any consistent assignment to k-1 can be extended to the kth node.
- Higher k more expensive to compute



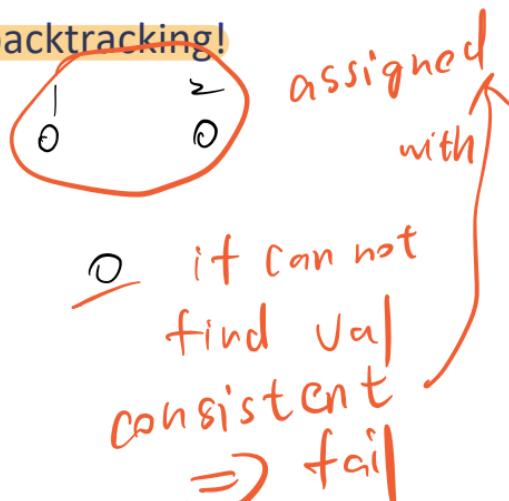
0
↑



Some.

Strong K-Consistency

- Strong k-consistency: also k-1, k-2, ... 1 consistent
- Claim: strong n-consistency means we can solve without backtracking!
- Why? **NP Hard**
 - Choose any assignment to any variable
 - Choose a new variable
 - By 2-consistency, there is a choice consistent with the first
 - Choose a new variable
 - By 3-consistency, there is a choice consistent with the first 2
 - ...
- Lots of middle ground between arc consistency and n-consistency! (e.g. k=3, called path consistency)



Ordering



MRV

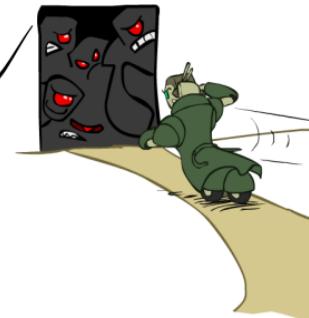
Ordering: Minimum Remaining Values

- Variable Ordering: Minimum remaining values (MRV):
 - Choose the variable with the fewest legal left values in its domain
 - Also called “most constrained variable”

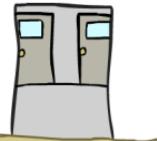


3 color
2 color
1 color ✓

“Fail Fast” easy to correct



choose variable

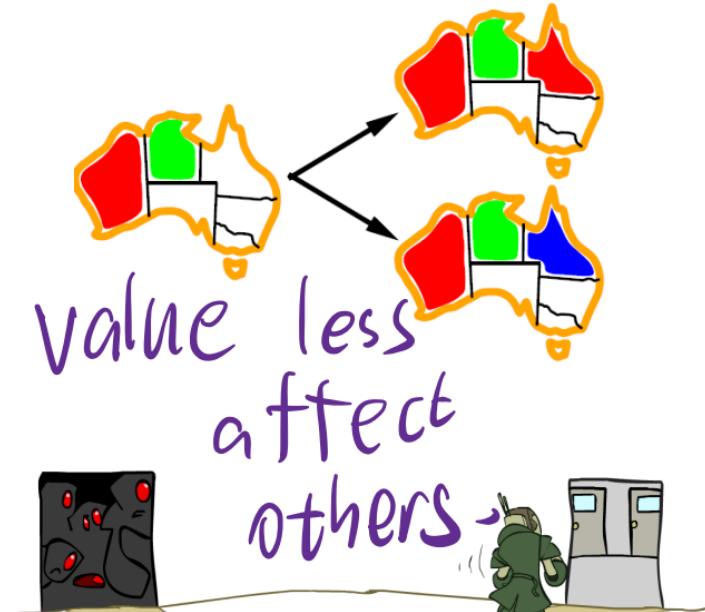


LCV

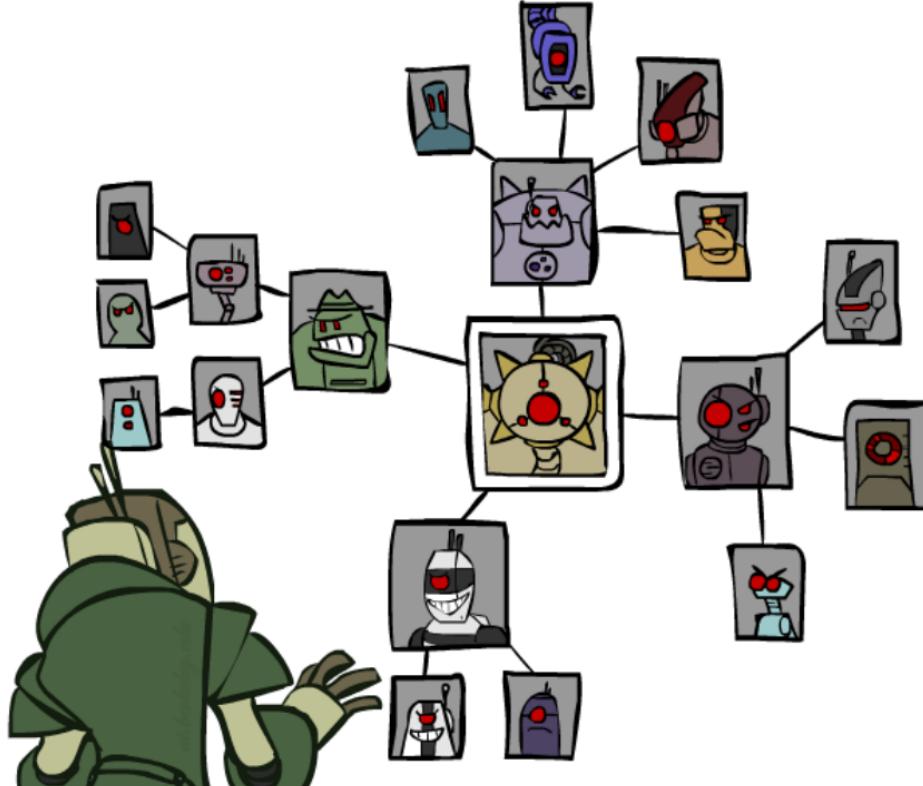
Ordering: Least Constraining Value

- Value Ordering: Least Constraining Value
 - Given a choice of variable, choose the *least constraining value* 排除
 - I.e., the one that rules out the fewest values in the remaining variables
 - Note that it may take some computation to determine this! (E.g., rerunning filtering)
- Combining these ordering ideas makes 1000 queens feasible

Given Variable

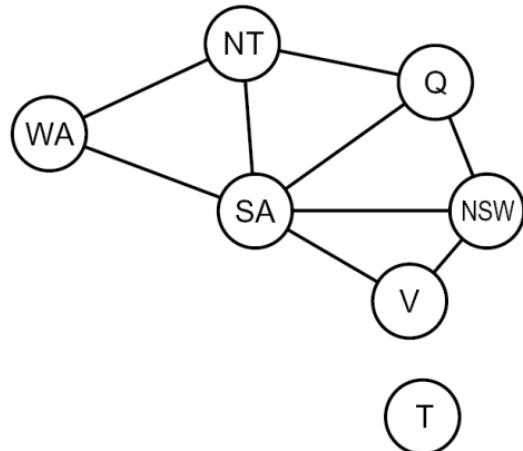


choose Value Structure

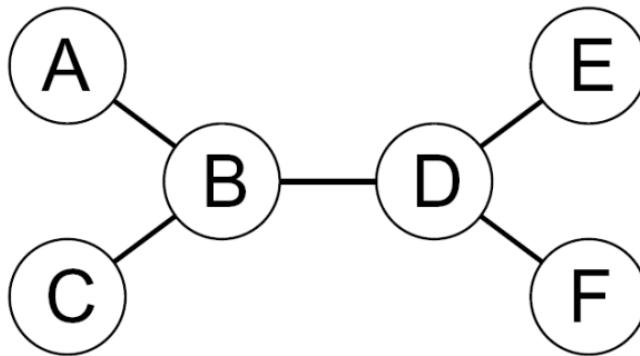


Problem Structure

- Extreme case: independent subproblems
 - Example: Tasmania and mainland do not interact
可并行
- Independent subproblems are identifiable as connected components of constraint graph
- Suppose a graph of n variables can be broken into subproblems of only c variables:
 - Worst-case solution cost is $O((n/c)(d^c))$, linear in n
 - E.g., $n = 80$, $d = 2$, $c = 20$
 - $2^{80} = 4$ billion years at 10 million nodes/sec
 - $(4)(2^{20}) = 0.4$ seconds at 10 million nodes/sec



Tree-Structured CSPs



(linear)

Time -

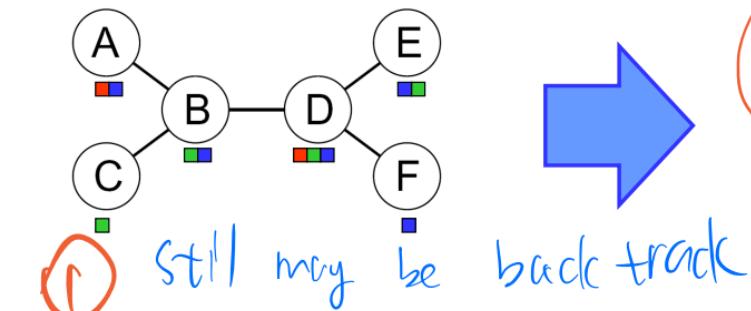
- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(n d^2)$ time
 - Compare to general CSPs, where worst-case time is $O(d^n)$
- This property also applies to probabilistic reasoning (later)
- An example of the relation between syntactic restrictions and the complexity of reasoning

Make it a tree!

Tree-Structured CSPs

- Algorithm for tree-structured CSPs:

- Order: Choose a root variable, order variables so that parents precede children

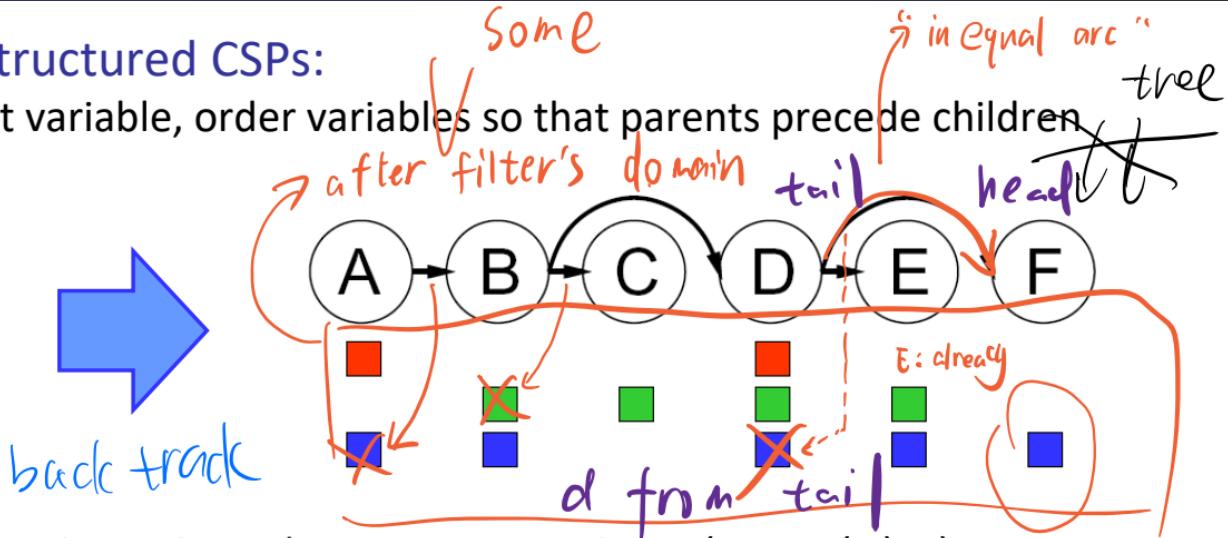


① Still may be back track

- ② Remove backward: For $i = n : 2$, apply RemoveInconsistent($\text{Parent}(X_i), X_i$)
③ Assign forward: For $i = 1 : n$, assign X_i consistently with $\text{Parent}(X_i)$

- Runtime: $O(n d^2)$

when do this
when consistent



assign values

A	B	C	D	E	F
R	B	G	G	B	B

for each node
make arc pointing
to it consistent

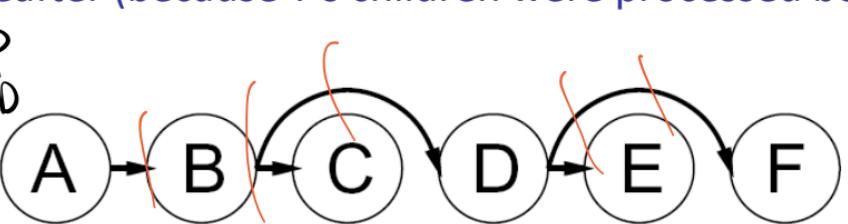
with
Just choose
things work
out

R Tree-Structured CSPs

delete from tail
to make the rest all

- Claim 1: After backward pass, all root-to-leaf arcs are consistent
- Proof: Each $X \rightarrow Y$ was made consistent at one point and Y's domain could not have been reduced thereafter (because Y's children were processed before Y)

Exemple : $A \rightarrow B$
 $A \rightarrow B$
r r ok! ✗
b b not ok!



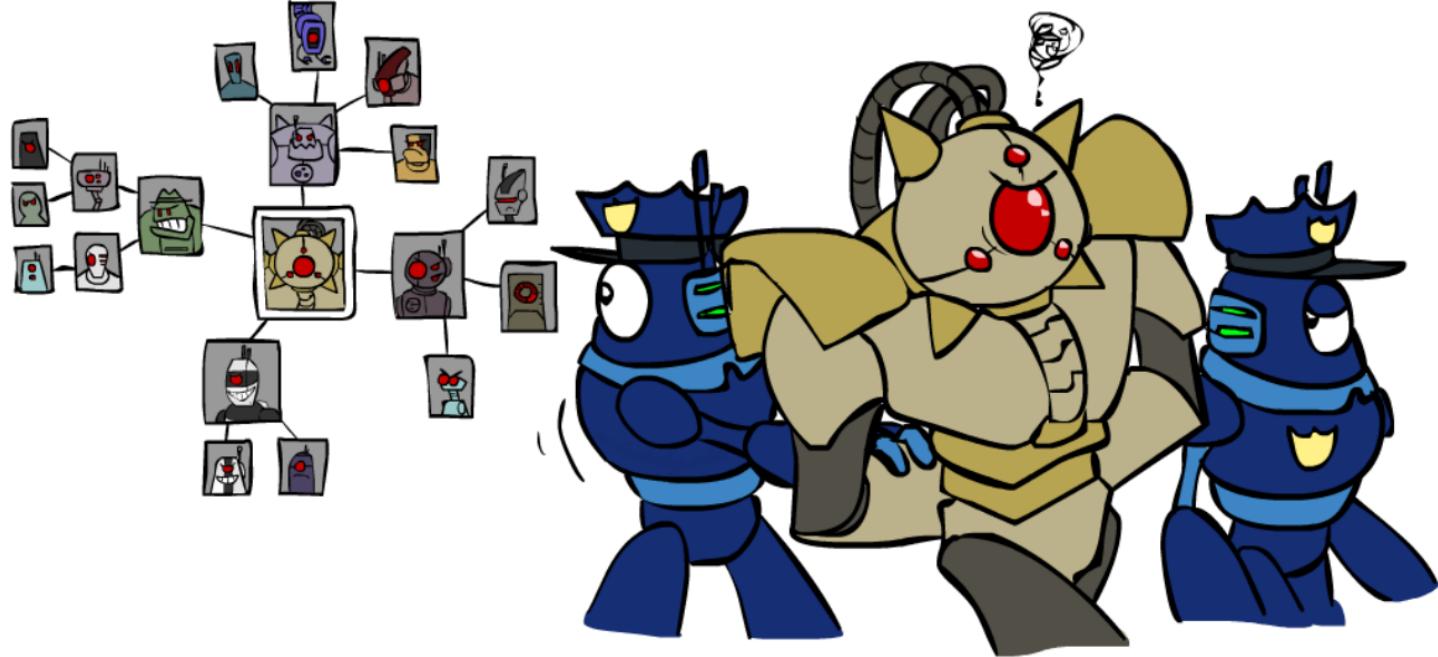
has some in head
consist with it



- Claim 2: If root-to-leaf arcs are consistent, forward assignment will not backtrack
- Easy to prove
- Why doesn't this algorithm work with cycles in the constraint graph?
- Note: we'll see this basic idea again with Bayes' nets

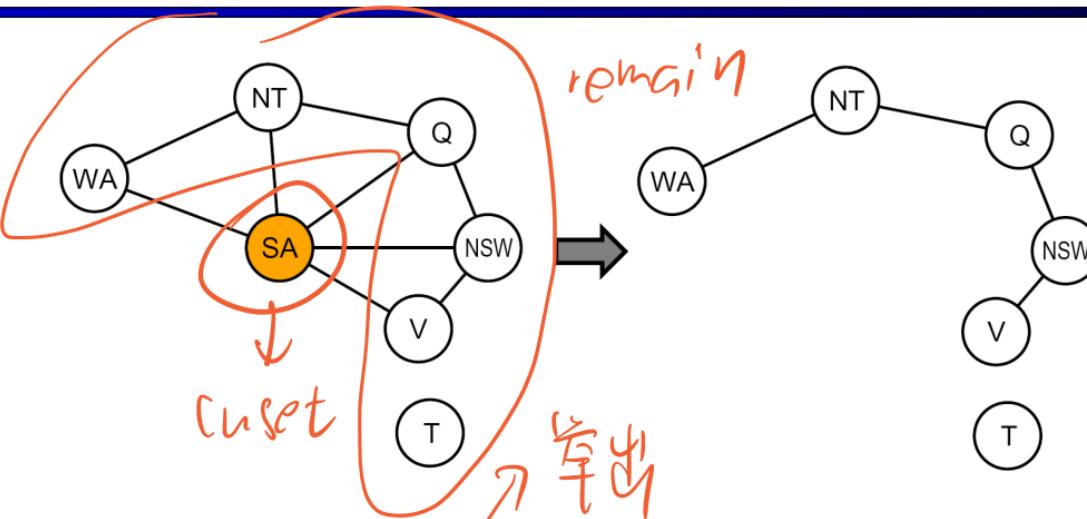
= is legal

Cutset Conditioning



To make it a tree

Nearly Tree-Structured CSPs



- Cutset: a set of variables s.t. the remaining constraint graph is a tree
- Cutset conditioning: instantiate (in all ways) the cutset and solve the remaining tree-structured CSP
 - Cutset size c gives runtime $O((d^c) (n-c) d^2)$, very fast for small c

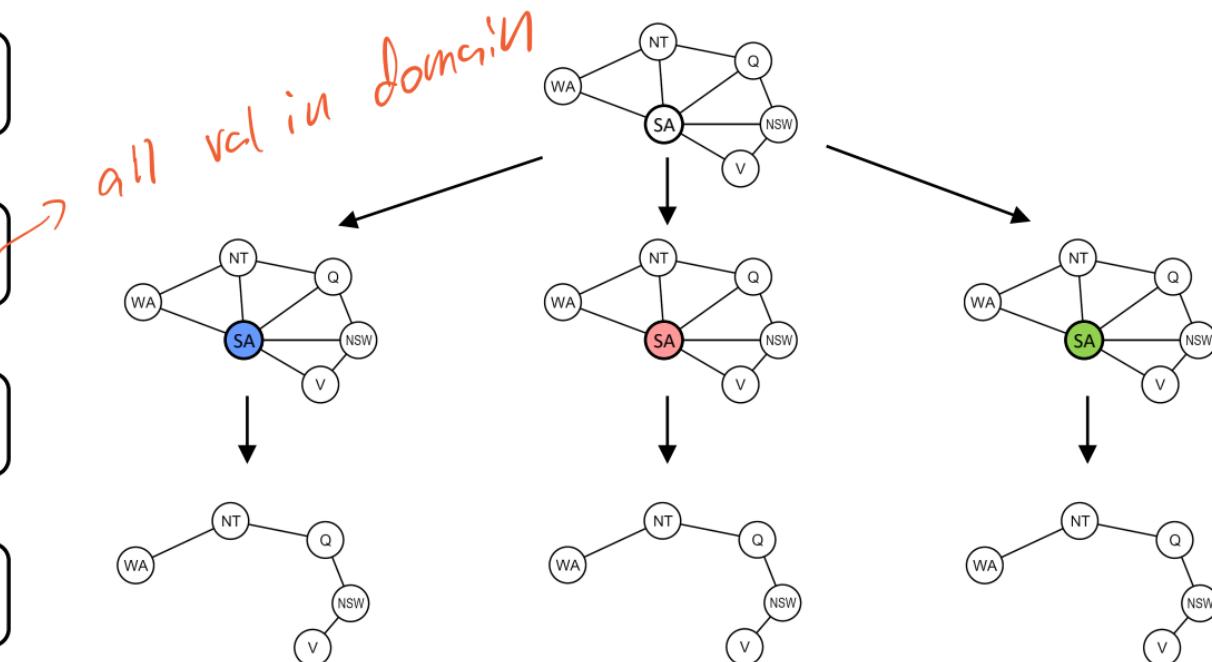
Cutset Conditioning

Choose a cutset

Instantiate the cutset
(all possible ways)

Compute residual CSP
for each assignment

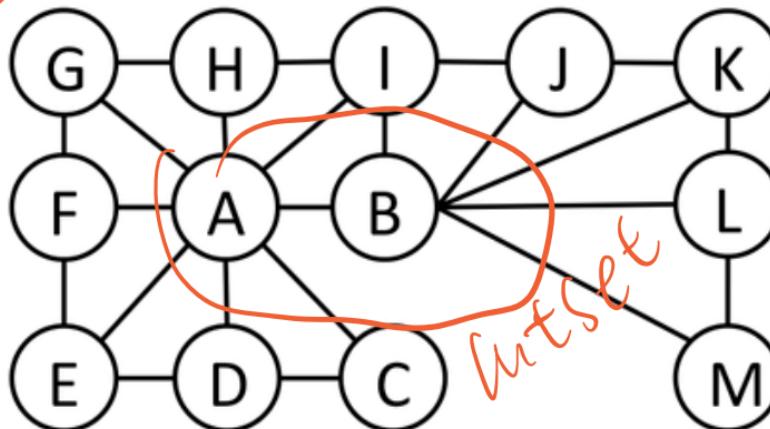
Solve the residual CSPs
(tree structured)



less dc

Finding Cutset

- Find the **smallest** cutset for the graph below.



- Finding the *smallest* cutset is **NP-hard**
- But there are efficient approximation algorithms