


Announcement


- Programming Assignment 1B: adversarial search
 - Instructions at Blackboard -> “Programming Assignments”
 - Submission at AutoLab
 - Due: Mar 10, 11:59pm

Adversarial Search



Adversarial Search (Minimax)

- **Deterministic, zero-sum games:**
 - Tic-tac-toe, chess, checkers
 - Players alternate turns
 - One player maximizes result
 - The other minimizes result
- **Minimax search:**
 - A state-space search tree
 - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



Minimax Implementation

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

def max-value(state):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

def min-value(state):

initialize $v = +\infty$


for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$

return v


$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example




Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?



Resource Limits



Resource Limits

- Problem: In realistic games, cannot search to leaves!

- Solution: Depth-limited search


- Instead, search only to a limited depth in the tree
- Replace terminal utilities with an evaluation function for non-terminal positions

- Example:

- Suppose we have 100 seconds, can explore 10K nodes / sec
- So can check 1M nodes per move
- α - β reaches about depth 8 – decent chess program


- Guarantee of optimal play is gone

- More depth makes a BIG difference



Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position
- A simple solution in practice: weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$


- e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.

Evaluation Functions

- Recent advances
 - Monte Carlo Tree Search
 - Randomly choose moves until the end of game
 - Repeat for many many times
 - Evaluate the state based on these simulations, e.g., the winning rate
 - Convolutional Neural Network (value network in AlphaGo)
 - Trained from records of game plays to predict a score of the state

Branching Factor

- Chess



Branching Factor


- Go




Branching Factor

- Go has a branching factor of up to 361
- Idea: limit the branching factor by considering only good moves
 - AlphaGo uses a Convolutional Neural Network (policy network)
 - Trained from records of game plays
 - Trained using reinforcement learning
 - AlphaGo Zero uses RL only

Game Tree Pruning



Minimax Pruning




Alpha-Beta Pruning

- General configuration (MIN version)

- We're computing the MIN-VALUE at some node n
- We're looping over n 's children, so n 's estimate is decreasing
- Let a be the best value that MAX can get at any choice point along the current path from the root
- If n becomes worse than a , then we can stop considering n 's other children
- Reason: if n is eventually chosen, then the nodes along the path shall all have the value of n , but n is worse than a and hence the path shall not be chosen at the MAX

- MAX version is symmetric




Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root


```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```


Alpha-Beta Example




Alpha-Beta Example 2




Alpha-Beta Example 2



Alpha-Beta Example 2




Alpha-Beta Example 2




Alpha-Beta Pruning Properties

- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!





Other Game Types






Mixed Layer Types

- Backgammon
- Expectiminimax
 - Environment is an extra “random agent” player that moves after each min/max agent
 - Each node computes the appropriate combination of its children



Multi-Agent Utilities

- What if the game is not zero-sum, or has multiple players?
- Generalization of minimax:
 - Terminals have **utility tuples**
 - Node values are also utility tuples
 - Each player maximizes its own component
 - Can give rise to cooperation and competition dynamically...



Summary

- Adversarial Games
- Adversarial Search
 - Minimax
- Resource Limits
 - Depth-limited search, limiting branching factor
- Game Tree Pruning (alpha-beta pruning)
- Uncertain Outcomes
 - Expectimax
- Other Game Types

