


Three types of (strong) AI approaches





Propositional Logic

AIMA Chapter 7

Outline


- Logic
- Propositional logic
 - Syntax
 - Semantics
 - Inference
- Horn logic
 - Inference
- An example application

Logic-based Symbolic AI

- Logic
 - Formal language in which knowledge can be expressed
 - A means of carrying out reasoning in the language



Logic-based Symbolic AI

- Logic (Knowledge-Based) AI
 - Knowledge base
 - set of **sentences** in a formal language to represent knowledge about the “world”
 - Inference engine
 - answers any answerable question following the knowledge base



Formal Language

- Components of a formal language in a logic
 - **Syntax**: What sentences are allowed?
 - **Semantics**:
 - Which sentences are true/false in each **model** (possible world)?



Formal Language

- Example: the language of arithmetic

- Syntax

- $x+2 \geq y$ is a sentence
 - $x^2+y > \{\}$ is not a sentence (not allowed)

cuz
no meaning





- Semantics

- $x+2 \geq y$ is true in a world where $x = 7, y = 1$
 - $x+2 \geq y$ is false in a world where $x = 0, y = 6$

Propositional Logic

Propositional logic: Syntax

- **Propositional logic** is the “simplest” logic
 - The proposition symbols P_1, P_2 , etc. are sentences
 - If S is a sentence, $\neg S$ is a sentence (**negation**)
 - If S_1 and S_2 are sentences, $S_1 \wedge S_2$ is a sentence (**conjunction**) 
 - If S_1 and S_2 are sentences, $S_1 \vee S_2$ is a sentence (**disjunction**) 
 - If S_1 and S_2 are sentences, $S_1 \Rightarrow S_2$ is a sentence (**implication**)
 - If S_1 and S_2 are sentences, $S_1 \Leftrightarrow S_2$ is a sentence (**biconditional**)

$\neg, \wedge, \vee, \Rightarrow, \Leftrightarrow$ are called logic connectives or operators



Sometimes \rightarrow and \leftrightarrow are used

propositional logic

Examples of PL sentences

- P means “It is hot.”
- Q means “It is humid.”
- R means “It is raining.”
- $(P \wedge Q) \Rightarrow R$
 - “If it is hot and humid, then it is raining”
- $Q \Rightarrow P$
 - “If it is humid, then it is hot”

Propositional logic: Semantics

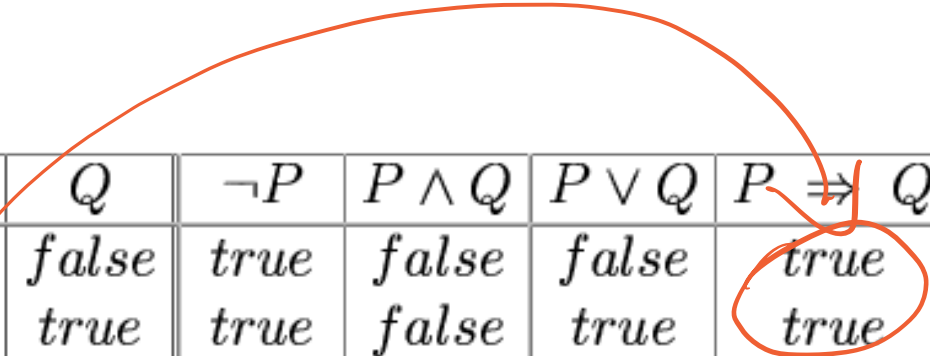
- Each model specifies true/false for each proposition symbol

– E.g. P_1 P_2 P_3
 false true false

- Rules for evaluating truth with respect to a model m :

- $\neg S$ is true iff S is false
- $S_1 \wedge S_2$ is true iff S_1 is true and S_2 is true
- $S_1 \vee S_2$ is true iff S_1 is true or S_2 is true
- $S_1 \Rightarrow S_2$ is true iff S_1 is false or S_2 is true
- $S_1 \Leftrightarrow S_2$ is true iff $S_1 \Rightarrow S_2$ is true and $S_2 \Rightarrow S_1$ is true


Truth tables for connectives




P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
<u>false</u>	false	true	false	false	true	true
<u>false</u>	true	true	false	true	true	false
true	false	false	false	true	false	false
true	true	false	true	true	true	true

Venn Diagrams


$P \vee Q$




$P \wedge Q$



$P \Rightarrow Q$



$P \Leftrightarrow Q$



Material Implication

- $S1 \Rightarrow S2$ is true iff $S1$ is false or $S2$ is true
- Given the following propositions, is “ $S1 \Rightarrow S2$ ” true?
 - $S1$ means “the moon is made of green cheese”
 - $S2$ means “the world is coming to an end”
- Material implication does not capture the meaning of “if... then”.
- See “[Paradoxes of material implication](#)” in Wikipedia

no natural
language

Q

$(\neg B) \wedge (A) \times$

$(\neg \beta) \wedge (\neg \alpha) \times$

Logical equivalence

- Two sentences are **logically equivalent** iff true in same models

$$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha) \quad \text{commutativity of } \wedge$$

$$(\alpha \vee \beta) \equiv (\beta \vee \alpha) \quad \text{commutativity of } \vee$$

$$((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma)) \quad \text{associativity of } \wedge$$

$$((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma)) \quad \text{associativity of } \vee$$

$$\neg(\neg\alpha) \equiv \alpha \quad \text{double-negation elimination}$$

$$(\alpha \Rightarrow \beta) \equiv (\neg\beta \Rightarrow \neg\alpha) \quad \text{contraposition}$$

$$(\alpha \Rightarrow \beta) \equiv (\neg\alpha \vee \beta) \quad \text{implication elimination}$$

$$(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)) \quad \text{biconditional elimination}$$

$$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta) \quad \text{de Morgan}$$

$$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta) \quad \text{de Morgan}$$

$$(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma)) \quad \text{distributivity of } \wedge \text{ over } \vee$$

$$(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma)) \quad \text{distributivity of } \vee \text{ over } \wedge$$

! \Rightarrow

$$(A \vee B \vee C) \vee (A \wedge B \wedge \neg C)$$

$$= (A \vee B \vee C) \vee (A \wedge (B \wedge \neg C))$$

(. . .)

/

Validity and satisfiability



- A sentence is **valid** if it is true in all models
 - e.g., $A \vee \neg A$, $A \Rightarrow A$, $(A \wedge (A \Rightarrow B)) \Rightarrow B$
- A sentence is **satisfiable** if it is true in some model
 - e.g., $A \vee B$, C
- A sentence is **unsatisfiable** if it is true in no models
 - e.g., $A \wedge \neg A$
- Obviously, S is valid iff. $\neg S$ is unsatisfiable

$$\begin{aligned}
 & A \wedge (A \Rightarrow B) \\
 &= A \wedge (\neg A \vee B) \\
 &= (A \wedge \neg A) \vee (A \wedge B) \\
 &= \text{false} \vee (A \wedge B) \\
 &= A \wedge B
 \end{aligned}$$

$A \wedge B$ true
 A true
 B true

Inference: entailment

- **Entailment:** $\alpha \models \beta$ (使 β 为真) (“ α entails β ” or “ β follows from α ”) means in every world where α is true, β is also true
 - i.e., the α -worlds are a subset of the β -worlds [$\text{models}(\alpha) \subseteq \text{models}(\beta)$]
 - In the example, $\alpha_2 \models \alpha_1$ world = model
- why



set 内

同值性: $A \cap B$ true, A true

" $\text{True} \models \text{False}$ " is indeed false: every model makes "true" true, but no model makes "false" true, so *every* model provides a counterexample.

However, since no model makes "false" true, " $\text{False} \models \text{True}$ " is actually *true*! It's an instance of *vacuous implication*: think of it as being true for the same reason the statement

"If $0 = 1$, then I'm the president"

is true.

- (a) α is valid if and only if $\text{True} \models \alpha$.

Forward direction: If $\text{True} \models \alpha$, then α is valid.

By definition, $\text{True} \models \alpha$ means that α is true in all worlds where True is True; this is all worlds. Thus α is true in all worlds, which is exactly the definition of validity. So α is in this case valid.

Backward direction: If α is valid, then $\text{True} \models \alpha$.

By definition, if α is valid then it is **true in all worlds**. In this case *anything* entails α , so clearly True entails α .

- (b) For any α , $\text{False} \models \alpha$.

Recall the definition of entailment: $p \models q$ means that in all worlds in which p is true, q is true as well. So, $\text{False} \models \alpha$ means that in all worlds in which False is true, α is true. But there are no worlds in which False is true! Clearly if there are no worlds in a set, then that set satisfies the condition that α be true in all worlds in that set.

- (c) $\alpha \models \beta$ if and only if the sentence $(\alpha \Rightarrow \beta)$ is valid.

Forward direction: If $\alpha \models \beta$, then the sentence $(\alpha \Rightarrow \beta)$ is valid.

By definition, $\alpha \models \beta$ means that β is true in all worlds in which α is true. Thus, in all worlds in which α is true $\alpha \Rightarrow \beta$ holds because both α and β will be true. We must also consider worlds in which α is false; in these worlds, $\alpha \Rightarrow \beta$ also holds by definition of the falsehood of α .

Backward direction: If the sentence $(\alpha \Rightarrow \beta)$ is valid, then $\alpha \models \beta$.

If the sentence $(\alpha \Rightarrow \beta)$ is valid, then it is true in all worlds. Thus, for every world, it must be the case that either both α and β are true, *or* α is false. This is enough to tell us that in every world in which α is true, β is also true, which is the definition of entailment.

Inference: proof

- A **proof** ($\alpha \models \beta$) is a demonstration of entailment from α to β
 - Method 1: model checking 枚举
 - Truth table enumeration to check if $\text{models}(\alpha) \subseteq \text{models}(\beta)$
 - Time complexity always exponential in n ☹️

P1	P2	...	Pn	α	β
F	F	...	F	F	T
F	F	...	T	T	T
.....					
T	T	...	F	T	T
T	T	...	T	F	F

CNF 判定

Examples and non-examples [\[edit \]](#)

All of the following formulas in the variables A, B, C, D, E , and F are in conjunctive normal form:

- $(A \vee \neg B \vee \neg C) \wedge (\neg D \vee E \vee F)$
- $(A \vee B) \wedge (C)$
- $(A \vee B)$
- (A)

For clarity, the disjunctive clauses are written inside parentheses above. In [disjunctive normal form](#) with parenthesized conjunctive clauses, the last case is the same, but the next to last is $(A) \vee (B)$. The constants *true* and *false* are denoted by the empty conjunct and one clause consisting of the empty disjunct, but are normally written explicitly.^[1]

The following formulas are **not** in conjunctive normal form:

- $\neg(B \vee C)$, since an OR is nested within a NOT
- $(A \wedge B) \vee C$
- $A \wedge (B \vee (D \wedge E))$, since an AND is nested within an OR

Every formula can be equivalently written as a formula in conjunctive normal form. The three non-examples in CNF are:

- $(\neg B) \wedge (\neg C)$
- $(A \vee C) \wedge (B \vee C)$
- $(A) \wedge (B \vee D) \wedge (B \vee E)$.

都可

Conversion into CNF [\[edit \]](#)

^[2]Every [propositional formula](#) can be converted into an [equivalent](#) formula that is in CNF. This transformation is based on rules about [logical equivalences](#): [double negation elimination](#), [De Morgan's laws](#), and the [distributive law](#).

Since all propositional formulas can be converted into an equivalent formula in conjunctive normal form, proofs are often based on the assumption that all formulae are CNF. However, in some cases this conversion to CNF can lead to an exponential explosion of the formula. For example, translating the following non-CNF formula into CNF produces a formula with 2^n clauses:

$$(X_1 \wedge Y_1) \vee (X_2 \wedge Y_2) \vee \cdots \vee (X_n \wedge Y_n).$$

In particular, the generated formula is:

$$(X_1 \vee X_2 \vee \cdots \vee X_n) \wedge (Y_1 \vee X_2 \vee \cdots \vee X_n) \wedge (X_1 \vee Y_2 \vee \cdots \vee X_n) \wedge (Y_1 \vee Y_2 \vee \cdots \vee X_n) \wedge \cdots \wedge (Y_1 \vee Y_2 \vee \cdots \vee Y_n).$$

This formula contains 2^n clauses; each clause contains either X_i or Y_i for each i .

There exist transformations into CNF that avoid an exponential increase in size by preserving [satisfiability](#) rather than [equivalence](#).^{[3][4]} These transformations are guaranteed to only linearly increase the size of the formula, but introduce new variables. For example, the above formula can be transformed into CNF by adding variables Z_1, \dots, Z_n as follows:

$$(Z_1 \vee \cdots \vee Z_n) \wedge (\neg Z_1 \vee X_1) \wedge (\neg Z_1 \vee Y_1) \wedge \cdots \wedge (\neg Z_n \vee X_n) \wedge (\neg Z_n \vee Y_n).$$

An [interpretation](#) satisfies this formula only if at least one of the new variables is true. If this variable is Z_i , then both X_i and Y_i are true as well. This means that every [model](#) that satisfies this formula also satisfies the original one. On the other hand, only some of the models of the original formula satisfy this one: since the Z_i are not mentioned in the original formula, their values are irrelevant to satisfaction of it, which is not the case in the last formula. This means that the original formula and the result of the translation are [equisatisfiable](#) but not [equivalent](#).

An alternative translation, the [Tseitin transformation](#), includes also the clauses $Z_i \vee \neg X_i \vee \neg Y_i$. With these clauses, the formula implies $Z_i \equiv X_i \wedge Y_i$; this formula is often regarded to "define" Z_i to be a name for $X_i \wedge Y_i$.


推断 Inference: proof

- A **proof** ($\alpha \vdash \beta$) is a demonstration of entailment from α to β
 - Method 2: application of inference rules
 - Search for a finite sequence of sentences each of which is an **axiom** or follows from the preceding sentences by a **rule of inference**
 - Axiom: a sentence known to be true
 - Rule of inference: a function that takes one or more sentences (premises) and returns a sentence (conclusion)


自明
之理

先前所

假设



Inference: soundness & completeness

- **Sound** inference 
 - everything that can be proved is in fact entailed
- **Complete** inference
 - everything that is entailed can be proved
- Method 1 (enumeration) is obviously sound and complete
- For method 2 (applying inference rules), it is much less obvious
 - Example: arithmetic is found to be not complete! (Gödel's theorem, 1931)

Quiz

- What's the connection between complete inference algorithms and complete search algorithms?
- Answer 1: they both have the words “complete...algorithm”
- Answer 2: Formulate inference $\alpha \vdash \beta$ as a search problem
 - Initial state: KB contains α
 - Actions: apply any inference rule that matches KB, add conclusion
 - Goal test: KB contains β

Q
(< B)

Hence any complete search algorithm can be used to produce a complete inference algorithm

↓
guarantee to find a solution
if exist

Resolution: an inference rule in PL

- **Conjunctive Normal Form** (CNF) 子句
 - conjunction of disjunctions of literals (clauses)
 - Ex
 - $(A \vee \neg B) \wedge (B \vee \neg C \vee \neg D)$
 - $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$

≥ 2 entry

子句

Conversion to CNF

$$B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$$

step

1. Eliminate \Leftrightarrow , replacing $\alpha \Leftrightarrow \beta$ with $(\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)$.

$$(B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})) \wedge ((P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1})$$

2. Eliminate \Rightarrow , replacing $\alpha \Rightarrow \beta$ with $\neg \alpha \vee \beta$.

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1})$$

3. Move \neg inwards using de Morgan's rules and double-negation:

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge ((\neg P_{1,2} \wedge \neg P_{2,1}) \vee B_{1,1})$$


4. Apply distributivity law (\wedge over \vee) and flatten:

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$$

Resolution: an inference rule in PL

- Resolution** inference rule (for CNF):

Suppose l_i is $\neg m_j$

没有矛盾也成立。


entail

$$\begin{array}{c}
 l_1 \vee \dots \vee l_k, \quad (VC) \quad m_1 \vee \dots \vee m_n \quad (VC) \\
 \hline
 l_1 \vee \dots \vee l_{i-1} \vee l_{i+1} \vee \dots \vee l_k \vee m_1 \vee \dots \vee m_{j-1} \vee m_{j+1} \vee \dots \vee m_n
 \end{array}$$

(Note: The diagram shows a horizontal line separating the premises from the conclusion. The premises are $l_1 \vee \dots \vee l_k$ and $m_1 \vee \dots \vee m_n$, both with handwritten red '(VC)' next to them. The conclusion is the disjunction of all literals except the resolved pair l_i and m_j . A handwritten red arrow points from the text '没有矛盾也成立' to the first premise, and another red arrow points from the text 'entail' to the conclusion line.)

Examples:

$$\frac{P_{1,3} \vee P_{2,2}, \quad P_{2,3} \vee \neg P_{2,2}}{P_{1,3} \vee P_{2,3}}$$

(Note: In the original image, the terms $P_{2,2}$ and $\neg P_{2,2}$ in the numerator are circled in red.)

$$\frac{P_1, \neg P_1}{\{\}}$$

- Resolution is sound and complete for propositional logic

The **resolution rule** in propositional logic is a single valid inference rule that produces a new clause implied by two clauses containing complementary literals. A literal is a propositional variable or the negation of a propositional variable. Two literals are said to be complements if one is the negation of the other (in the following, $\neg c$ is taken to be the complement to c). The resulting clause contains all the literals that do not have complements. Formally:

$$\frac{a_1 \vee a_2 \vee \dots \vee c, \quad b_1 \vee b_2 \vee \dots \vee \neg c}{a_1 \vee a_2 \vee \dots \vee b_1 \vee b_2 \vee \dots}$$

where

all a_i , b_i , and c are literals,

the dividing line stands for "entails".

$$\frac{P \rightarrow Q, P}{Q}$$

$$\frac{\neg P \vee Q, P}{Q}$$

The above may also be written as:

$$\frac{(\neg a_1 \wedge \neg a_2 \wedge \dots) \rightarrow c, \quad c \rightarrow (b_1 \vee b_2 \vee \dots)}{(\neg a_1 \wedge \neg a_2 \wedge \dots) \rightarrow (b_1 \vee b_2 \vee \dots)}$$

Or schematically as:

$$\frac{\Gamma_1 \cup \{l\} \quad \Gamma_2 \cup \{\bar{l}\}}{\Gamma_1 \cup \Gamma_2} |l|$$

key

Example OF Propositional Resolution

Consider the following Knowledge Base:

1. The humidity is high or the sky is cloudy. $P \vee Q$
2. If the sky is cloudy, then it will rain. $Q \rightarrow R$
3. If the humidity is high, then it is hot. $P \rightarrow S$
4. It is not hot. $\neg S$

$$KB \text{ true} = \bigwedge_{i=1}^n \text{true} \quad (CNF)$$

Goal: It will rain.

Use propositional logic and apply resolution method to prove that the goal is derivable from the given knowledge base.

Solution: Let's construct propositions of the given sentences one by one:

1. Let, P: Humidity is high.

Q: Sky is cloudy.

It will be represented as $P \vee Q$.

- 2) Q: Sky is cloudy. ...from(1)

Let, R: It will rain.

It will be represented as $Q \rightarrow R$.

- 3) P: Humidity is high. ...from(1)

KB | x

Let, S: It is hot.

It will be represented as $P \rightarrow S$.

4) $\neg S$: It is not hot.


Applying resolution method:

In (2), $Q \rightarrow R$ will be converted as $(\neg Q \vee R)$

In (3), $P \rightarrow S$ will be converted as $(\neg P \vee S)$

Negation of Goal ($\neg R$): It will not rain.

Finally, apply the rule as shown below:



After applying Proof by Refutation (Contradiction) on the goal, the problem is solved, and it has terminated with a **Null clause** (\emptyset). Hence, the goal is achieved. ~~Thus, it is not raining.~~

LEMMA 3 $KB \vdash \alpha$

target: prove $KB \vdash \alpha$

$\alpha_1 \wedge \alpha_2 \wedge \dots \wedge \alpha_k$
is true.

KB: knowledge base. \uparrow
... of sentences $\alpha_1, \dots, \alpha_k$.

Collection of sentences we know are true.


Resolution algorithm

KB true, α true

- The best way to prove $KB \models \alpha$?
 - **Proof by contradiction**, i.e., show $KB \wedge \neg \alpha$ is unsatisfiable
- 1. Convert $KB \wedge \neg \alpha$ to **CNF**
- 2. Repeatedly apply the resolution rule to add new clauses, until one of the two things happens $\left. \begin{array}{l} \text{a)} \text{ Two clauses resolve to yield the empty clause, in which case } KB \text{ entails } \alpha \\ \text{b)} \text{ There is no new clause that can be added, in which case } KB \text{ does not entail } \alpha \end{array} \right\}$

Knowledge Base

A knowledge base is a collection of sentences $\alpha_1, \alpha_2, \dots, \alpha_k$ that we know are true, i.e., $\alpha_1 \wedge \alpha_2 \wedge \dots \wedge \alpha_k$. The sentences in the knowledge base allows us to compactly specify the allowable states of the world. By adding new sentences, the number of possible worlds consistent with our knowledge base could decrease (if we gain new information), go to zero (if the new sentence is inconsistent with the existing knowledge base), or remain the same (if the new sentence is entailed by existing sentences).




Resolution example

$$KB = (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1}$$

$$\alpha = \neg P_{1,2}$$

CNF

negative
target



Q: 没用到吧 ↗

↗

$$(B \Rightarrow P_{1,2} \vee P_{2,1}) \wedge (P_{1,2} \vee P_{2,1} \Rightarrow B)$$

$$\begin{aligned}
&= (\neg B \vee P_{1,2} \vee P_{2,1}) \wedge (\neg (P_{1,2} \vee P_{2,1}) \vee B) \\
&= () \wedge (\neg P_{1,2} \wedge \neg P_{2,1}) \vee B \\
&\quad \text{分配!} \\
&= () \wedge (\neg P_{1,2} \vee B) \wedge (\neg P_{2,1} \vee B)
\end{aligned}$$