

# CS 188: Artificial Intelligence

## Hidden Markov Models



Instructor: Anca Dragan --- University of California, Berkeley

[These slides were created by Dan Klein, Pieter Abbeel, and Anca. [http://ai.berkeley.edu.\]](http://ai.berkeley.edu.)

# Probability Recap

---

- Conditional probability

$$P(x|y) = \frac{P(x,y)}{P(y)}$$

- Product rule

$$P(x,y) = P(x|y)P(y)$$

- Chain rule

$$\begin{aligned} P(X_1, X_2, \dots, X_n) &= P(X_1)P(X_2|X_1)P(X_3|X_1, X_2)\dots \\ &= \prod_{i=1}^n P(X_i|X_1, \dots, X_{i-1}) \end{aligned}$$

- X, Y independent if and only if:  $\forall x, y : P(x,y) = P(x)P(y)$

- X and Y are conditionally independent given Z if and only if:

$$\forall x, y, z : P(x, y|z) = P(x|z)P(y|z)$$



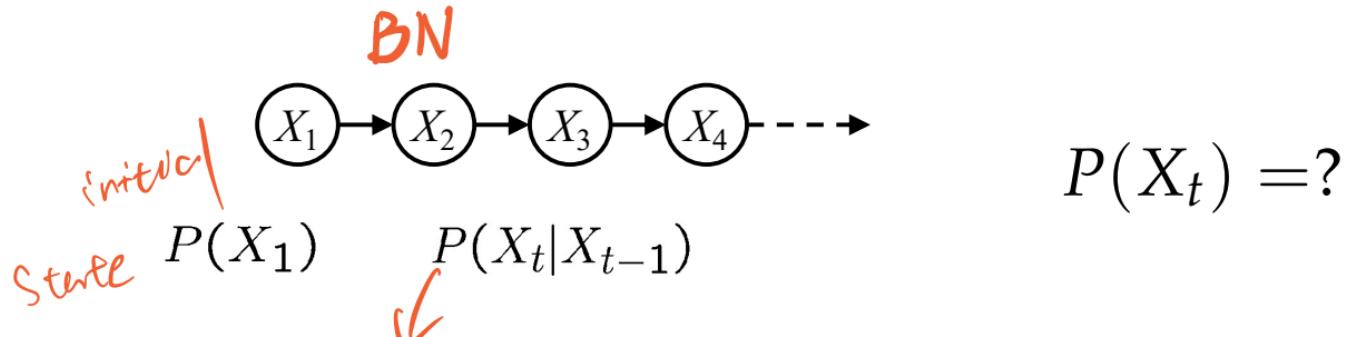
# Reasoning over Time or Space

---

- Often, we want to **reason about a sequence** of observations
  - Speech recognition
  - Robot localization
  - User attention
  - Medical monitoring
- Need to introduce **time (or space)** into our models

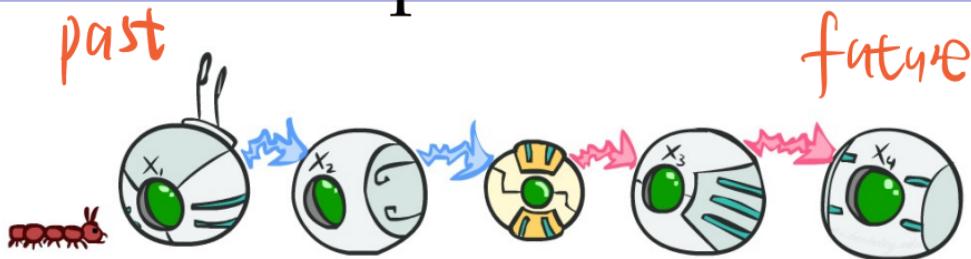
# Markov Models

- Value of  $X$  at a given time is called the **state**



- Parameters: called **transition probabilities** or dynamics, specify how the state evolves over time (also, initial state probabilities)
- Stationarity assumption: transition probabilities the same at all times
- Same as MDP transition model, but no choice of action
- A (growable) BN: We can always use generic BN reasoning on it if we truncate the chain at a fixed length

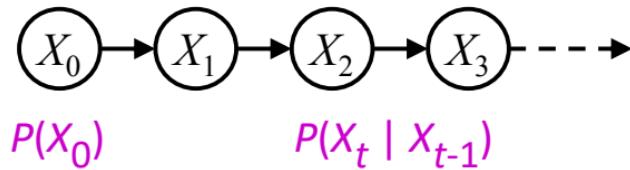
# Markov Assumption: Conditional Independence



- Basic conditional independence:
  - Past and future independent given the present
  - Each time step only depends on the previous
  - This is called the (first order) Markov property

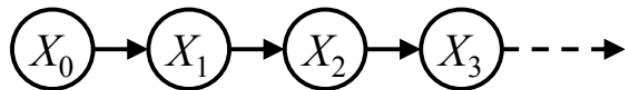
# Markov Models (aka Markov chain/process)

- Value of  $X$  at a given time is called the **state** (usually discrete, finite)



- The **transition model**  $P(X_t | X_{t-1})$  specifies how the state evolves over time
- Stationarity** assumption: same transition probabilities at all time steps
- Joint distribution  $P(X_0, \dots, X_T) = P(X_0) \prod_t P(X_t | X_{t-1})$

# Quiz: are Markov models a special case of Bayes nets?



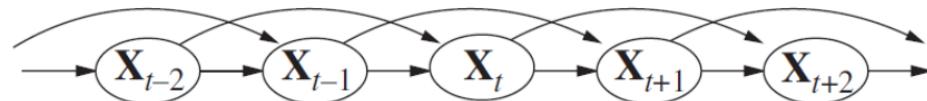
- Yes and no!
- Yes:
  - Directed acyclic graph, joint = product of conditionals
- No:
  - Infinitely many variables (unless we truncate)
  - Repetition of transition model not part of standard Bayes net syntax

截断

# Markov Assumption: Conditional Independence



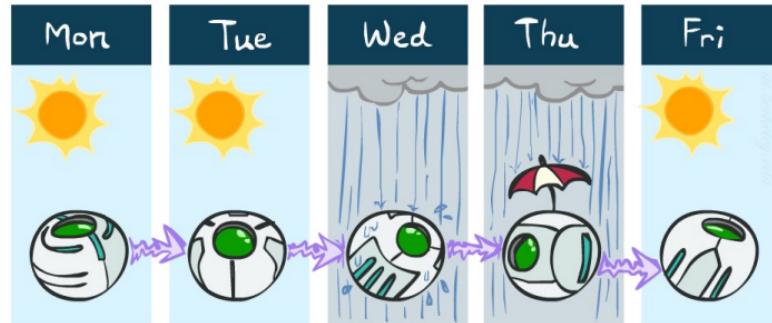
- **Markov** assumption:  $X_{t+1}, \dots$  are independent of  $X_0, \dots, X_{t-1}$  given  $X_t$ 
  - Past and future independent given the present
  - Each time step only depends on the previous
- This is a **first-order** Markov model
- A  $k$ th-order model allows dependencies on  $k$  earlier steps



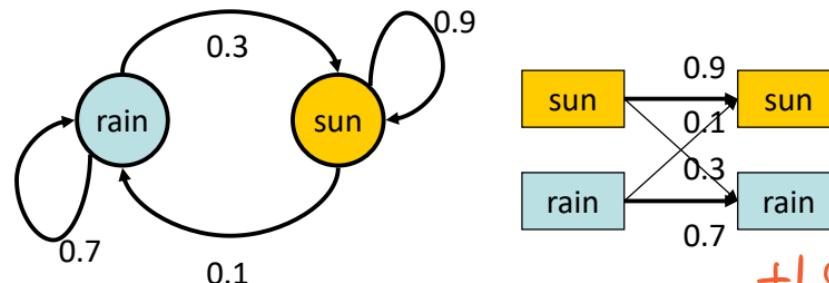
# Example Markov Chain: Weather

- States:  $X = \{\text{rain, sun}\}$
- Initial distribution: 1.0 sun
- CPT  $P(X_t | X_{t-1})$ :

$X_{t-1}$	$X_t$	$P(X_t   X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7

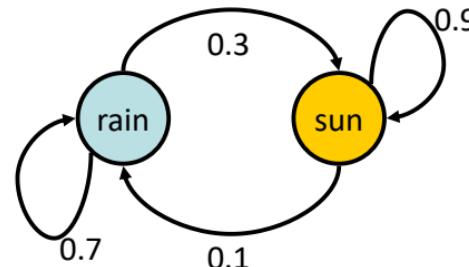


Two new ways of representing the same CPT



# Example Markov Chain: Weather

- Initial distribution: 1.0 sun



- What is the probability distribution after one step?

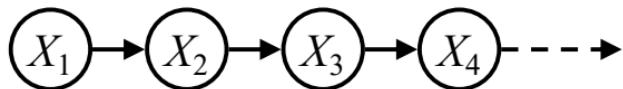
$$P(X_2 = \text{sun}) = \sum_{x_1} P(x_1, X_2 = \text{sun}) = \sum_{x_1} P(X_2 = \text{sun}|x_1)P(x_1)$$

$$\begin{aligned} P(X_2 = \text{sun}) = & P(X_2 = \text{sun}|X_1 = \text{sun})P(X_1 = \text{sun}) + \\ & P(X_2 = \text{sun}|X_1 = \text{rain})P(X_1 = \text{rain}) \end{aligned}$$

$$0.9 \cdot 1.0 + 0.3 \cdot 0.0 = 0.9$$

# Mini-Forward Algorithm

- Question: What's  $P(X)$  on some day  $t$ ?

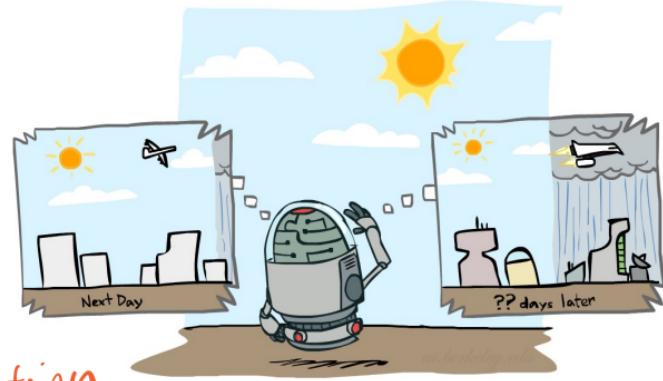


$P(x_1)$  = known

$$\begin{aligned} P(x_t) &= \sum_{x_{t-1}} P(x_{t-1}, x_t) \\ &= \sum_{x_{t-1}} P(x_t | x_{t-1}) P(x_{t-1}) \end{aligned}$$

update equation.

Forward simulation



# Weather prediction

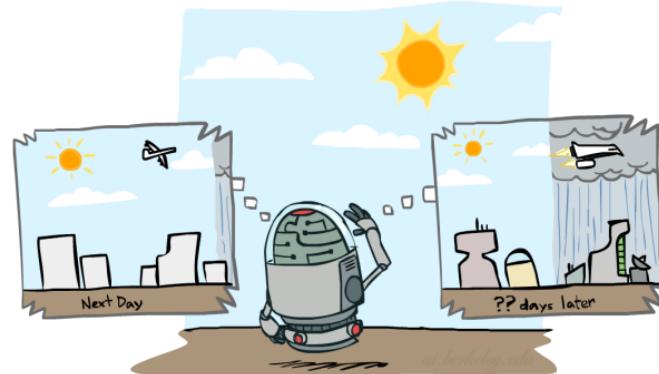
- Time 0:  $\langle 0.5, 0.5 \rangle$

$x_{t-1}$	$P(x_t   x_{t-1})$	
	sun	rain
sun	0.9	0.1
rain	0.3	0.7

- What is the weather like at time 1?

$$\begin{aligned} P(X_1) &= \sum_{x_0} P(X_1, X_0=x_0) \\ &= \sum_{x_0} P(X_0=x_0) P(X_1 | X_0=x_0) \\ &= 0.5 \langle 0.9, 0.1 \rangle + 0.5 \langle 0.3, 0.7 \rangle = \langle 0.6, 0.4 \rangle \end{aligned}$$

T, F ...



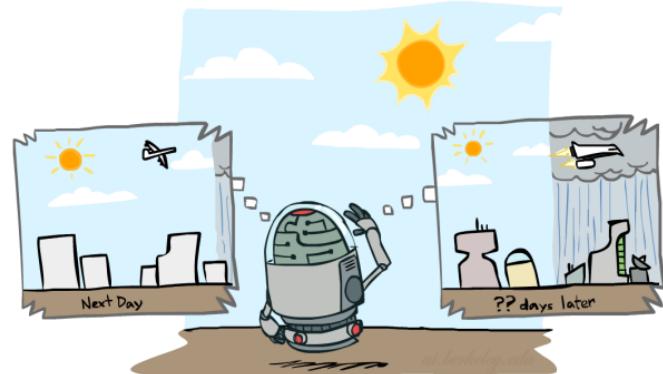
# Weather prediction, contd.

- Time 1:  $\langle 0.6, 0.4 \rangle$

$x_{t-1}$	$P(x_t   x_{t-1})$	
	sun	rain
sun	0.9	0.1
rain	0.3	0.7

- What is the weather like at time 2?

- $$P(X_2) = \sum_{x_1} P(X_2, X_1=x_1)$$
- $$= \sum_{x_1} P(X_1=x_1) P(X_2 | X_1=x_1)$$
- $$= 0.6 \langle 0.9, 0.1 \rangle + 0.4 \langle 0.3, 0.7 \rangle = \langle 0.66, 0.34 \rangle$$



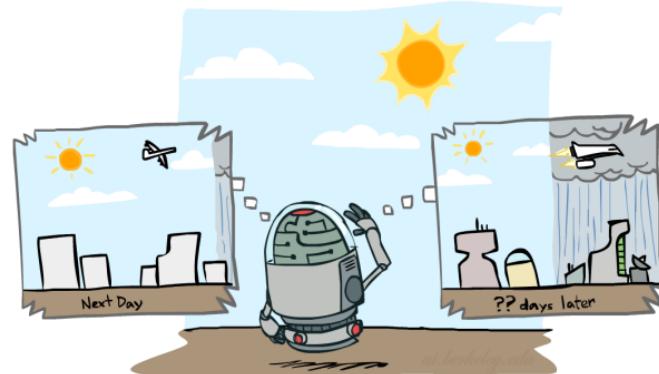
# Weather prediction, contd.

- Time 2:  $<0.66, 0.34>$

$x_{t-1}$	$P(x_t   x_{t-1})$	
	sun	rain
sun	0.9	0.1
rain	0.3	0.7

- What is the weather like at time 3?

- $$P(X_3) = \sum_{x_2} P(X_3, X_2=x_2)$$
- $$= \sum_{x_2} P(X_2=x_2) P(X_3 | X_2=x_2)$$
- $$= 0.66<0.9, 0.1> + 0.34<0.3, 0.7> = <0.696, 0.304>$$

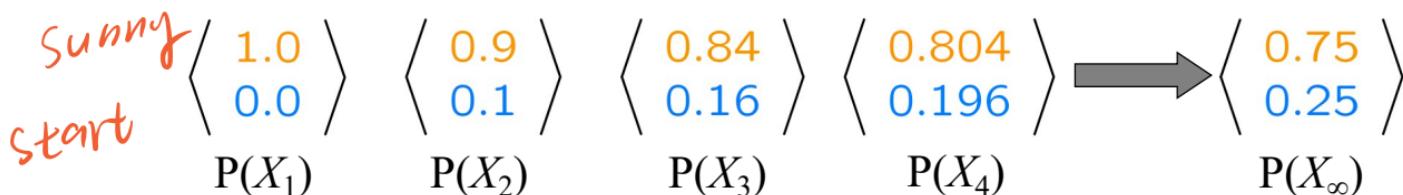


# Forward algorithm (simple form)

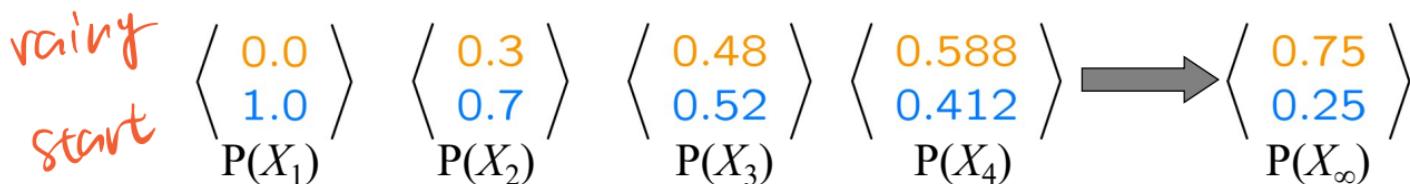
- What is the state at time  $t$  (given an initial distribution  $P(X_0)$ )?
  - $$P(X_t) = \sum_{X_{t-1}} P(X_t | X_{t-1}=x_{t-1})$$
  - $$= \sum_{X_{t-1}} P(X_{t-1}=x_{t-1}) P(X_t | X_{t-1}=x_{t-1})$$
- Iterate this update starting at  $t=0$

# Example Run of Mini-Forward Algorithm

- From initial observation of sun



- From initial observation of rain



- From yet another initial distribution  $P(X_1)$ :



# Stationary Distributions

- For most chains:
  - Influence of the initial distribution gets less and less over time.
  - The distribution we end up in is independent of the initial distribution

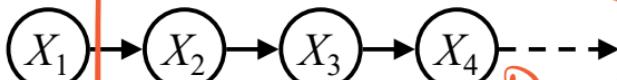
- Stationary distribution:
    - The distribution we end up with is called the **stationary distribution**  $P_\infty$  of the chain  $P(X) = \sum_x P(x'|x) \cdot P(x)$
    - It satisfies  $\sum P(X) = 1$
- $$P_\infty(X) = P_{\infty+1}(X) = \sum_x P(X|x) P_\infty(x)$$

Solve it to find  $\infty$



# Example: Stationary Distributions

Question: What's  $P(X)$  at time  $t = \infty$ ?



$$P_{\infty}(\text{sun}) = P(\text{sun}|\text{sun})P_{\infty}(\text{sun}) + P(\text{sun}|\text{rain})P_{\infty}(\text{rain})$$

$$P_{\infty}(\text{rain}) = P(\text{rain}|\text{sun})P_{\infty}(\text{sun}) + P(\text{rain}|\text{rain})P_{\infty}(\text{rain})$$

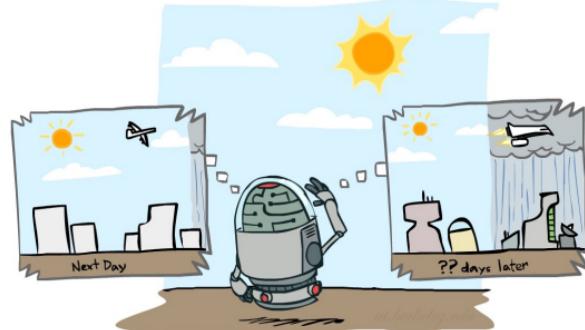
$$P_{\infty}(\text{sun}) = 0.9P_{\infty}(\text{sun}) + 0.3P_{\infty}(\text{rain})$$

$$P_{\infty}(\text{rain}) = 0.1P_{\infty}(\text{sun}) + 0.7P_{\infty}(\text{rain})$$

$$P_{\infty}(\text{sun}) = 3P_{\infty}(\text{rain})$$

$$P_{\infty}(\text{rain}) = 1/3P_{\infty}(\text{sun})$$

Also:  $\checkmark P_{\infty}(\text{sun}) + P_{\infty}(\text{rain}) = 1$



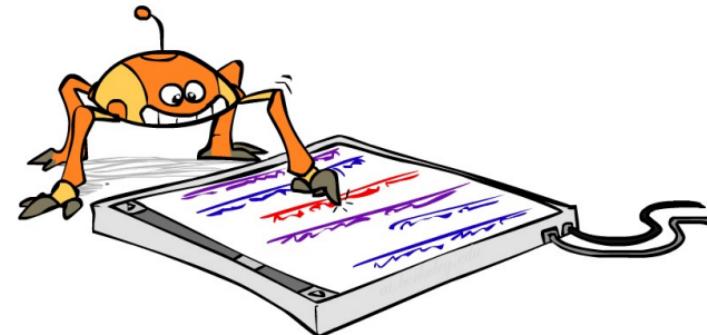
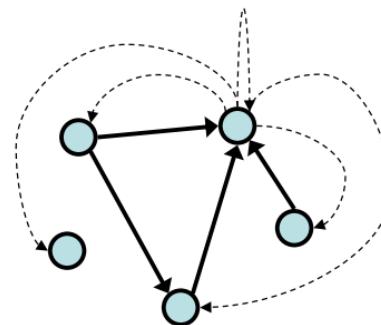
$X_{t-1}$	$X_t$	$P(X_t   X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7

$$P_{\infty}(\text{sun}) = 3/4$$

$$P_{\infty}(\text{rain}) = 1/4$$

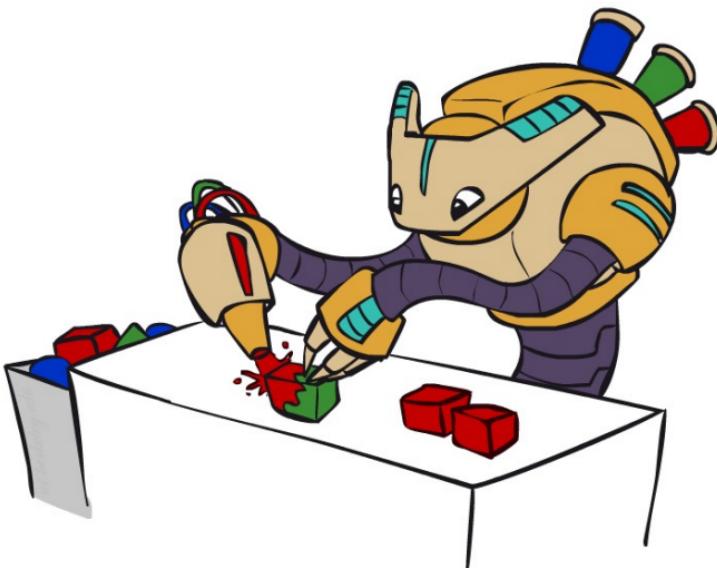
# Application of Stationary Distribution: Web Link Analysis

- PageRank over a web graph
  - Each web page is a possible value of a state
  - Initial distribution: uniform over pages
  - Transitions:
    - With prob.  $c$ , uniform jump to a random page (dotted lines, not all shown)
    - With prob.  $1-c$ , follow a random outlink (solid lines)
- Stationary distribution
  - Will spend more time on highly reachable pages
  - E.g. many ways to get to the Acrobat Reader download page
  - Somewhat robust to link spam.
  - Google 1.0 returned the set of pages containing all your keywords in decreasing rank, now all search engines use link analysis along with many other factors (rank actually getting less important over time)



# Application of Stationary Distributions: Gibbs Sampling\*

- Each joint instantiation over all hidden and query variables is a state:  $\{X_1, \dots, X_n\} = H \cup Q$
- Transitions:
  - With probability  $1/n$  resample variable  $X_j$  according to
$$P(X_j | x_1, x_2, \dots, x_{j-1}, x_{j+1}, \dots, x_n, e_1, \dots, e_m)$$
- Stationary distribution:
  - Conditional distribution  $P(X_1, X_2, \dots, X_n | e_1, \dots, e_m)$
  - Means that when running Gibbs sampling long enough we get a sample from the desired distribution
  - Requires some proof to show this is true!



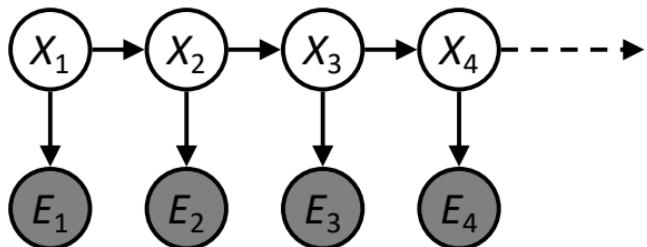
# Hidden Markov Models

HM

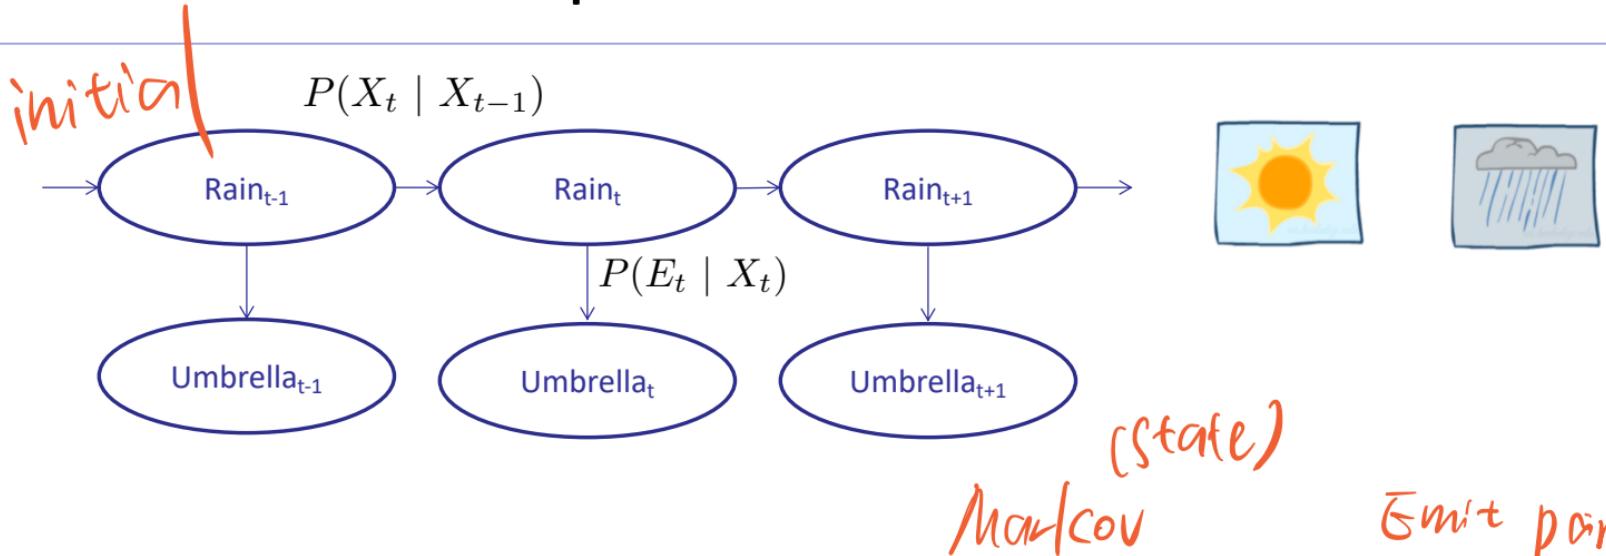


# Hidden Markov Models

- Markov chains not so useful for most agents
  - Need observations to update your beliefs
- Hidden Markov models (HMMs)
  - Underlying Markov chain over states X
  - You observe outputs (effects) at each time step



# Example: Weather HMM



- An HMM is defined by:
  - Initial distribution:  $P(X_1)$
  - Transitions:  $P(X_t | X_{t-1})$
  - Emissions:  $P(E_t | X_t)$

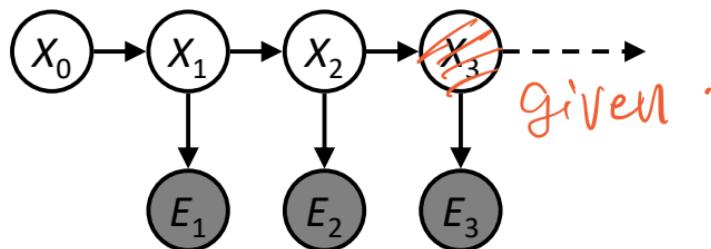
$R_{t-1}$	$R_t$	$P(R_t   R_{t-1})$
+r	+r	0.7
+r	-r	0.3
-r	+r	0.3
-r	-r	0.7

$R_t$	$U_t$	$P(U_t   R_t)$
+r	+u	0.9
+r	-u	0.1
-r	+u	0.2
-r	-u	0.8

# HMM as probability model

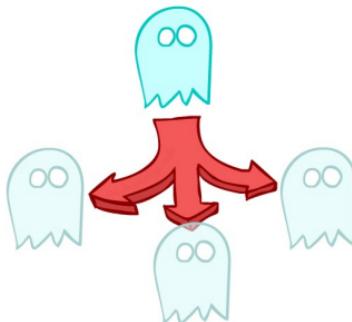
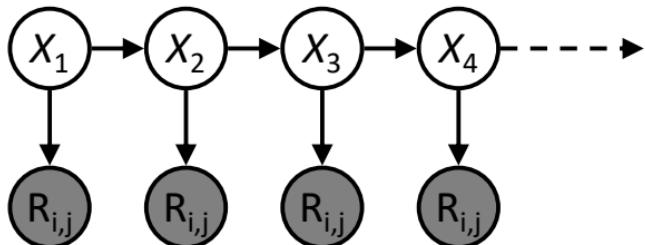
- Joint distribution for Markov model:  $P(X_0, \dots, X_T) = P(X_0) \prod_{t=1:T} P(X_t | X_{t-1})$
- Joint distribution for hidden Markov model:  
 $P(X_0, X_1, E_1, \dots, X_T, E_T) = P(X_0) \prod_{t=1:T} P(X_t | X_{t-1}) P(E_t | X_t)$
- Independence in HMM
  - Future states are independent of the past given the present
  - Current evidence is independent of everything else given the current state

Basic



# Example: Ghostbusters HMM

- $P(X_1) = \text{uniform}$
- $P(X|X')$  = usually move clockwise, but sometimes move in a random direction or stay in place
- $P(R_{ij}|X)$  = same sensor model as before: red means close, green means far away.



1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

$P(X_1)$

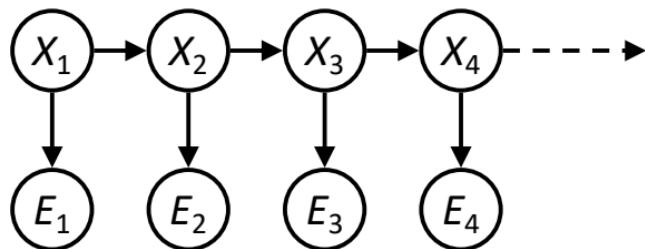


1/6	1/6	1/2
0	1/6	0
0	0	0

$P(X|X' = <1,2>)$

# Conditional Independence

- HMMs have two important independence properties:
  - Markov hidden process: future depends on past via the present
  - Current observation independent of all else given current state



$E_1$  inde of  $E_4$ .  
given  $X_2$ .

- Does this mean that evidence variables are guaranteed to be independent?
  - [No, they tend to correlate by the hidden state]

# Real HMM Examples

---

- Robot tracking:

- $E$  ○ Observations are range readings (continuous)
- $S$  ○ States are positions on a map (continuous)

- Speech recognition HMMs:

- $E$  ○ Observations are acoustic signals (continuous valued)
- $S$  ○ States are specific positions in specific words (so, tens of thousands)

- Machine translation HMMs:

- $E$  ○ Observations are words (tens of thousands)
- $S$  ○ States are translation options

# Inference tasks

- Useful notation:  $X_{a:b} = X_a, X_{a+1}, \dots, X_b$
- **Filtering**:  $P(X_t | e_{1:t})$ 
  - **belief state** — posterior distribution over the most recent state given all evidence
- **Prediction**:  $P(X_{t+k} | e_{1:t})$  for  $k > 0$ 
  - posterior distribution over a future state given all evidence
- **Smoothing**:  $P(X_k | e_{1:t})$  for  $0 \leq k < t$ 
  - posterior distribution over a past state given all evidence
- **Most likely explanation**:  $\arg \max_{x_{0:t}} P(x_{0:t} | e_{1:t})$ 
  - Ex: speech recognition, decoding with a noisy channel

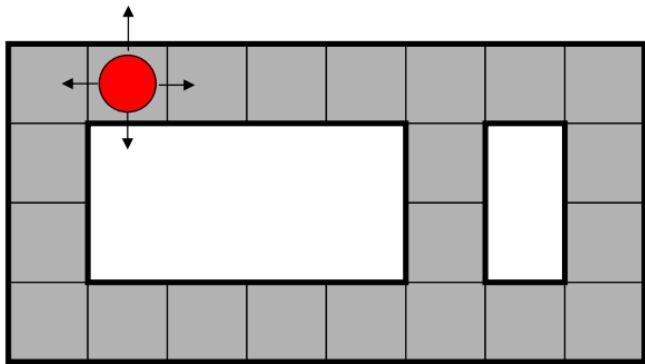
# Filtering / Monitoring

---

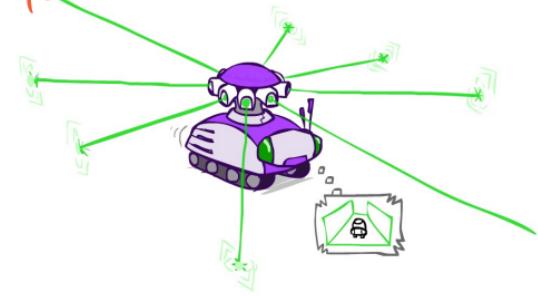
- Filtering, or monitoring, is the task of tracking the distribution  $B_t(X) = P_t(X_t \mid \underline{e_1, \dots, e_t})$  (the belief state) over time  
*evidence*
- We start with  $B_1(X)$  in an initial setting, usually uniform
- As time passes, or we get observations, we update  $B(X)$
- The Kalman filter was invented in the 60's and first implemented as a method of trajectory estimation for the Apollo program

# Example: Robot Localization

Example from  
Michael Pfeiffer



Hidden state:  
robot location



感知

Prob



t=0

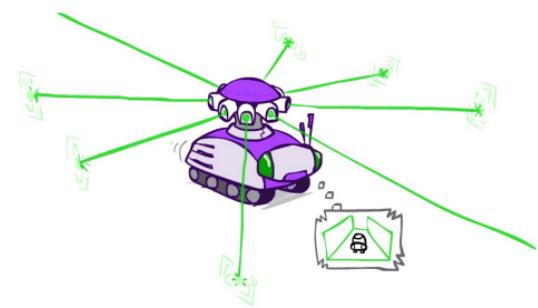
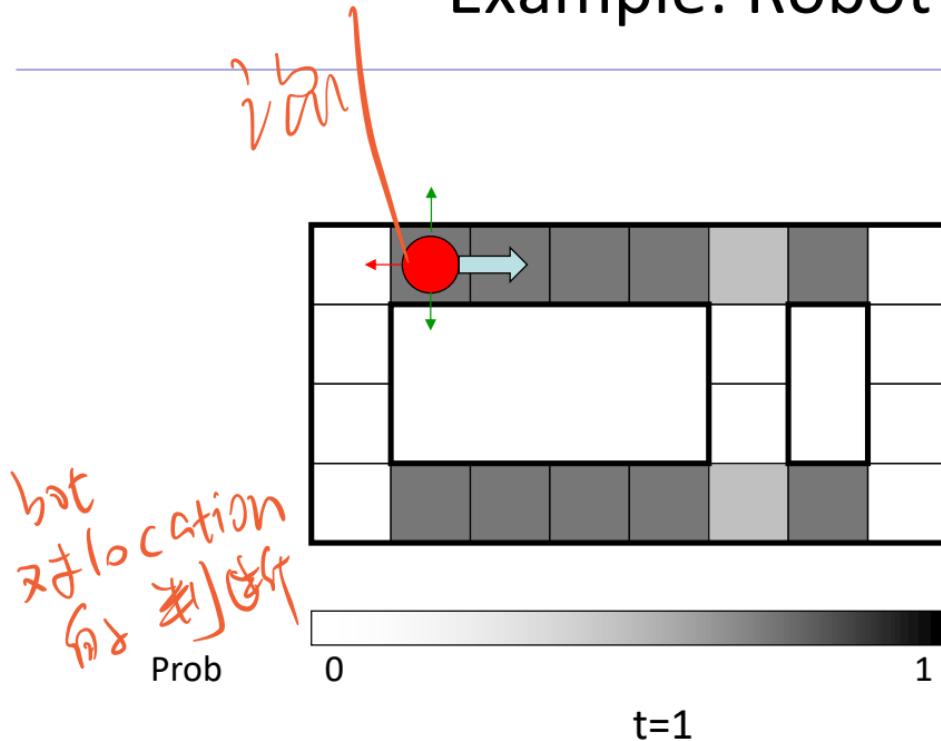
4 bits wall/no wall for 4 direction

Sensor model: can read in which directions there is a wall,  
never more than 1 mistake

运动

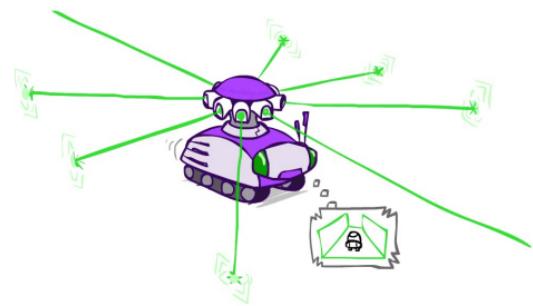
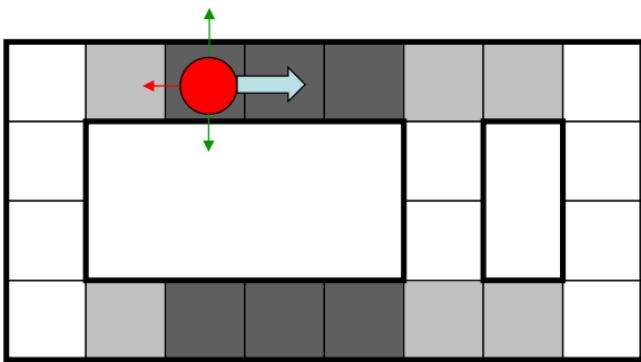
Motion model: may not execute action with small prob.

# Example: Robot Localization



# Example: Robot Localization

np.



Prob

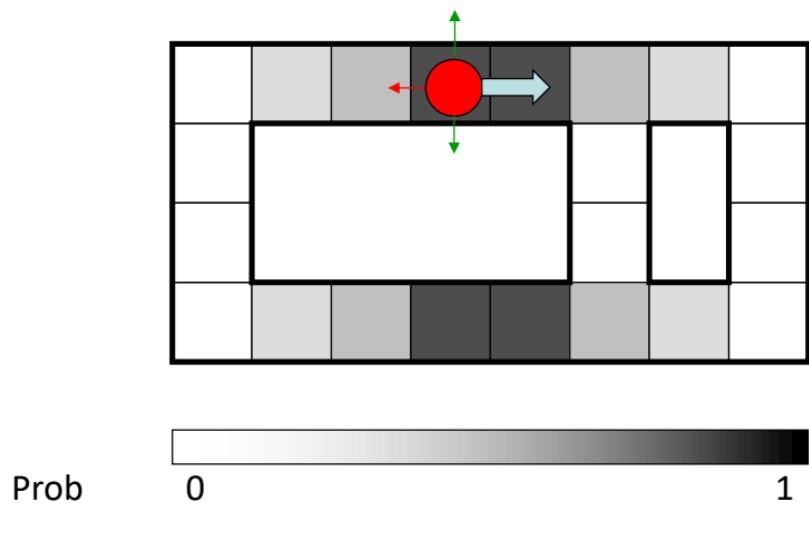
0

1

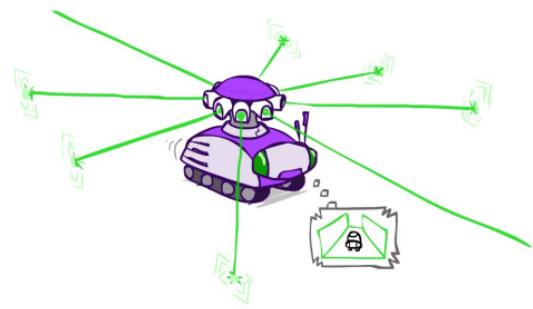
$t=2$

# Example: Robot Localization

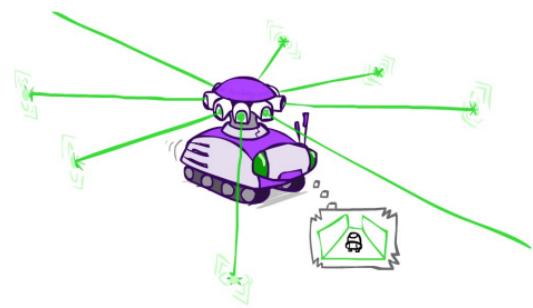
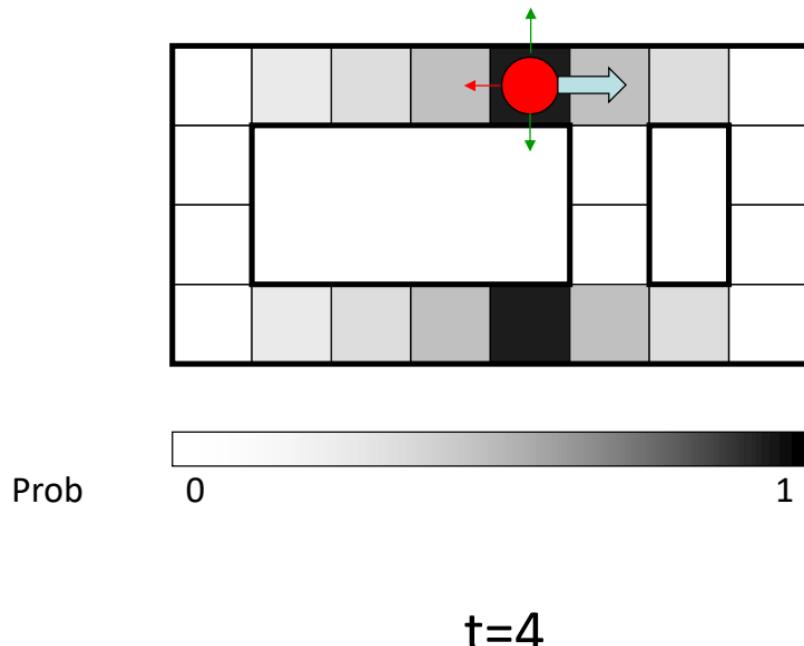
---



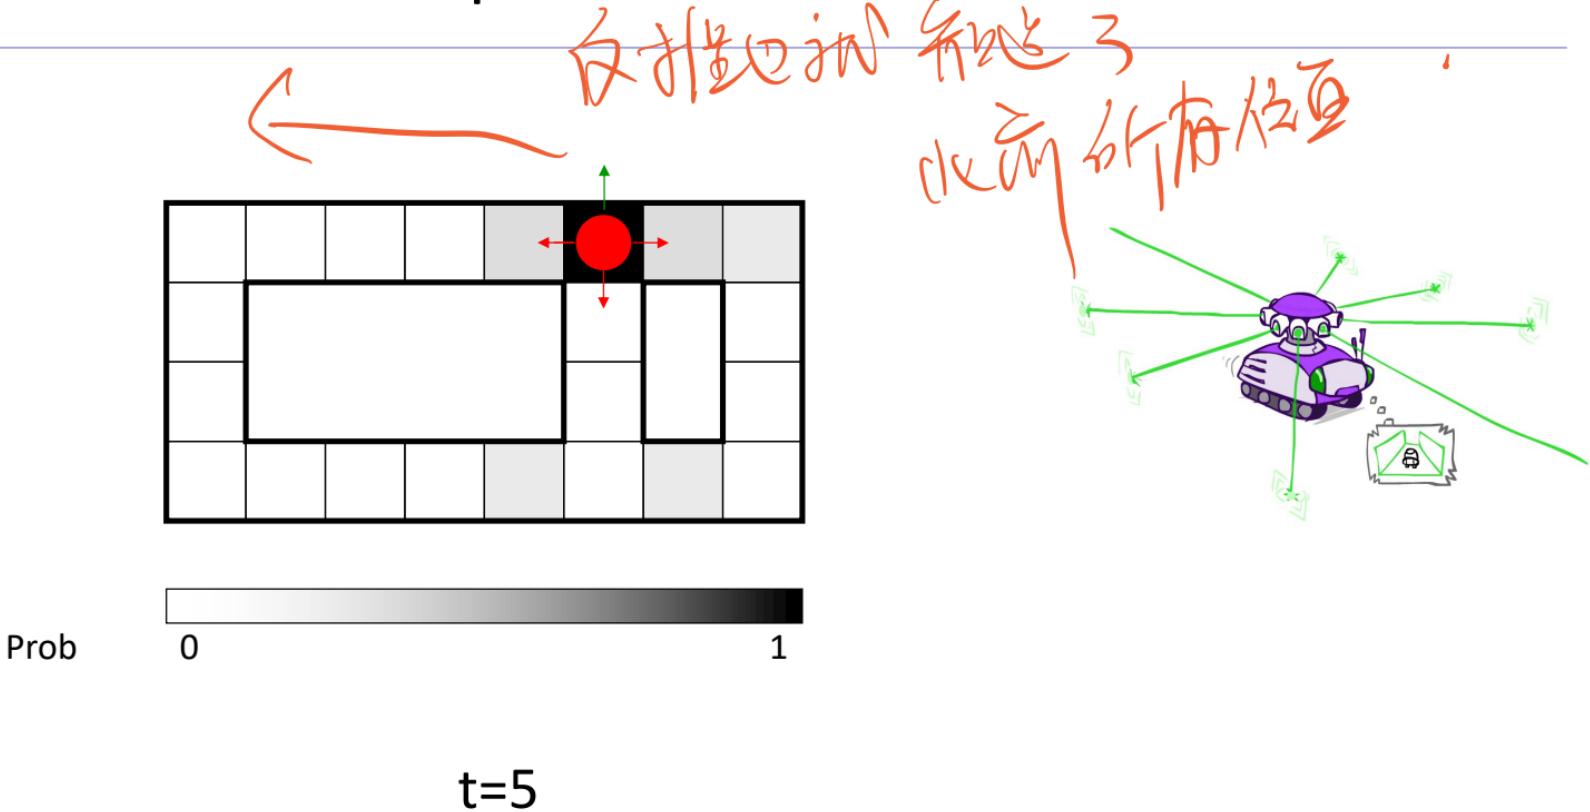
$t=3$



# Example: Robot Localization



# Example: Robot Localization



# Inference: Find State Given Evidence

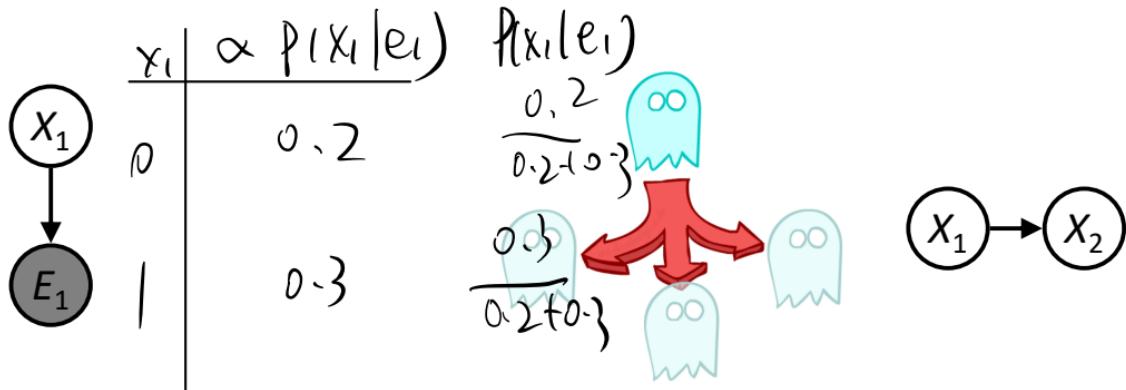
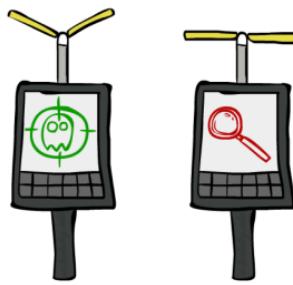
---

- We are given evidence at each time and want to know

$$B_t(X) = P(X_t | e_{1:t})$$

- Idea: start with  $P(X_1)$  and derive  $B_t$  in terms of  $B_{t-1}$ 
  - equivalently, derive  $B_{t+1}$  in terms of  $B_t$

# Inference: Base Cases



$$P(X_1|e_1)$$

$$P(x_1|e_1) = P(x_1, e_1) / \underbrace{P(e_1)}_{\text{constant}}$$

$$\propto_{X_1} P(x_1, e_1) = P(e_1, x_1)$$

$$= P(x_1)P(e_1|x_1)$$

mean: when consider only  $x$  as variable proportion

$$P(X_2)$$

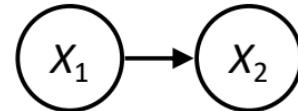
$$P(x_2) = \sum_{x_1} P(x_1, x_2)$$

$$= \sum_{x_1} P(x_1)P(x_2|x_1)$$

# Passage of Time

- Assume we have current belief  $P(X \mid \text{evidence to date})$

$$B(X_t) = P(X_t | e_{1:t})$$



- Then, after one time step passes:

$$P(X_{t+1} | e_{1:t}) = \sum_{x_t} P(X_{t+1}, x_t | e_{1:t})$$

Some independent. (from Marlov)

$$\begin{aligned} P(a, b | c) &= \sum_{x_t} P(X_{t+1} | x_t, e_{1:t}) P(x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t) P(x_t | e_{1:t}) \end{aligned}$$

known • Or compactly:

$$B'(X_{t+1}) = \sum_{x_t} P(X' | x_t) B(x_t)$$

- Basic idea: beliefs get “pushed” through the transitions
  - With the “B” notation, we have to be careful about what time step  $t$  the belief is about, and what evidence it includes

# Example: Passage of Time

- As time passes, uncertainty “accumulates”

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	1.00	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01

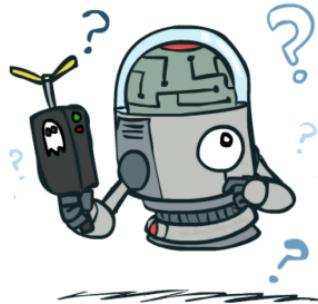
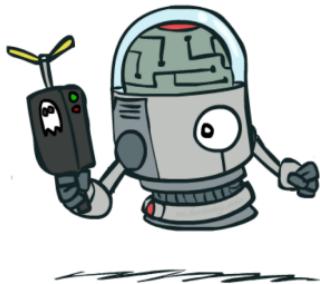
T = 1

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01

T = 2

0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01

T = 5



(Transition model: ghosts usually go clockwise)

“Believe”

# Observation

- Assume we have current belief  $P(X \mid \text{previous evidence})$ :

$$B'(X_{t+1}) = P(X_{t+1} | e_{1:t})$$

- Then, after evidence comes in:

$$P(X_{t+1} | e_{1:t+1}) = P(X_{t+1}, e_{t+1} | e_{1:t}) / P(e_{t+1} | e_{1:t})$$

$$\propto_{X_{t+1}} P(X_{t+1}, e_{t+1} | e_{1:t}) \quad \text{normalize factor}$$

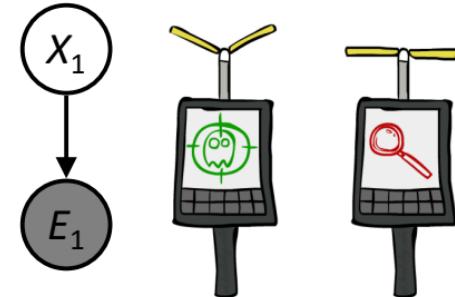
$$= P(e_{t+1} | e_{1:t}, X_{t+1}) P(X_{t+1} | e_{1:t})$$

$$= P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t})$$

present

Or, compactly:

$$\underline{B(X_{t+1})} \propto_{X_{t+1}} \underline{P(e_{t+1} | X_{t+1})} \underline{B'(X_{t+1})} \text{ past}$$



- Basic idea: beliefs “reweighted” by likelihood of evidence
- Unlike passage of time, we have to renormalize

# (new) Example: Observation

- As we get observations, beliefs get reweighted, uncertainty “decreases”

More known

More accurate



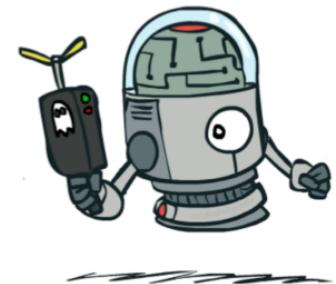
0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01

Before observation

<0.01	<0.01	<0.01	<0.01	0.02	<0.01
<0.01	<0.01	<0.01	0.83	0.02	<0.01
<0.01	<0.01	0.11	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01

After observation

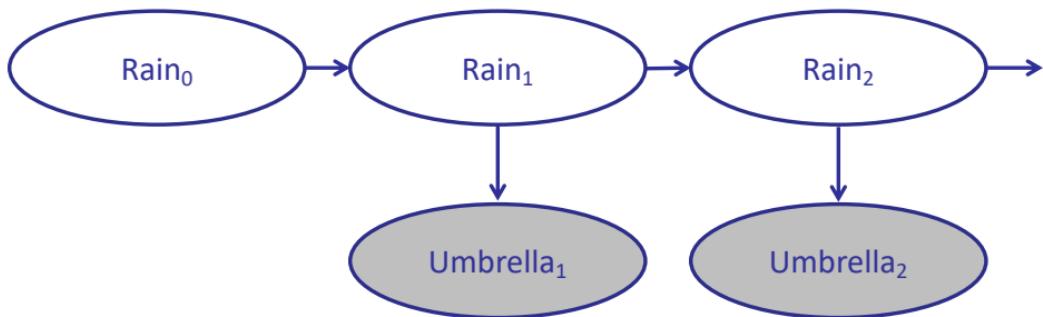
$$B(X) \propto P(e|X)B'(X)$$



# Example: Weather HMM



$$\begin{array}{ll} \text{B}(+r) = 0.5 & \text{B}'(+r) = 0.5 \\ \text{B}(-r) = 0.5 & \text{B}'(-r) = 0.5 \\ & \downarrow \\ \text{B}(+r) = 0.818 & \text{B}'(+r) = 0.627 \\ \text{B}(-r) = 0.182 & \text{B}'(-r) = 0.373 \\ & \downarrow \\ \text{B}(+r) = 0.883 & \\ \text{B}(-r) = 0.117 & \end{array}$$



R <sub>t</sub>	R <sub>t+1</sub>	P(R <sub>t+1</sub>  R <sub>t</sub> )
+r	+r	0.7
+r	-r	0.3
-r	+r	0.3
-r	-r	0.7

R <sub>t</sub>	U <sub>t</sub>	P(U <sub>t</sub>  R <sub>t</sub> )
+r	+u	0.9
+r	-u	0.1
-r	+u	0.2
-r	-u	0.8

# The Forward Algorithm

- We are given evidence at each time and want to know

$$B_t(X) = P(X_t | e_{1:t})$$

- We can derive the following updates

$$\begin{aligned} P(x_t | e_{1:t}) &\propto_X P(x_t, e_{1:t}) \xrightarrow{\text{P}(e|t)} \\ &= \sum_{x_{t-1}} P(x_{t-1}, x_t, e_{1:t}) \\ &= \sum_{x_{t-1}} P(x_{t-1}, e_{1:t-1}) \underbrace{P(x_t | x_{t-1})}_{\text{transition}} \underbrace{P(e_t | x_t)}_{\text{evidence}} \\ &= P(e_t | x_t) \sum_{x_{t-1}} P(x_t | x_{t-1}) P(x_{t-1}, e_{1:t-1}) \end{aligned}$$

We can normalize as we go if we want to have  $P(x|e)$  at each time step, or just once at the end...

# Online Belief Updates

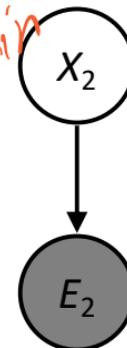
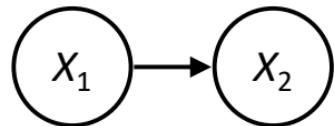
- Every time step, we start with current  $P(X \mid \text{evidence})$
- We update for time:

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} \underbrace{P(x_{t-1} | e_{1:t-1})}_{\text{trans.}} \cdot \underbrace{P(x_t | x_{t-1})}_{\text{}}$$

Can do time update  
evidence again and again  
until find evidence

$$P(x_t | e_{1:t}) \propto_X P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$

- The forward algorithm does both at once (and doesn't normalize)



# Filtering algorithm

- Filtering: infer current state given all evidence
- Aim: a **recursive filtering** algorithm of the form

- $P(X_{t+1}|e_{1:t+1}) = g(e_{t+1}, P(X_t|e_{1:t}))$

Apply Bayes' rule

- $$\begin{aligned} P(X_{t+1}|e_{1:t+1}) &= P(\underline{X_{t+1}} | \underline{e_{1:t}}, \underline{e_{t+1}}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

$$\alpha = 1 / P(e_{t+1} | e_{1:t})$$

# (Observation) Filtering algorithm

- Filtering: infer current state given all evidence  $X_{t+1}$ ,  $e_{1:t+1}$
  - Aim: a **recursive filtering** algorithm of the form
    - $P(X_{t+1}|e_{1:t+1}) = g(e_{t+1}, P(X_t|e_{1:t}))$   
↑ func      observe evidence  
want:  $P(e|X)$
  - $P(X_{t+1}|e_{1:t+1}) = P(X_{t+1}|e_{1:t}, e_{t+1})$   
Apply conditional independence
  - $= \alpha P(e_{t+1}|X_{t+1}, e_{1:t}) P(X_{t+1}|e_{1:t})$
  - $= \alpha P(e_{t+1}|X_{t+1}) P(X_{t+1}|e_{1:t})$
- $X = \eta(e_{t+1}|e_t)$
- Normalize      Update      Predict

# Bayes' Rule

- Two ways to factor a joint distribution over two variables:

$$P(x, y) = P(x|y)P(y) = P(y|x)P(x)$$

- Dividing, we get:

$$P(x|y) = \frac{P(y|x)}{P(y)}P(x)$$

- Why is this at all helpful?

- Lets us build one conditional from its reverse
- Often one conditional is tricky but the other one is simple
- Foundation of many systems we'll see later



- In the running for most important AI equation!

That's my rule!



# Filtering algorithm

- Filtering: infer current state given all evidence
- Aim: a **recursive filtering** algorithm of the form
  - $P(X_{t+1}|e_{1:t+1}) = g(e_{t+1}, P(X_t|e_{1:t}))$

$$\begin{aligned} \boxed{\quad} \quad & P(X_{t+1}|e_{1:t+1}) = P(X_{t+1}|e_{1:t}, e_{t+1}) \\ & = \alpha P(e_{t+1}|X_{t+1}, e_{1:t}) P(X_{t+1}|e_{1:t}) \quad \text{Condition on } X_t \\ & = \alpha P(e_{t+1}|X_{t+1}) \underline{P(X_{t+1}|e_{1:t})} \\ & = \alpha P(e_{t+1}|X_{t+1}) \sum_{x_t} P(x_t|e_{1:t}) P(X_{t+1}|x_t, e_{1:t}) \end{aligned}$$

“Believe”

# Observation

- Assume we have current belief  $P(X \mid \text{previous evidence})$ :

$$B'(X_{t+1}) = P(X_{t+1} | e_{1:t})$$

- Then, after evidence comes in:

$$P(X_{t+1} | e_{1:t+1}) = P(X_{t+1}, e_{t+1} | e_{1:t}) / P(e_{t+1} | e_{1:t})$$

$$\propto_{X_{t+1}} P(X_{t+1}, e_{t+1} | e_{1:t}) \quad \text{normalize factor}$$

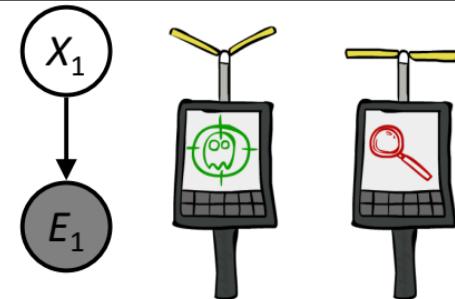
$$= P(e_{t+1} | e_{1:t}, X_{t+1}) P(X_{t+1} | e_{1:t})$$

$$= P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t})$$

present

Or, compactly:

$$\underline{B(X_{t+1})} \propto_{X_{t+1}} \underline{P(e_{t+1} | X_{t+1})} \underline{B'(X_{t+1})} \text{ past}$$



- Basic idea: beliefs “reweighted” by likelihood of evidence
- Unlike passage of time, we have to renormalize

# Filtering algorithm

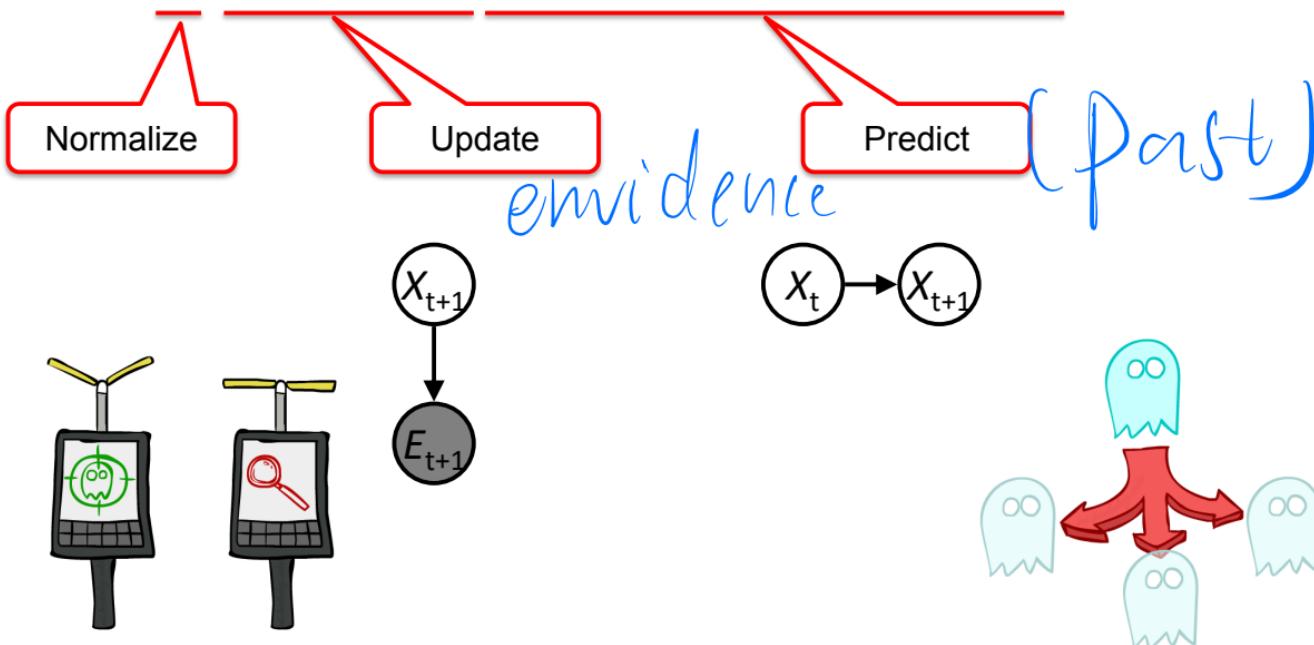
- Filtering: infer current state given all evidence
- Aim: a **recursive filtering** algorithm of the form
  - $P(X_{t+1}|e_{1:t+1}) = g(e_{t+1}, P(X_t|e_{1:t}))$

$$\begin{aligned} P(X_{t+1}|e_{1:t+1}) &= P(X_{t+1}|e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1}|X_{t+1}, e_{1:t}) P(X_{t+1}|e_{1:t}) \\ &= \alpha P(e_{t+1}|X_{t+1}) P(X_{t+1}|e_{1:t}) \\ &= \alpha P(e_{t+1}|X_{t+1}) \sum_{x_t} P(x_t|e_{1:t}) P(X_{t+1}|x_t, e_{1:t}) \\ &= \alpha P(e_{t+1}|X_{t+1}) \sum_{x_t} P(x_t|e_{1:t}) P(X_{t+1}|x_t) \end{aligned}$$

Apply conditional independence

# Filtering algorithm

$$P(X_{t+1} | e_{1:t+1}) = \alpha P(e_{t+1} | X_{t+1}) \sum_{x_t} P(x_t | e_{1:t}) P(X_{t+1} | x_t)$$



$$\begin{aligned}
 P(X_{t+1} | e_{1:t+1}) &= \underbrace{\alpha}_{\text{chain law}} \frac{P(e_{t+1} | e_{1:t})}{P(X_{t+1} | e_{1:t})} \cdot P(X_{t+1}, e_{t+1} | e_{1:t}) \\
 &= \alpha P(X_{t+1}, e_{t+1} | e_{1:t}) \xrightarrow{\text{known}} \\
 &= \alpha P(e_{t+1} | e_{1:t}, X_{t+1}) \underbrace{P(X_{t+1} | e_{1:t})}_{\text{independent Markov}} \\
 &= \alpha P(e_{t+1} | X_{t+1}) \underbrace{P(X_{t+1} | e_{1:t})}_{\sum_{X_t} P(X_t | e_{1:t}) \cdot P(X_{t+1} | X_t)} \\
 &\quad \cancel{\downarrow X_t | e_{1:t}}
 \end{aligned}$$

inde

# Filtering algorithm

$$\boxed{P(X_{t+1} | e_{1:t+1}) = \alpha P(e_{t+1} | X_{t+1}) \sum_{x_t} P(x_t | e_{1:t}) P(X_{t+1} | x_t)}$$

Normalize

Update

Predict



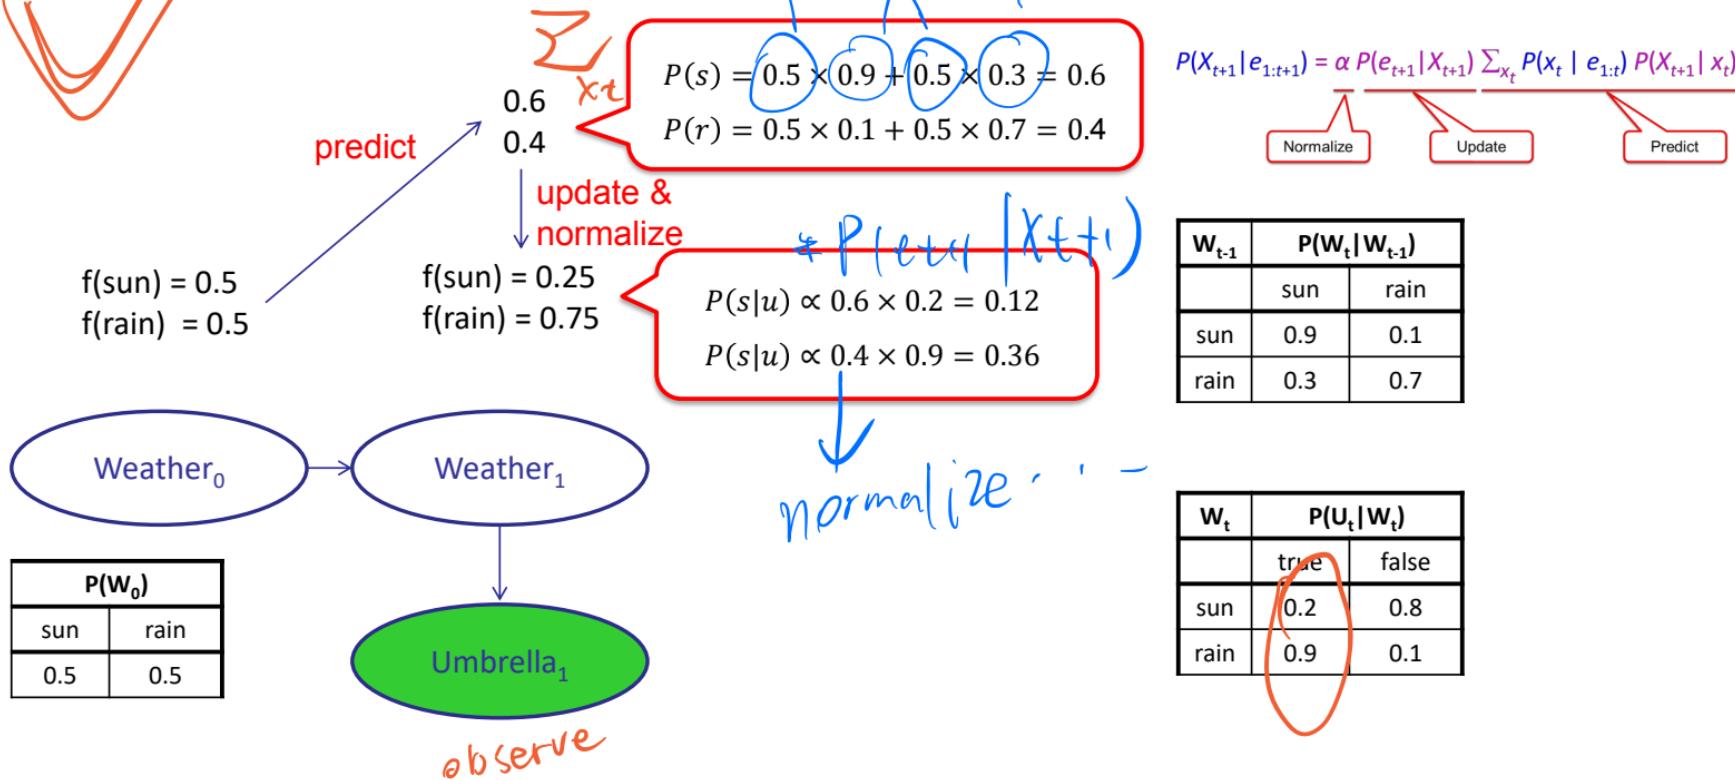
$$\boxed{f_{1:t+1} = \text{FORWARD}(f_{1:t}, e_{t+1})}$$

We start with  $f_{1:0} = P(X_0)$  and then iterate

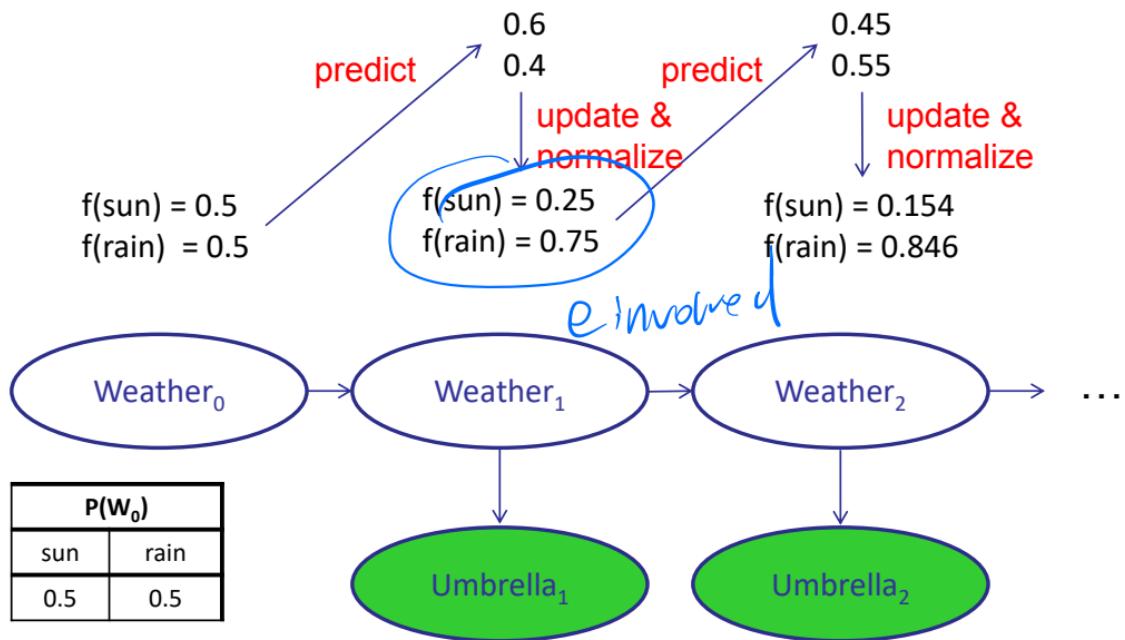
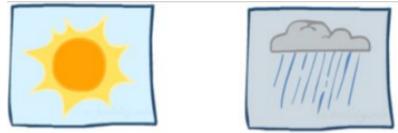
Cost per time step:  $\underline{\mathcal{O}(|X|^2)}$  where  $|X|$  is the number of states

上车的概率 (正果)  
 $(x_t | e_{1:t}) \quad P(X_{t+1} | x_t)$

# Example: Weather HMM



# Example: Weather HMM



$$P(X_{t+1} | e_{1:t+1}) = \alpha P(e_{t+1} | X_{t+1}) \sum_{x_t} P(x_t | e_{1:t}) P(X_{t+1} | x_t)$$

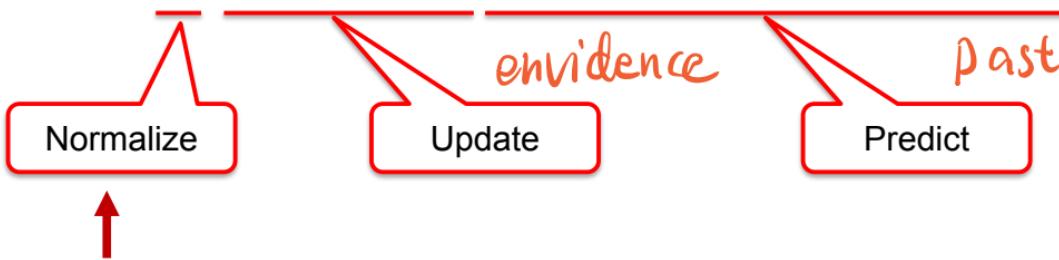
Normalize      Update      Predict

$W_{t-1}$	$P(W_t   W_{t-1})$	
	sun	rain
sun	0.9	0.1
rain	0.3	0.7

$W_t$	$P(U_t   W_t)$	
	true	false
sun	0.2	0.8
rain	0.9	0.1

# Filtering algorithm

$$\bullet P(X_{t+1} | e_{1:t+1}) = \alpha P(e_{t+1} | X_{t+1}) \sum_{x_t} P(x_t | e_{1:t}) P(X_{t+1} | x_t)$$

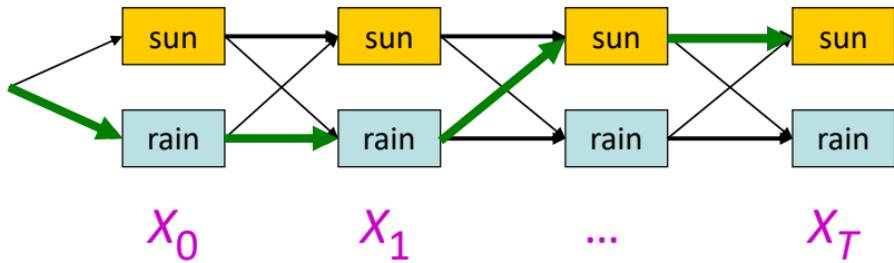


$\alpha$  is a constant. So if we only want to compute  $P(x_t | e_{1:t})$ , then we can skip normalization when computing  $P(x_1 | e_1), P(x_2 | e_{1:2}), \dots, P(x_{t-1} | e_{1:t-1})$

final normalization  
Can just do

# Another view of the algorithm

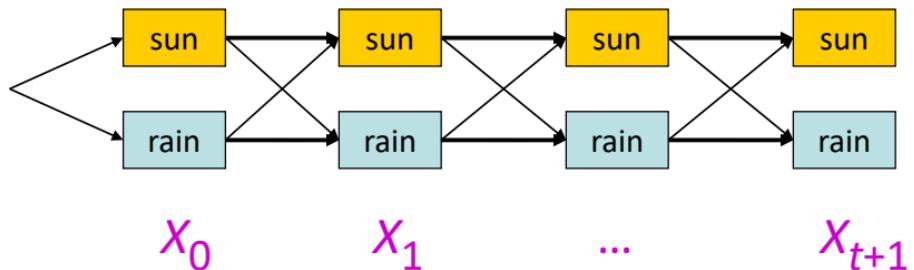
- **State trellis**: graph of states and transitions over time



- Each arc represents some transition  $x_{t-1} \rightarrow x_t$
- Each arc has weight  $P(x_t | x_{t-1}) P(e_t | x_t)$  (arcs to initial states have weight  $P(x_0)$ )
- Each path is a sequence of states
- The **product** of weights on a path is proportional to that state sequence's probability

One path:  $P(x_0) \prod_t P(x_t | x_{t-1}) P(e_t | x_t) = \underbrace{P(x_{1:t}, e_{1:t})}_{\text{path}} \propto \underbrace{P(x_{1:t} | e_{1:t})}_{\text{filtering}}$

# Another view of the algorithm



- Forward algorithm computes sum over all possible paths

$$P(X_{t+1} | e_{1:t+1}) = \sum_{x_{1:t}} P(x_{1:t+1} | e_{1:t+1})$$

- It uses dynamic programming to sum over all paths

- For each state at time  $t$ , keep track of the total probability of all paths to it

$$\mathbf{f}_{1:t+1} = \text{FORWARD}(\mathbf{f}_{1:t}, e_{t+1})$$

$$= \alpha P(e_{t+1} | X_{t+1}) \sum_{x_t} P(X_{t+1} | x_t) \mathbf{f}_{1:t}[x_t]$$

DP :

# Most Likely Explanation

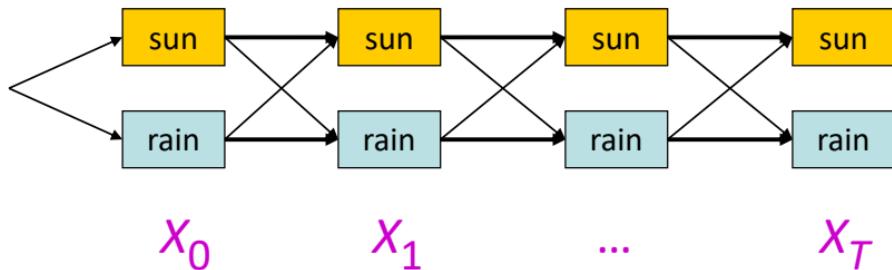


# Inference tasks

- **Filtering:**  $P(X_t | e_{1:t})$ 
  - **belief state**—input to the decision process of a rational agent
- **Prediction:**  $P(X_{t+k} | e_{1:t})$  for  $k > 0$ 
  - evaluation of possible action sequences; like filtering without the evidence
- **Smoothing:**  $P(X_k | e_{1:t})$  for  $0 \leq k < t$ 
  - better estimate of past states, essential for learning
- **Most likely explanation:**  $\arg \max_{x_{0:t}} P(x_{0:t} | e_{1:t})$ 
  - speech recognition, decoding with a noisy channel

# Most likely explanation = most probable path

- **State trellis**: graph of states and transitions over time



- The **product** of weights on a path is proportional to that state sequence's probability

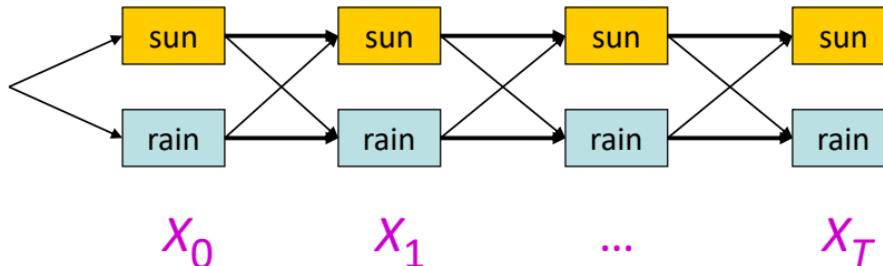
$$P(x_0) \prod_t P(x_t | x_{t-1}) P(e_t | x_t) = P(x_{0:t}, e_{1:t}) \propto P(x_{0:t} | e_{1:t})$$

- **Viterbi algorithm** computes best paths

$$\arg \max_{x_{0:t}} P(x_{0:t} | e_{1:t})$$

finding most possible path

# Forward / Viterbi algorithms



## Viterbi Algorithm (max)

For each state at time  $t$ , keep track of the **maximum probability of any path** to it

$$\mathbf{m}_{1:t+1} = \text{VITERBI}(\mathbf{m}_{1:t}, \mathbf{e}_{t+1})$$

$$= P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \max_{x_t} P(X_{t+1} | x_t) \mathbf{m}_{1:t}[x_t]$$

〃

$$\text{sun} | \begin{matrix} \text{sun} = 0.5 \\ \text{rain} = 0.2 \end{matrix} \quad 0.5 \times 0.2$$

## Forward Algorithm (sum)

For each state at time  $t$ , keep track of the **total probability of all paths** to it

$$\mathbf{f}_{1:t+1} = \text{FORWARD}(\mathbf{f}_{1:t}, \mathbf{e}_{t+1})$$

$$= \alpha P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \sum_{x_t} P(X_{t+1} | x_t) \mathbf{f}_{1:t}[x_t]$$

# Viterbi algorithm contd.

$P(W_0)$	
sun	rain
0.5	0.5

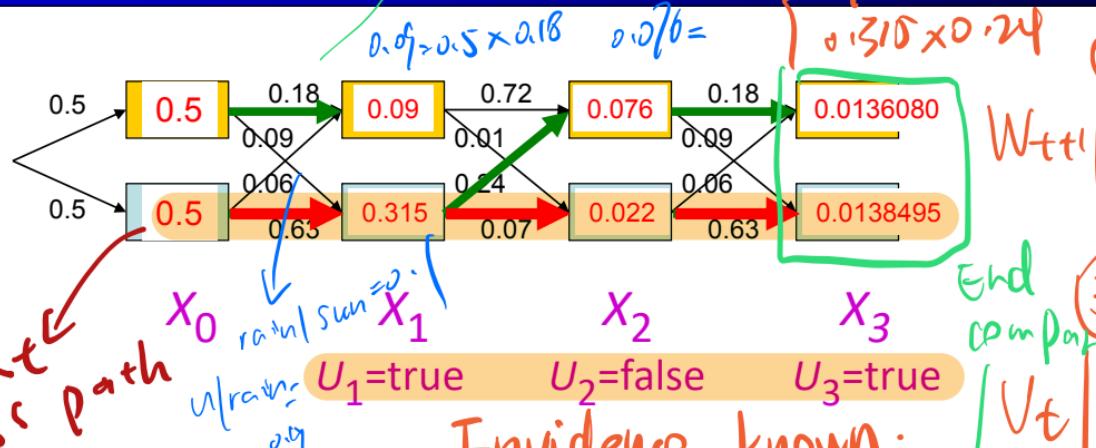
① init

most pos<sup>c</sup> path

$$m_{1:t+1} = P(e_{t+1}|X_{t+1}) \max_{x_t} P(X_{t+1}|x_t) m_{1:t}[x_t]$$

- Time complexity:  $O(|X|^2 T)$
- Space complexity:  $O(|X| T)$

e.g.: find most possible path to both sun / rain at  $X_1$



Evidence known

$$\text{weight} = P(X_t|X_{t-1}) P(e_t|X_t)$$

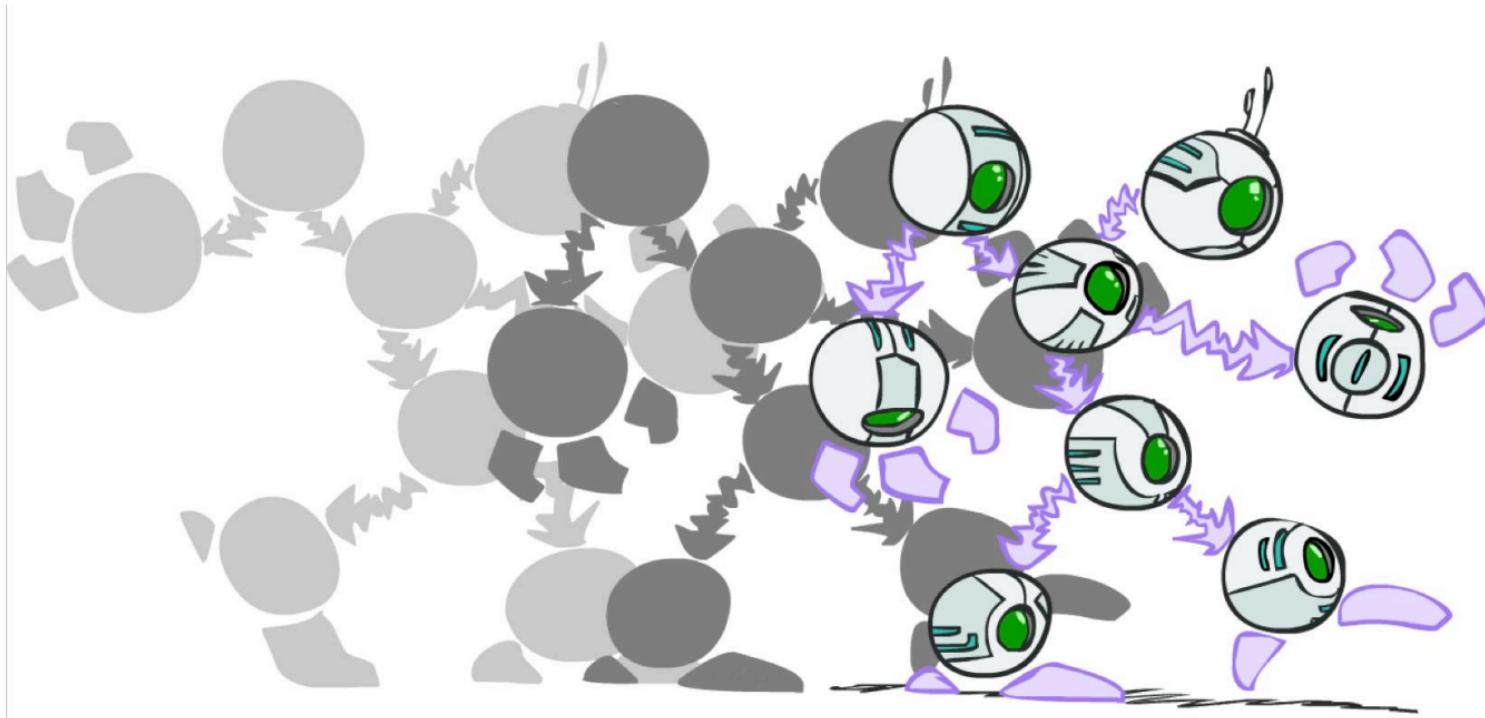
T T T

at  $X_1$   
sun / rain.

DBN

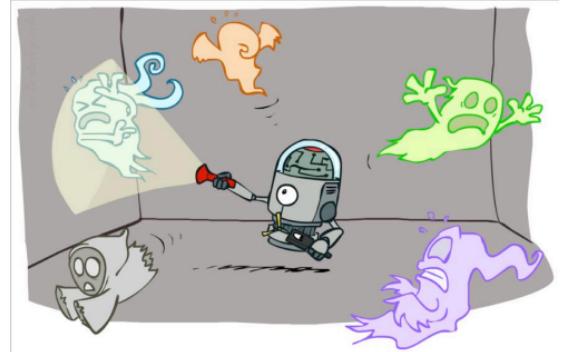
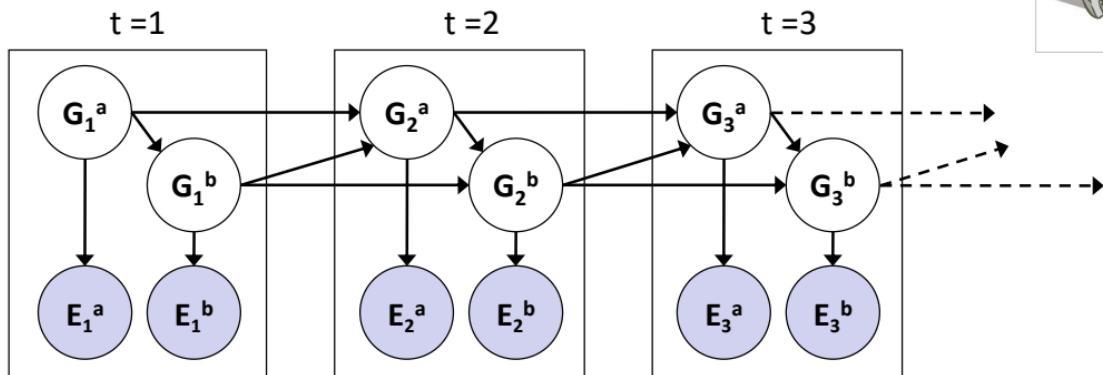
# Dynamic Bayes Nets

(intro)



# Dynamic Bayes Nets (DBNs)

- We want to track multiple variables over time, using multiple sources of evidence
- Idea: Repeat a fixed Bayes net structure at each time
- Variables from time  $t$  can condition on those from  $t-1$

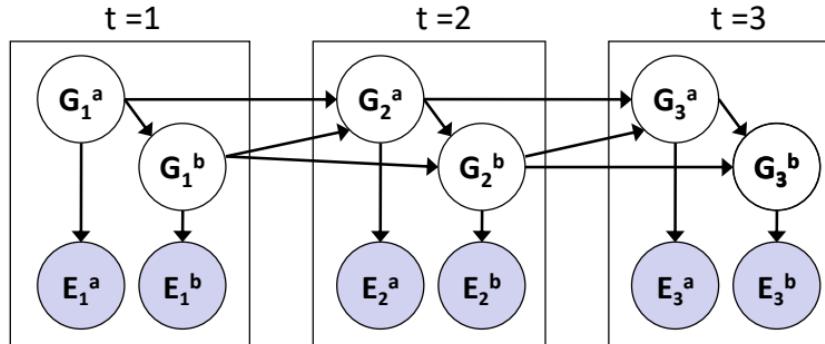


# DBNs and HMMs

- Every HMM is a DBN
- Every discrete DBN can be represented by a HMM
  - Each HMM state is Cartesian product of DBN state variables
    - E.g., 3 binary state variables => one state variable with  $2^3$  possible values
  - Advantage of DBN vs. HMM?
    - Sparse dependencies => exponentially fewer parameters
    - E.g., 20 binary state variables, 2 parents each;  
DBN has  $20 \times 2^{2+1} = 160$  parameters, HMM has  $2^{20} \times 2^{20} = \sim 10^{12}$  parameters

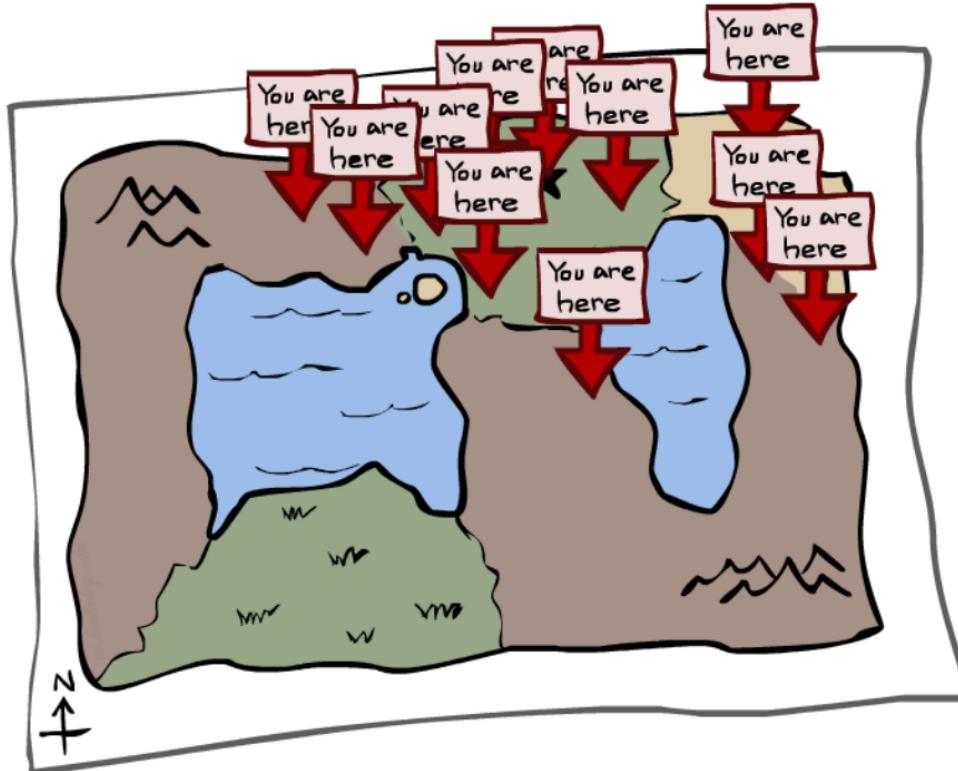
# Exact Inference in DBNs

- Variable elimination applies to dynamic Bayes nets
- Offline: “unroll” the network for  $T$  time steps, then eliminate variables to find  $P(X_T | e_{1:T})$ 
  - Problem: results in very large BN



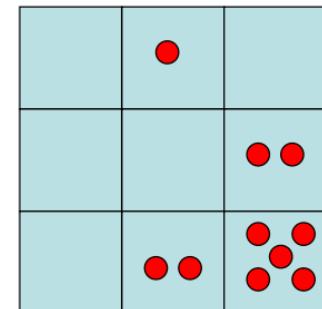
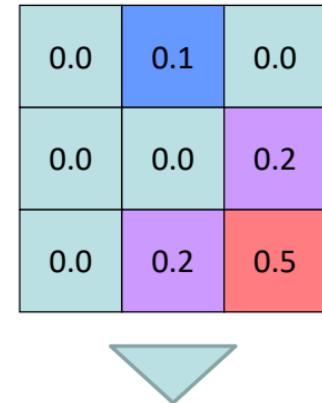
- Can we do better?
  - Do we need to unroll for many steps? What is the best variable order of elimination?
- Online: unroll as we go, **eliminate all variables from the previous time step each time**

# Particle Filtering



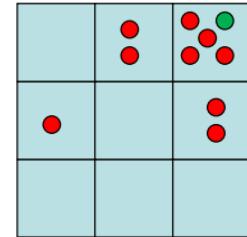
# Particle Filtering

- Filtering: approximate solution
- Sometimes  $|X|$  is too big to use exact inference
- Solution: approximate inference
  - Track samples of  $X$ , not all values
  - Samples are called **particles**
- Our representation of  $P(X)$  is now a list of **N** particles (samples)
  - $P(x)$  approximated by number of particles with value  $x$ 
    - So, many  $x$  may have  $P(x) = 0$
  - Generally,  $N \ll |X|$ 
    - More particles, more accuracy; but a large  $N$  would defeat the point.



# Representation: Particles

- At first, all particles have a weight of 1



Particles:

(3,3)

(2,3)

(3,3)

(3,2)

(3,3)

(3,2)

(1,2)

(3,3)

(3,3)

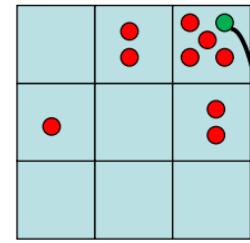
(2,3)

# Particle Filtering: Propagate forward

- Each particle is moved by sampling its next position from the transition model:
  - $x_{t+1} \sim P(X_{t+1} | x_t)$
- This captures the passage of time
  - If enough samples, close to exact probabilities (consistent)

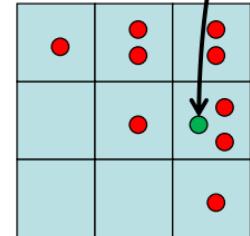
Particles:

- (3,3)
- (2,3)
- (3,3)
- (3,2)
- (3,3)
- (3,2)
- (1,2)
- (3,3)
- (3,3)
- (2,3)



Particles:

- (3,2)
- (2,3)
- (3,2)
- (3,1)
- (3,3)
- (3,2)
- (1,3)
- (2,3)
- (3,2)
- (2,2)

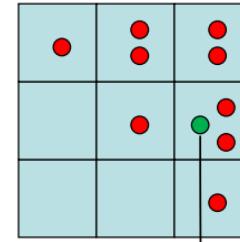


# Particle Filtering: Observe

- Similar to likelihood weighting, weight samples based on the evidence
  - $W = P(e_t | x_t)$
  - Particles that fit the evidence better get higher weights, others get lower weights
- What happens if we repeat the Propagate-Observe procedure over time?
  - It is exactly likelihood weighting (if we multiply the weights)
  - Weights drop quickly...

Particles:

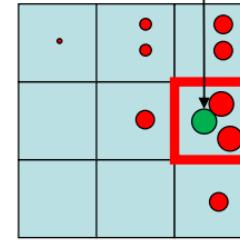
(3,2)  
(2,3)  
(3,2)  
(3,1)  
(3,3)  
(3,2)  
(1,3)  
(2,3)  
(3,2)  
(2,2)



weight

Particles:

(3,2) w=.9  
(2,3) w=.2  
(3,2) w=.9  
(3,1) w=.4  
(3,3) w=.4  
(3,2) w=.9  
(1,3) w=.1  
(2,3) w=.2  
(3,2) w=.9  
(2,2) w=.4

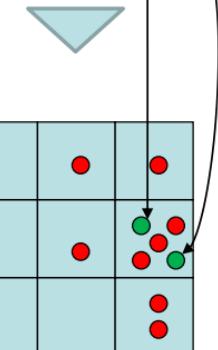
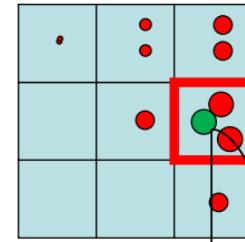


# Particle Filtering: Resample

- Rather than tracking weighted samples, we *resample*
  - Generate  $N$  new samples from our weighted sample distribution
  - Each new sample is selected from the current population of samples; the probability is proportional to its weight.
  - The new samples have weight of 1
- Now the update is complete for this time step, continue with the next one

Particles:

(3,2) w=.9  
(2,3) w=.2  
(3,2) w=.9  
(3,1) w=.4  
(3,3) w=.4  
(3,2) w=.9  
(1,3) w=.1  
(2,3) w=.2  
(3,2) w=.9  
(2,2) w=.4

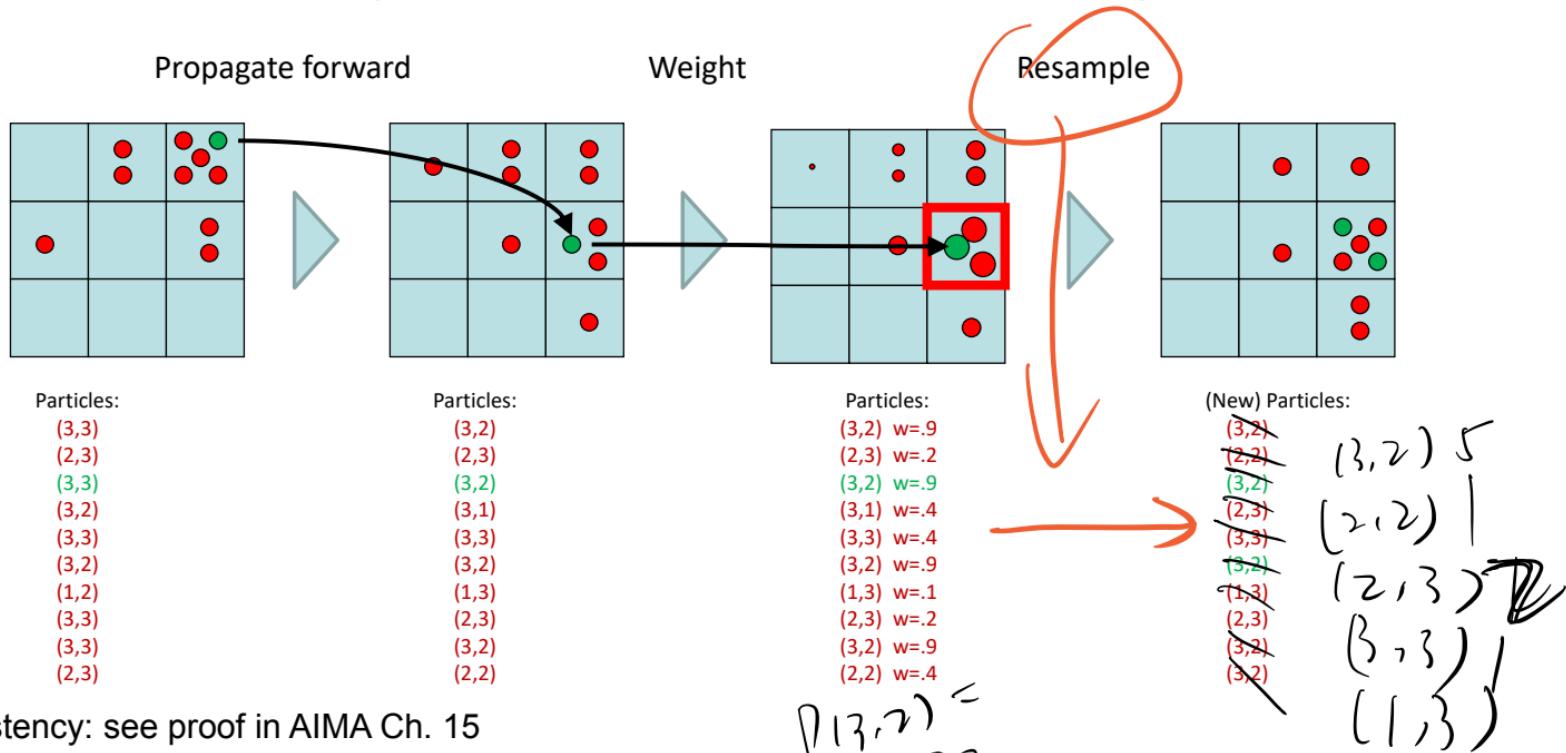


(New) Particles:

(3,2)  
(2,2)  
(3,2)  
(2,3)  
(3,3)  
(3,2)  
(1,3)  
(2,3)  
(3,2)  
(3,2)

# Summary: Particle Filtering

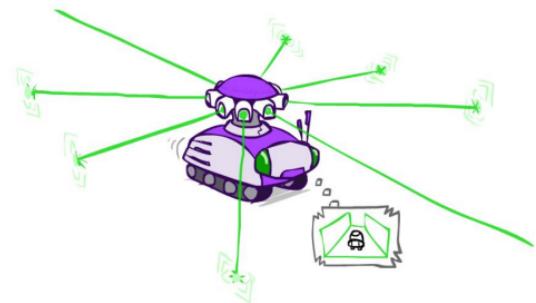
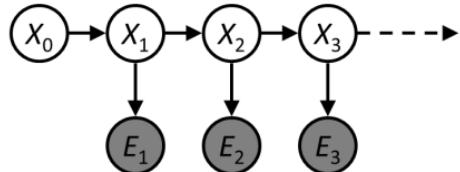
- Particles: track samples of states rather than an explicit distribution



# Summary ↴

- Probabilistic temporal models

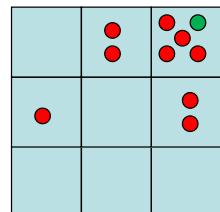
- Markov model
- Hidden Markov model
  - Filtering: forward algorithm
  - MLE: Viterbi algorithm
- Dynamic Bayesian network
- Approximate inference by particle filtering



↓ just  
for rex

# Representation: Particles

- Our representation of  $P(X)$  is now a list of  $N$  particles (samples)
  - Generally,  $N \ll |X|$
  - Storing map from  $X$  to counts would defeat the point
- $P(x)$  approximated by number of particles with value  $x$ 
  - So, many  $x$  may have  $P(x) = 0$ !
  - More particles, more accuracy
- For now, all particles have a weight of 1



Particles:  
(3,3)  
(2,3)  
(3,3)  
(3,2)  
(3,3)  
(3,2)  
(1,2)  
(3,3)  
(3,3)  
(2,3)

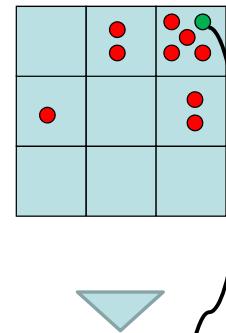
# Particle Filtering: Elapse Time

- Each particle is moved by sampling its next position from the transition model

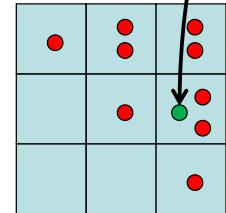
$$x' = \text{sample}(P(X'|x))$$

- This is like prior sampling – samples' frequencies reflect the transition probabilities
- Here, most samples move clockwise, but some move in another direction or stay in place
- This captures the passage of time
  - If enough samples, close to exact values before and after (consistent)

Particles:  
(3,3)  
(2,3)  
(3,3)  
(3,2)  
(3,3)  
(3,2)  
(1,2)  
(3,3)  
(3,3)  
(2,3)



Particles:  
(3,2)  
(2,3)  
(3,2)  
(3,1)  
(3,3)  
(3,2)  
(1,3)  
(2,3)  
(3,2)  
(2,2)



# Particle Filtering: Observe

- Slightly trickier:

- Don't sample observation, fix it
- Similar to likelihood weighting, downweight samples based on the evidence

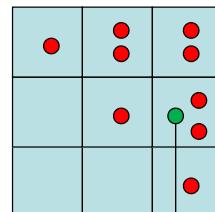
$$w(x) = P(e|x)$$

$$B(X) \propto P(e|X)B'(X)$$

- As before, the probabilities don't sum to one, since all have been downweighted (in fact they now sum to (N times) an approximation of  $P(e)$ )

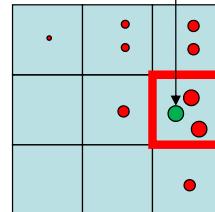
Particles:

(3,2)  
(2,3)  
(3,2)  
(3,1)  
(3,3)  
(3,2)  
(1,3)  
(2,3)  
(3,2)  
(2,2)



Particles:

(3,2) w=.9  
(2,3) w=.2  
(3,2) w=.9  
(3,1) w=.4  
(3,3) w=.4  
(3,2) w=.9  
(1,3) w=.1  
(2,3) w=.2  
(3,2) w=.9  
(2,2) w=.4

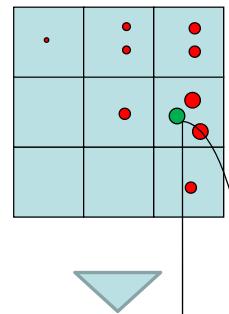


# Particle Filtering: Resample

- Rather than tracking weighted samples, we resample
- N times, we choose from our weighted sample distribution (i.e. draw with replacement)
- This is equivalent to renormalizing the distribution
- Now the update is complete for this time step, continue with the next one

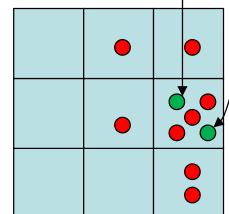
Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4



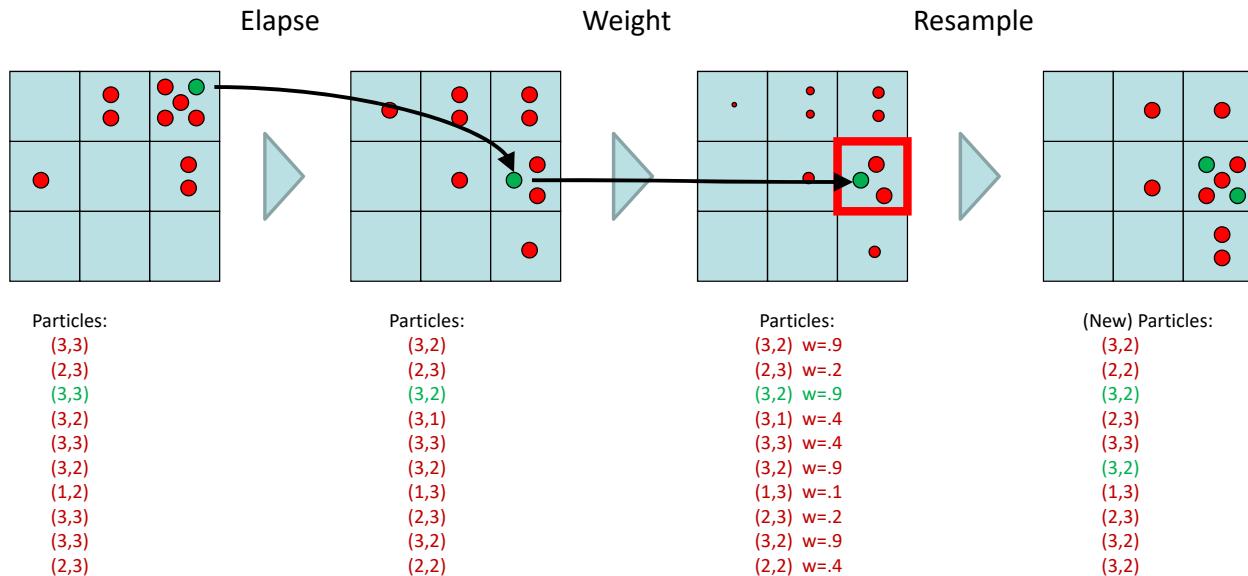
(New) Particles:

(3,2)
(2,2)
(3,2)
(2,3)
(3,3)
(3,2)
(3,2)
(1,3)
(2,3)
(3,2)
(3,2)



# Recap: Particle Filtering

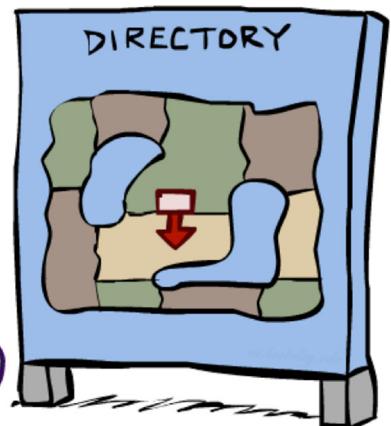
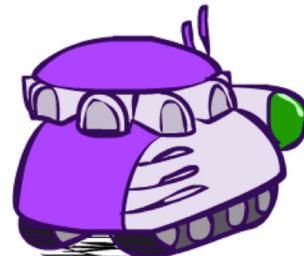
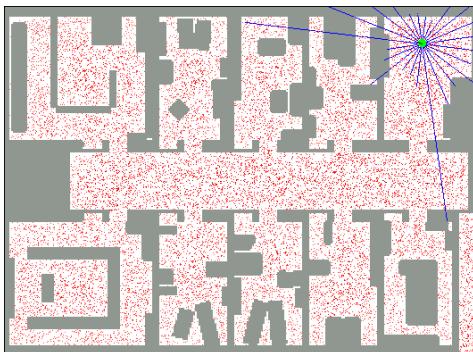
- Particles: track samples of states rather than an explicit distribution



[Demos: ghostbusters particle filtering (L15D3,4,5)]

# Robot Localization

- In robot localization:
  - We know the map, but not the robot's position
  - Observations may be vectors of range finder readings
  - State space and readings are typically continuous (works basically like a very fine grid) and so we cannot store  $B(X)$
  - Particle filtering is a main technique



# Particle Filter Localization (Sonar)



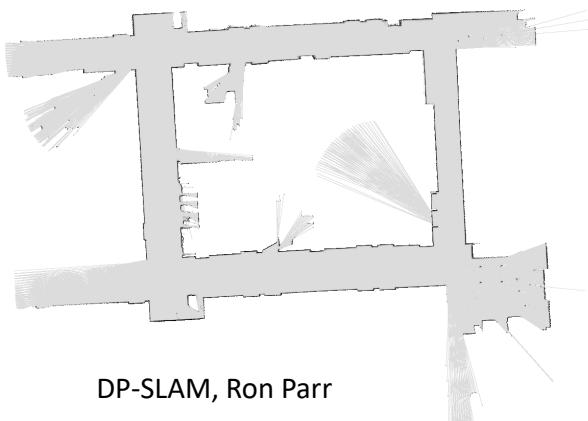
The image shows a simulation environment for particle filter localization. A grey rectangular arena contains a complex, dark grey obstacle. A cluster of red particles is scattered across the floor, with a higher density near the center. A green dot represents the estimated robot position, located near the center of the particle cloud. In the bottom-left corner, a blue rectangular box displays the number "40000".

**Global localization with  
sonar sensors**

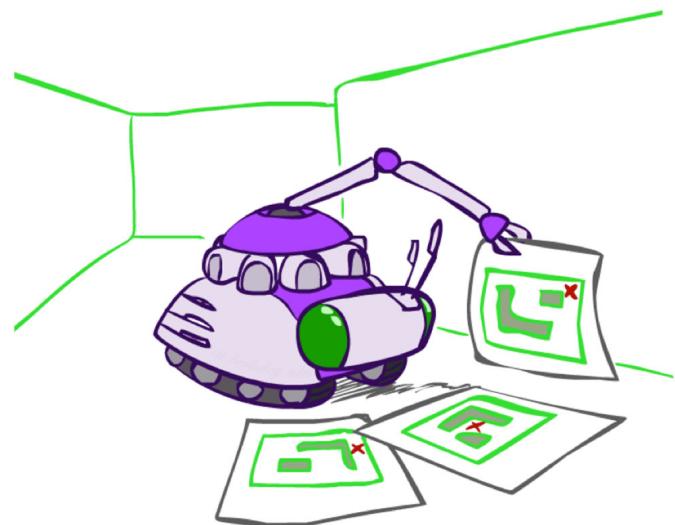
[Video: global-sonar-uw-annotated.avi]

# Robot Mapping

- SLAM: Simultaneous Localization And Mapping
  - We do not know the map or our location
  - State consists of position AND map!
  - Main techniques: Kalman filtering (Gaussian HMMs) and particle methods



DP-SLAM, Ron Parr



[Demo: PARTICLES-SLAM-mapping1-new.avi]